

OpenVGA: An Open-Source PCI Graphics Adapter

Patrick Suggate

May 1, 2009

Abstract

This thesis describes the development of OpenVGA, an open-source graphics adapter. OpenVGA uses programmable logic for its core functionality, and features a processor logic core for data processing tasks. A novel transport-triggered processor architecture, and a more traditional RISC processor, were designed and evaluated for the data processing tasks within OpenVGA.

OpenVGA is a free and open-source hardware project. The PCB artwork and the source code, written in Verilog, Assembly, Python, and C/C++, can be freely distributed and modified in accordance to the conditions of the GPL. OpenVGA was designed to use a two-layer PCB and readily-available, low-cost ICs. The hardware is therefore easy to fabricate by individuals wishing to participate in future development.

The project produced many Wishbone-compliant logic cores including: a small, fast, 16-bit TTA processor which operates up to 190 MHz when synthesised for a Spartan-3 FPGA; a 16-bit RISC processor that operates at up to 140 MHz; a data cache which has a dual-clock, 2-way set-associative architecture and operates at >150 MHz; a small, fast PCI-to-Wishbone bridge that supports multiple clock domains, by using asynchronous FIFOs; a VGA and DVI compatible display-redraw logic core; and a SDRAM controller that operates at up to 120 MHz. The Wishbone-interconnect standard allows all the logic cores to be used within other open-source projects too.

The TTA16 processor is significantly faster than existing FPGA processor cores. The RISC16 processor also has very good performance when compared with existing FPGA processors. These results show that open-source, FPGA-based, graphics adapters are feasible, and in particular that TTA processors can be used as high-performance, small-footprint general-purpose processors within FPGAs.

Acknowledgements

This thesis would have been impossible without the help and support of many people. I would like to thank my supervisor, Dr. Tim Molteno, for all his time and enthusiasm. He contributed considerable effort and advice to keep the OpenVGA project on-track.

I would also like to thank Dr. Neil Thomson for encouraging me to return to university and undertake postgraduate studies. Also a thank you to all of the other University of Otago Department of Physics staff and postgraduate students.

This project used many components developed by the open-source community and without these OpenVGA would also have been impossible. Two pieces of software were especially important for OpenVGA development, Icarus Verilog, developed by Stephen Williams, and the GtkWave waveform viewer, with the project's development led by Tony Bybell.

Contents

1	Introduction	14
1.1	Purpose of this Project	15
1.2	Relevance of OpenVGA	15
1.3	Limitations of this Project	15
1.4	Outline of this Thesis	17
2	Graphics Adapter Background	19
2.1	PC Graphics Adapter History	20
2.2	The Original IBM VGA	20
2.2.1	The Video Graphics Array	21
2.2.2	Video DAC	23
2.2.3	VGA BIOS	24
2.2.4	Display Memory	25
2.3	Non-VGA Graphics Adapters	26
2.4	Related Projects	26
2.4.1	Open Graphics Project	26
2.4.2	Project VGA	27
2.4.3	Manticore	28
3	OpenVGA Outline	29
3.1	Hardware Development	29
3.2	Important Design Elements within the FPGA	31
3.2.1	Logic Core Interconnection	32
3.2.2	FPGA-Optimised Processor Cores	32
3.2.3	Processor Data Cache	34
3.2.4	Clock Domains and Domain Crossing	34
3.2.5	Memory Controller	35
3.2.6	PCI-to-Wishbone Bridge	35

3.2.7	VGA/LCD Display Controller	35
3.2.8	Additional Logic Cores	36
4	The TTA16 and RISC16 Processor Cores	37
4.1	Processor Architectures	37
4.2	Shift Registers as Program Counters	39
4.3	Processor Performance	39
4.4	Tools & Testing	40
4.4.1	Object Code Generation	41
4.4.2	Processor Testing and Testbenches	44
4.5	TTA16	45
4.5.1	Introduction to TTA Processors	46
4.5.2	Internal Tri-States vs. Multiplexers	48
4.5.3	The TTA16 Pipeline	49
4.5.4	Functional Units	49
4.5.5	Instruction Format	56
4.5.6	Instruction Memory	57
4.6	RISC16	58
4.6.1	General RISC16 Characteristics	58
4.6.2	Design Choices	59
4.6.3	The RISC16 Pipeline	60
4.6.4	RISC16 Instruction Set	61
4.6.5	Instruction Memory	63
4.6.6	Functional Units	63
4.6.7	Data Forwarding	64
4.7	Processor Logic Cores Summary	65
5	Local Memory and Caching	67
5.1	SDRAM Controller	68
5.1.1	Controller Overview	68
5.1.2	Testing and Synthesis	71
5.2	Data Cache	74
5.2.1	Cache Justification	75
5.2.2	Implementation Considerations	76
5.2.3	Set Associativity	77
5.2.4	Cache Size	79
5.2.5	Cache Line Size	80

5.2.6	Replacement Policy	81
5.2.7	Write Policy	81
5.2.8	Pipeline Length	82
5.2.9	Testing	83
5.3	DMA Controller	84
5.4	Serial PROM	84
5.4.1	Serial PROM Read-Back Considerations	85
6	I/O Logic Cores and Data Synchronisation	86
6.1	Clock Architecture	86
6.1.1	Spartan-3 Clocking Resources	86
6.1.2	Clock Domains	88
6.1.3	Asynchronous FIFOs	90
6.2	PCI-to-Wishbone Bridge	94
6.2.1	PCI-to-Wishbone Bridge Logic Core	94
6.2.2	PCI Testing and Summary	96
6.3	VGA and DVI Controller	97
6.3.1	Display Modes	98
6.3.2	Pixel Data Prefetch	99
6.3.3	CRT Controller	101
6.3.4	VGA Output	101
6.3.5	DVI Output	101
6.4	Miscellaneous Peripherals	102
6.4.1	LEDs	102
6.4.2	USB Serial Port	102
7	Summary, Conclusions, and Future Work	103
7.1	Summary of the OpenVGA Project	103
7.2	OpenVGA Logic Cores	104
7.2.1	Summary of Important OpenVGA Logic Cores	104
7.2.2	Processor Comparison and Summary	107
7.2.3	Known Bugs	108
7.3	Future Work	108
7.3.1	VGA Compatibility	109
7.3.2	Software	109
7.3.3	Hardware Improvements	109
7.3.4	Logic Core Improvements	110

7.4	Conclusion	110
A	Source Code Layout and Overview	115
A.1	Top-Level	115
A.2	HDL Hierarchy	115
A.3	Simulation Test-Harness Hierarchy	116
A.4	Tools	117
B	OpenVGA Components and PCB	118
B.1	Off-Board Connections	119
B.2	Integrated Circuits	119
B.2.1	The Xilinx Spartan-3 FPGA and SPROM	119
B.2.2	The DVI and VGA Video Encoders	120
B.2.3	USB UART	120
B.2.4	PCI Bus Switches	120
C	Wishbone Interconnect Overview	122
C.1	Wishbone Interconnect Introduction	122
C.1.1	Features	122
C.1.2	A Simple Example	123
C.2	Description of Wishbone Signals	123
D	TTA16 Assembly Language Programming Guide	126
D.1	TTA16 Overview	126
D.1.1	TTA16 Datatypes	127
D.1.2	Instruction Format	127
D.1.3	Data Transport Scheduling	129
D.1.4	Exceptions	131
D.2	A Simple Example	131
D.2.1	Functional Units	133
D.3	OpenVGA Memory-Mapped I/O Peripherals	134
D.3.1	Cache Flush Peripheral	135
D.3.2	CRT Controller	135
D.3.3	DMA Unit	136
D.4	Assembly Coding Conventions	137
D.5	Executing the Generated Object Code	138
D.5.1	TTA16 Programming Example: ‘memcpy’	138

E	RISC16 Assembly Language Programming Guide	141
E.1	Tools: The RISC16 Assembler	141
E.2	Instruction Format Overview	142
E.2.1	RISC16 Registers	144
E.2.2	ALU Instructions	144
E.2.3	Processor Condition Code Flags	145
E.2.4	Multiply Instructions	145
E.2.5	Program Flow Control	146
E.3	Loads and Stores	146
E.3.1	Memory Segment Registers	147
E.3.2	OpenVGA MMIO Device Segments	147
E.4	Programming Example: ‘memcpy’	147

List of Tables

2.1	VGA I/O Ports	25
4.1	A simple addition operation to demonstrate the TTA concept	47
4.2	TTA Accumulate Example	48
4.3	Synthesis results comparing multiplexers and internal pull-ups	49
4.4	RISC16 addition operation	59
4.5	Processor properties	66
5.1	Cache size vs. performance	80
5.2	Cache line-size vs. performance	80
5.3	DEC Alpha 21264 cache eviction policy vs. performance	81
5.4	Cache eviction policy vs. performance	81
5.5	Cache Write Policies vs. Performance	82
6.1	OpenVGA clock frequencies	88
6.2	Comparison of Decimal, Binary, and Gray Encoding.	92
6.3	OpenVGA video modes	100
7.1	Summary of significant OpenVGA logic cores	105
7.2	Evaluated processors and a comparison with TTA16 and RISC16	108
B.1	OpenVGA integrated circuits	119
D.1	TTA16 registers and the transports to which they have access	129
D.2	Functional unit summary	133
D.3	TTA16 memory-mapped I/O modules and segments	134
D.4	CRTC ports and their defaults values	136
D.5	Register usage conventions	137
E.1	ALU instructions that can optionally set condition codes	145
E.2	RISC16 branch and jump instructions	146

List of Figures

2.1	ATi Rage graphics adapter	19
2.2	VGA display redraw timing and signals	22
2.3	ASCII and extended-ASCII character set	24
2.4	OpenGraphics OGD1	27
3.1	OpenVGA Hardware Features	30
3.2	PCI development stub board	31
3.3	OpenVGA PCB version 2 with DDR SDRAM	32
3.4	OpenVGA Architecture Block Diagram	33
4.1	The division of responsibilities between hardware and compiler	38
4.2	An 8-bit MFSR with a cycle size of 255	39
4.3	TTA16 synthesis timing report	45
4.4	TTA16 transports and functional units overview	46
4.5	TTA16 pipeline overview	49
4.6	The PC increment and branch unit	50
4.7	3-Bit, Bitwise Operations Unit	53
4.8	TTA16 instruction format	56
4.9	RISC16 pipeline	61
4.10	RISC16 instruction formats	62
4.11	Data Dependency.	64
4.12	Data Bypassing Scheme Used by RISC16	65
5.1	OpenVGA memory hierarchy	68
5.2	Wishbone Clock Domain State Machine	69
5.3	Spartan-3 DDR IOB Primitive	71
5.4	A simulation of a PCI to SDRAM data transfer	72
5.5	SDRAM Memory Test Results	73
5.6	Text-mode simulator without a cache	78

5.7	Text-mode simulator with a cache	78
5.8	Simplified Schematic of the Cache Sense Logic	83
6.1	OpenVGA pixel clock generation	88
6.2	Single-bit synchronisers	89
6.3	Simplified Asynchronous FIFO Schematic	91
6.4	PCI-to-Wishbone bridge block diagram	94
6.5	The PCI bridge state machines	95
6.6	PCI configuration space write transaction	97
6.7	Video controller block diagram	98
6.8	Emulated text-mode display data produced by a simulation	99
B.1	OpenVGA PCB artwork, both sides	118
B.2	OpenVGA solder-side components	120
C.1	A Wishbone Point-to-Point Connection Scheme	123
C.2	A Sample Wishbone Transaction	124
D.1	TTA16 instruction format	128
D.2	TTA16 instruction shared bit-fields	130
D.3	C ‘memcpy’ Example	139
D.4	TTA16 Assembly ‘memcpy’ Example	140
E.1	RISC16 instruction formats	142

Glossary

ALU	Arithmetic Logic Unit, 40
ASCII	American Standard Code for Information Inter- change, 10
ASIC	Application Specific Integrated Circuit, 7
BGA	Ball Grid Array, 8
BIOS	Basic Input/Output System, 13
BRAM	Xilinx Block RAM, 38
CISC	Complex Instruction Set Computer, 24
CPI	Clocks Per Instruction, 52
CRT	Cathode Ray Tube, 10
CRTC	CRT Controller, 11
DAC	Digital to Analogue Converter, 18
DDR	Double Data Rate, 20
DFF	D-type Flip-Flop, 22
DMA	Direct Memory Access, 46
DRAM	Dynamic Random Access Memory, 8
DVI	Digital Visual Interface, 8
EDA	Electronic Design Automation, 29
FIFO	First-In, First-Out queue, 22
FOSS	Free and Open-Source Software, 7
FPGA	Field-Programmable Gate Array, 7
FU	Functional Unit, 24

GCC	GNU Compiler Collection, 39
GPL	GNU General Public License, 7
GUI	Graphical User Interface, 15
hsync	Horizontal synchronisation, 11
I/O	Input/Output, 8
IBMTM	International Business Machines, 7
ILP	Instruction Level Parallelism, 27
IOB	Input/Output Block, 23
ISA	Instruction Set Architecture, 24
JTAG	Joint Test Action Group, 82
LSB	Least Significant Bit, 71
LUT	Look-Up Table, 25, 52
MDA	Monochrome Display Adapter, 10
MFSR	Multiple Feedback Shift Register, 25
MMIO	Memory-Mapped Input/Output, 23
MSB	Most Significant Bit, 64
NRE	Non-Recurring Engineering, 7
opcode	Condensed form of operation-code, 28
OpenGL	The Open Graphics Library, 15
OpenVGA	The free and open-source hardware graphics adapter created by this project, 7
OS	Operating System, 7
PC	Personal Computer, 7
PC	Program Counter, 9, 25
PCB	Printed Circuit Board, 7
PCI	Peripheral Component Interconnect, 8

pixel	A condensed form of the words “picture element”, 8
PLD	Programmable Logic Device, 7
RF	Register File, 31
RGB	Red Green Blue, 12
RISC	Reduced Instruction Set Computer, 21
RISC16	The 16-bit RISC architecture processor developed for this project, 8
RTL	Register Transfer Level, 7
SDR	Single Data Rate, 47
SDRAM	Synchronous Dynamic RAM, 8
SoC	System-on-a-Chip, 21
SPROM	Serial Programmable Read Only Memory, 20
SRAM	Static RAM, 7
SVGA	Super VGA, 11
TTA	Transport Triggered Architecture, 9
TTA16	The 16-bit TTA processor developed for the project, 8
VCD	Value Change Dump, 29
VGA	Video Graphics Array, 7
VLIW	VHSIC (Very Long Instruction Word, 24
vsync	Vertical synchronisation, 11
x86	The general name for a computer architecture based upon the Intel 8086 architecture and derivatives, 13
XST	Xilinx Synthesis Tool, 35

Chapter 1

Introduction

This thesis presents the design and implementation of a FPGA-based, free and open-source hardware, graphics adapter. Free and Open-Source Software (FOSS) has achieved widespread use and acceptance [27] though open-source designs of the underlying hardware which comprise the systems on which the software is executed have not. While the interfaces between many components have been standardised, the implementation of these components is often proprietary, and considered a trade secret.

Graphics adapters are an example of proprietary hardware. They have to be compatible with the VGA (Video Graphics Array), or an earlier IBM specification, to function as a PC's primary adapter. This specification has never been released, all current implementations are proprietary. This leads to significant barriers for those wishing to create their own, whether for research or commercial purposes.

By implementing a design within programmable logic, such as an FPGA, it is possible for organisations, and even individuals, to design and implement complex digital circuits without having to pay the significant Non-Recurring Engineering costs to have an Application Specific Integrated Circuit (ASIC) fabricated. By using logic cores, advanced digital circuits can be constructed within Programmable Logic Devices (PLDs) by combining many smaller cores. With the advent of high logic capacity, low-cost FPGA families¹, and with freely available HDL EDA tools, the barriers to designing and implementing advanced logic cores are now significantly lower. And like FOSS, it is possible for the Register Transfer Level (RTL, the most popular method for describing digital logic circuits in HDL) descriptions of logic cores to be made freely available, resulting in free and open-source hardware designs.

¹High-capacity, low-cost FPGA product ranges that are readily available at the time of writing are Xilinx Spartan, Altera Cyclone, and Lattice EC, and other vendors have similar products too.

1.1 Purpose of this Project

The purpose of this project, OpenVGA, is to use low-cost FPGAs to implement an open-source, FPGA-based, computer graphics adapter. The HDL, driver and firmware source-code, and PCB design are to be freely available. The objective is to lower the obstacles that exist for developing PC display adapters. Additional goals for OpenVGA are to:

- Develop HDL logic cores and implement a very flexible display adapter, along with a collection of tools and testbenches. The design should allow a useful subset of the VGA specification to be emulated in the future.
- Produce Operating System (OS) device drivers so that OpenVGA can be recognised by the OS, and used by software applications.
- Use commonly available components to allow a moderately low-cost adapter to be produced.
- Have all source code and PCB artwork generated by this project released under the GNU General Public License (GPL), and available on the Internet. This will allow others to modify and extend OpenVGA, and contribute back the changes as well.

1.2 Relevance of OpenVGA

Since OpenVGA is an open graphics development platform, and with a small logic-use footprint, it could be ideal as the basis for other graphics projects (like a real-time ray tracer [38] or for data visualisation). It could possibly even serve as a low-cost development board for a hardware computation project, since it contains a FPGA, a Peripheral Component Interconnect (PCI) Local Bus connector, and some DRAM.

At the time of writing, the availability of high-quality, open-source logic cores is limited², especially when compared to the number of FOSS projects. OpenVGA adds to this pool of logic cores as it contains logic cores which can be of use to other projects. Possibly useful logic cores are: a couple of small and fast processors, TTA16 and RISC16; a simple data cache; a Synchronous Dynamic Random Access Memory (SDRAM) controller; a simple PCI to Wishbone bridge; and VGA and DVI redraw logic cores.

1.3 Limitations of this Project

OpenVGA was not intended to have advanced features and compete with graphics adapters developed by the major vendors; currently Intel, AMD, and Nvidia. These companies employ hundreds

²OpenCores seems to be the largest repository of open-source logic cores, and it hosts many interesting projects, but these are of variable quality and many more cores are needed.

of engineers to design and test their graphics adapters. These modern graphics adapters are extremely sophisticated, containing many millions of logic gates, and have 2D and 3D hardware acceleration functionality³. Programmable logic does not allow designs as large, or that can achieve the same operating frequencies, as those implemented as ASICs. OpenVGA is designed to be a simple, open-source, graphics adapter instead.

At the time of writing, OpenVGA cannot function as a PC's primary graphics adapter, since it would need to be compatible with the VGA specification. Meeting the VGA specification was beyond the scope of this project, though OpenVGA was designed to support VGA emulation, and this is a future goal of the open-source OpenVGA project. Already though, once an OS is loaded, OpenVGA can be accessed using a software device driver.

Another limitation of OpenVGA is that there are only two frequencies available as the dot-clock for the display. This is the clock used to generate display timing signals and draw pixels⁴ to the screen. This limits the range of supported screen resolutions. Due to the Spartan-3 architecture, an external clock generator IC would be needed to allow for a large range of supported display resolutions.

To meet the design objectives of low-cost and easy to fabricate meant using a two-layer PCB. This limited data-bus widths and frequencies between the Spartan-3 FPGA and its associated peripherals. The memory bus-width was constrained to 16-bits because of the limited number of Input/Output (I/O) connections available on non-Ball Grid Array (BGA) Spartan-3 FPGAs. Xilinx recommends a PCB with six or more layers for its BGA-packaged FPGAs.

The two-layer PCB also reduced performance within the high-frequency sections of the design as well. For example, the maximum stable SDRAM operating frequency achieved by the memory hardware testbench was 120 MHz, though the manufacturer specifies 133 MHz for the SDRAM IC used. This was due to limited options for placing ground and power planes, and also placing termination resistors is more difficult. The result was that signal integrity issues limited the speed of some signals to below those listed in the manufacturer's specifications.

Due to this combination of lower operating frequencies and narrower bus widths, compared with current commercial graphics adapters, OpenVGA has a relatively low memory bandwidth (the peak measured during testing was 240 MB/s). This limits available display modes and many possibilities for hardware acceleration⁵.

³A commercial graphics adapter like an AMD Radeon™HD 4890 claims to have floating-point calculation performance of over one teraFLOP, and can render millions of triangles per second, accelerate the play-back of encoded video streams, and has a memory data-bus that is 256-bits wide [11].

⁴Pixel is a simplification of the two words "picture element".

⁵There is an FPGA processor logic core implemented as part of the OpenVGA design. This is for data processing tasks and managing the state of the graphics adapters (setting video modes and initialisation). This processor is pipelined and operates at 150 MHz, features instruction-level parallelism, and has a high-speed data cache. A future version of OpenVGA firmware will use this processor for some 2D acceleration tasks as well as emulating VGA.

1.4 Outline of this Thesis

A review of PC graphics adapters, both past specifications and current related projects, is presented in Chapter 2. This introduces graphics adapters developed for the computer architectures derived from the original IBM PC. Included is a brief overview of the history and development of PC graphics adapters up until VGA. These adapters share many functional components, with each succeeding adapter generation building on the specifications of previous one. These functional components, and their operation, are detailed. This chapter then concludes with a review of other open-source graphics adapter projects.

The architecture of the graphics adapter presented within this thesis, OpenVGA, is introduced in Chapter 3. Topics covered are the OpenVGA hardware and an introduction to the logic cores that were developed. These logic cores are used to provide the functionality within the FPGA, which is the core component of this graphics adapter.

OpenVGA can use one of two custom, FPGA-based processors for initialisation, mode management, and data processing tasks. Chapter 4 first presents an overview of processor architectures and other processor design topics. Included is a discussion of a novel technique for incrementing a processor's Program Counter (PC), and an uncommon, but powerful, class of processor architecture, the Transport Triggered Architecture (TTA). The body of this chapter then covers the development of two processors with radically different architectures, TTA16 and RISC16. Concluding the chapter is a comparison of these two processors and how well suited they are for OpenVGA.

Graphics adapters have local memory and Chapter 5 presents the logic cores for accessing and caching this local memory. The local memory is used for storing display data, firmware, and adapter state information. OpenVGA has both ROM and RAM and a cache is used to give the OpenVGA processor fast, low-latency access to this memory. The design of the DMA controller used within OpenVGA is discussed as well. It was developed so that the processor can then write data back to the RAM using efficient burst transfers.

Interfacing the significant functional components of OpenVGA, both within the FPGA and the FPGA to external interfaces, the numerous modules needed for this, and drawing to the display, are the topics of Chapter 6. Since different components operate at different clock frequencies, data synchronisation problems had to be solved and the solutions are presented in Section 6.1.

The project summary and conclusions are discussed in Chapter 7. This is a summary of the current status of OpenVGA and what has been achieved, it includes comparisons of the two processors developed and how they compared to existing processors. Finally, areas for future work are covered, including topics such as software drivers, firmware, logic core improvements, DVI support and testing, and upgraded hardware.

This thesis includes several appendices that cover important information that is related to this

project. There is a source code overview in Appendix A. The code described here is available from an Internet open-source software repository SourceForge (<http://openvga.sourceforge.net/>). The hardware components and PCB artwork is included in Appendix B. The Wishbone interconnect is described in Appendix C, and Appendix D and Appendix E, are the assembly programming guides for the TTA16 and RISC16 processors.

Chapter 2

Graphics Adapter Background

To begin this chapter, PC graphics adapter history is explored. The history concludes with an overview of VGA as all earlier specifications are now considered obsolete. The primary graphics adapter in a PC must be VGA compatible in order for the system to boot. The main functional components of a VGA adapter are then introduced followed by a brief explanation of these components.

OpenVGA is not the first open graphics adapter project either so other similar projects are investigated. Completing this chapter, the similarities and differences with these other projects, relative to OpenVGA, are examined. Amongst others, the capable and prominent, though complex and expensive, OpenGraphics adapter is covered here.




Figure 2.1: ATi Rage graphics adapter with the significant components identified.

2.1 PC Graphics Adapter History

The original IBM PC featured a Monochrome Display Adapter (MDA) [14]. The MDA only had a very small local memory due to the high cost of DRAMs. This memory only stored ASCII (American Standard Code for Information Interchange, see Figure 2.3) and attribute data for several screenfuls of text (80 columns and 25 rows of ASCII characters). The MDA had dedicated hardware which converted ASCII characters from text-buffer into pixel values during display redraw periods. The pixel values were transmitted to a monochrome CRT monitor for display.

The original IBM PC was produced from readily available parts and was reverse engineered by other companies to produce numerous compatible versions. This means that the MDA had to be reverse engineered. Derivatives were produced that had more features and were backwards compatible, like the Hercules graphics adapter (sometimes called HGA).

The original IBM PC was the first in a line of computers, with future models having different hardware, including changes to the display adapters. The later graphics adapter standards: CGA, MCGA, EGA, VGA, 8514, XGA, and TIGA; all followed MDA. These later adapters were reverse engineered and cloned as well, though the display adapter standards after VGA never gained widespread use [14].

VGA was extended by many manufacturers, increasing framebuffer sizes, numbers of colours, higher resolution modes, and hardware acceleration for 2D and 3D functions. The umbrella term for all of these adapters was Super VGA, or SVGA. Every time a different manufacturer produced a VGA compatible display adapter, they essentially had to start from scratch, reimplementing VGA and adding their own extensions. Some of these companies were: Oak, Advance, Genoa, NCR, Integrated Info Tech, Tseng Labs, Weitek, Trident, Western Digital, Video7, VIA, SiS, Chips and Technologies, Cirrus Logic, ATI, Nvidia, Intel, S3, and 3DFX.

2.2 The Original IBM VGA

The major components of a VGA-compatible graphics adapter are shown in Figure 2. These components being: a VGA IC¹, some memory for a framebuffer, a ROM containing the VGA BIOS, connector for communicating with the host, connector(s) for attaching a display device, and some power supply circuitry.

¹The earlier VGA adapters typically used separate ICs for the RAM DAC and the VGA, but with this adapter, the functionality is combined into the main IC. Additionally, this adapter contains hardware support for 3D graphics acceleration.

2.2.1 The Video Graphics Array

Before VGA, the graphics adapter's functional components were contained within multiple ICs. VGA was IBM's first adapter which contained the majority of the logic functionality within a single IC. The functions that were combined into one ASIC were:

- CRT controller
- Attribute controller
- Sequencer
- Pixel data serialiser
- Graphics controller
- Miscellaneous control registers
- ASCII text to pixel converter
- Bus interface logic
- Memory controller

The CRT Controller (CRTC) generates the timing signals necessary to display an image on a VGA or DVI connected monitor. Images are sent to a monitor as a stream of pixels, scanning from left to right, and from top to bottom. The CRTC generates the signals which synchronise the pixel-stream being sent. These are the horizontal synchronisation (hsync) signal that indicates the end of the current horizontal line, and the vertical synchronisation signal that marks the end of the current screen (see Figure 2.2 for more detail).

Early VGA adapters, like their predecessors, had a very limited quantity of framebuffer memory to store the image to be displayed, due to the price of DRAM ICs. Rather than store every colour as RGB (Red Green Blue) components, the colours were stored as indices into a colour palette. The palette contained the corresponding colour for a given index, and the palette only contained indices for a subset of available colours. This technique allowed the use of smaller framebuffers.

In monochrome modes, for example, one bit encodes the colour, so it can only be one of two values. VGA supported 1-bit, 2-bit, 4-bit, and 8-bit colour modes. The attribute controller was responsible for converting these colour indices into the full 24-bit RGB colour format that the video DAC then converts to VGA compatible analogue signals.

The attribute controller contains two colour look-up tables, or palettes, and these are typically daisy-chained, the output of the first becoming the input to the second (except in 8-bit, packed-pixel modes, where the first table was not used). The first palette takes as input a 4-bit value and generates a 6-bit value. This 6-bit value is combined with some additional register values, and is

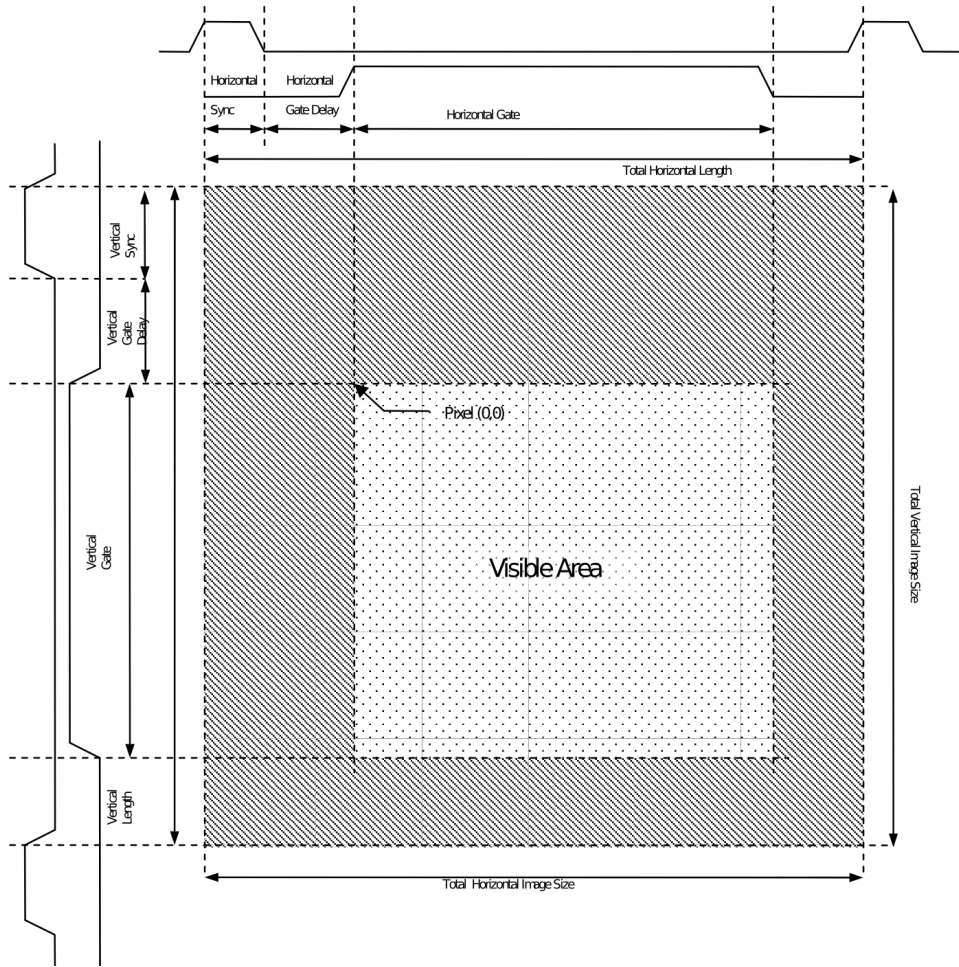


Figure 2.2: VGA display redraw timing and signals. (Image from [26].)

used to generate the an 8-bit index into the second table. This 8-bit index into the second palette retrieves an 18-bit output RGB value, 6-bits per colour component, and is then combined with additional attribute controller registers to extend it to the final 24-bit RGB format.

The attribute controller has other programmer accessible features too, it can also mask pixels from any of the four bit-planes, depending on the video mode, can enable pixel blinking, and can pan the display horizontally, which works in both graphics and alphanumeric modes.

To set a single pixel, in the 4-bit colour mode for example, a single bit from each of the four bit-planes (see the display memory explanation in Section 2.2.4) may need to be modified, therefore each of the four DRAMs needs to be accessed. To improve performance with these operations, the graphics controller allows various programmer-set, bitwise operations and shifts to be performed on the incoming data and combines these modified values with the data stored in the bit-planes. Afterwards, the modified values are written back to the DRAMs, and these operations are pipelined to achieve high performance.

All of the timing signals required for the VGA are controlled by the sequencer, and often generated from a single off-chip oscillator. The sequencer generated dot-clock and character-clock signals control the timing for nearly all of the VGA, with the exception of the external bus clock signal. One last feature of the sequencer allows the bit-plane to be enabled or disabled too.

Pixel colour-values are not stored within the display memory when the VGA is operating in alphanumeric modes. Only ASCII-encoded text along with its corresponding attributes (see Figure 2.3 to see how ASCII alphanumeric fonts were stored as bit-masks), and a font bit-mask, are stored in the display memory. The attributes values for each ASCII character are foreground and background text colour, and optionally a value that causes that particular character to blink². These alphanumeric modes were created since the memory storage requirements are far lower than graphics modes where every pixel is addressable. In the standard 80 column, 25 row, alphanumeric mode (from now on referred to as 80x25 text-mode), only requires 4 kB of DRAM. Due to the low memory requirements of alphanumeric modes, up to eight “pages” were supported, and could be dynamically exchanged by modifying user programmable registers.

VGA, and all earlier display adapters produced by IBM, includes dedicated hardware for performing ASCII-character to pixel conversion, since it needs to operate continuously while in alphanumeric modes. ASCII text is retrieved from the framebuffer at the rate of the character clock, and converted to pixel data at the rate of the dot-clock, which is 25.175 MHz in the default VGA 80x25 text-mode. The default VGA text-mode resolution is 640x400 pixels, with a redraw rate of 60 Hz, and for the system’s x86 processor (the generic name for the Intel 8086 CPU architecture and derivatives) to produce pixel values at this rate would have been difficult in 1987 .

All of the control registers of VGA were mapped to x86 I/O ports and these are listed in Table 2.1. For a device to be VGA-compatible, these control registers have to be emulated. These registers were extended by SVGA to allow more video modes, but these were not standardised between manufacturers. Subsequently, directly programming SVGA registers is complicated, but there are books which detail the programming of the more popular SVGA graphics cards [39, 14].

2.2.2 Video DAC

The video DAC converts 24-bit, RGB-encoded, colour values into the analogue signals required to drive a VGA display. The video DAC was originally a separate IC [14, 39] though modern graphics adapters can include them within the main ASIC (see Figure 2). The digital encoding consists of three 8-bit fields, representing each RGB colour component, and the DAC produces analogue values which range in magnitude between 0.7 and 1.4 V.

²Character blinking is achieved by periodically swapping foreground and background colour as the character is decoded into pixels.




Figure 2.3: ASCII and extended-ASCII character sets, and an example pixel representation.

2.2.3 VGA BIOS

Each VGA device requires a VGA Basic Input/Output System (BIOS)) ROM. This is a collection of routines for performing video functions. These routines are accessed on an x86 PC architecture by using software interrupts (using the instruction ‘int 10h’, register values are used as the arguments) [14, 39], and are typically written in just 16-bit, Intel 8086, assembly code.

The VGA BIOS contains routines for changing modes, setting pixels, cursor positioning, and many others. These routines are required for VGA compatibility. A separate ROM IC containing this code is usually present on a VGA board, as shown in Figure 2. A free and open source implementation of the VGA BIOS is available, and this is included within the OpenVGA project since it provides all the required functionality, complete with a VESA BIOS Extensions (VBE)

Port	Direction	Register Description
3B4h	R/W	CRTC Controller Address Register
3B5h	R/W	CRTC Controller Data Register
3BAh	Read	Input Status #1 Register
3BAh	Write	Feature Control Register
3C0h	R/W	Attribute Address/Data Register
3C1h	R/W	Attribute Data Read Register
3C2h	Read	Input Status #0 Register
3C2h	Write	Miscellaneous Output Register
3C4h	R/W	Sequencer Address Register
3C5h	R/W	Sequencer Data Register
3C7h	Read	DAC State Register
3C7h	Write	DAC Address Read Mode Register
3C8h	R/W	DAC Address Write Mode Register
3C9h	R/W	DAC Data Register
3CAh	Read	Feature Control Register
3CCh	Read	Miscellaneous Output Register
3CEh	R/W	Graphics Controller Address Register
3CFh	R/W	Graphics Controller Data Register
3D4h	R/W	CRTC Controller Address Register
3D5h	R/W	CRTC Controller Data Register
3DAh	Read	Input Status #1 Register
3DAh	Write	Feature Control Register

Table 2.1: VGA I/O Ports.

implementation.

2.2.4 Display Memory

The original IBM VGA had display memory that consisted of four 8-bit DRAMs, and each DRAM contained the data for one of the four bit-planes (except in packed-pixel modes, where the memory was not arranged as planes). The total quantity of display memory on original VGA adapters was typically 256 kB, though versions were available with only 64 kB or 128 kB of DRAM and could not support all video modes. The portion of display memory which is involved with redrawing the screen in the current mode is also called the framebuffer, as it stores (buffers) the data for the current frame.

A VGA display adapter has graphics modes, which are also called “all points addressable” modes, in which every pixel is mapped to a location in memory (or multiple “bit-planes” in some VGA modes). For the 640x480 resolution graphics mode, the maximum resolution officially supported by the original IBM VGA, there are 307,200 pixels, each can be set individually. Some graphics modes support multiple pages of image data. These can be exchanged, typically during the vertical-synchronisation period, to achieve a smooth animation effect.

2.3 Non-VGA Graphics Adapters

Graphics adapters have been developed and produced that are designed to be installed and used alongside a VGA adapter within a PC. Prominent examples are a couple of the 3D graphics accelerators produced in the 1990s, the PowerVR Series 1, and 3DFX Voodoo Graphics (later referred to as Voodoo 1) and Voodoo 2. The primary graphics adapter, typically a SVGA, is used for all 2D modes, like those commonly used by OS Graphical User Interfaces (GUIs). Once a 3D application is started, like a computer game, the 3D graphics accelerator is activated and is used to redraw the display instead.

The PowerVR and Voodoo 1 & 2 adapters used a pass-through architecture, where the primary graphics adapter had its output redirected to the 3D accelerator, in a daisy-chain configuration. This was achieved with a short “loopback” cable routed from the primary adapter to the input port of the secondary adapter. The output of the secondary adapter was connected to the display device, typically a SVGA CRT monitor.

2.4 Related Projects

There exists several other graphics adapter projects, and a brief overview of these is given. The most extensive of these is the Open Graphics Project, with aims to produce both FPGA and ASIC graphics adapters, and including support for multiple operating systems. Also planned is 3D acceleration, including support for OpenGL [25], an open-specification 3D graphics library.

2.4.1 Open Graphics Project

Open Graphics Project³ is a prominent open-source VGA project, which aims to develop a VGA compatible display adapter with hardware 3D acceleration support, and OpenGL drivers. Development is currently using an FPGA-based board, called OGD1 (See Figure 2.4), but the project’s stated goal is to raise enough money to have an ASIC fabricated.

³For more information see <http://opengraphics.org/>

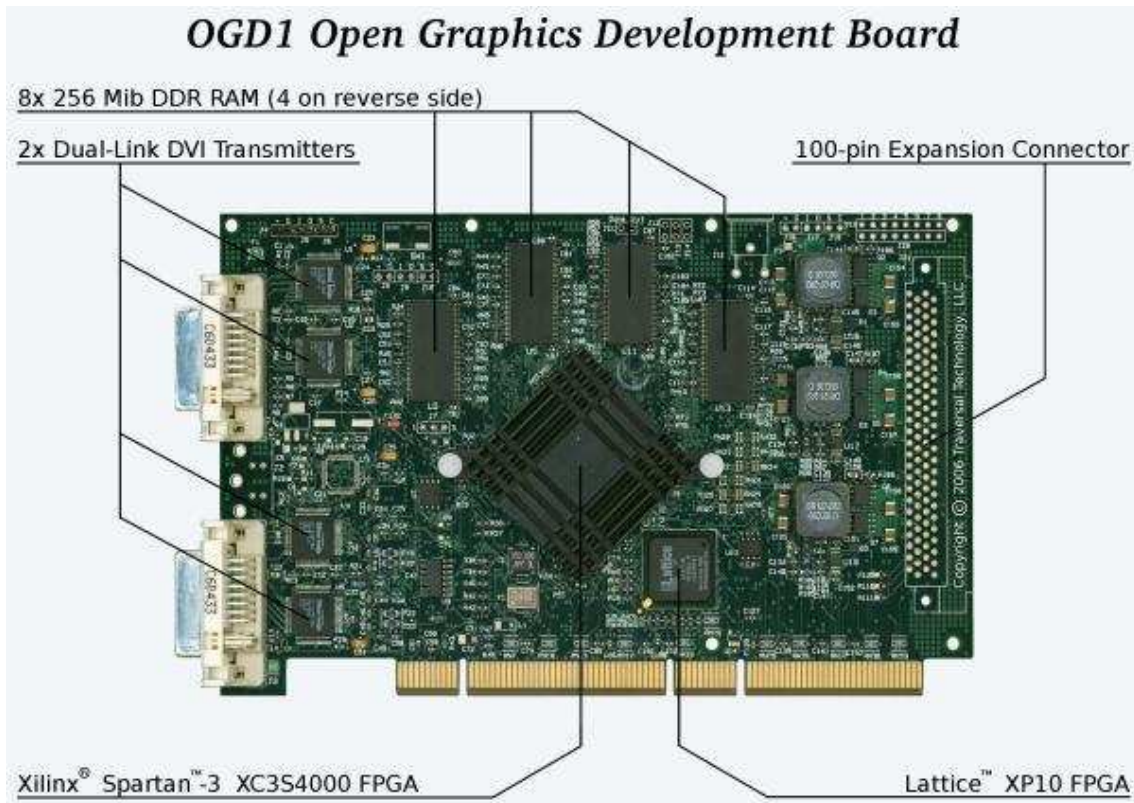


Figure 2.4: OpenGraphics OGD1 PCB and components. (Image from [30].)

OpenGraphics has an impressive feature list, allowing for a very capable graphics adapter, and is significantly more complex than what was planned for OpenVGA, but the cost of the development boards is around 1000 USD. Their plan is to emulate VGA using a soft processor core. The OGP processor core is a 32-bit RISC design, and the architecture is similar to a MIPS processor [30].

The OpenGraphics source code is freely available and licensed under the GPL, but all contributions back to the project require copyright be assigned to Traversal Technologies Incorporated so they can dual-license the code to sell commercial versions of it.

At the time of writing this, the development of OGP has reached the point where they have OGD1 running as a simple non-VGA framebuffer device, and now can run code on their soft-processor core too.

2.4.2 Project VGA

This project has stated aims similar to OpenVGA, it aims to be low-cost, simple, VGA compatible, open-source graphics adapter (see <http://wacco.mveas.com/>). Unfortunately though, the status of this project does not appear to have been updated since 28/02/08 .

2.4.3 Manticore

The goal of this project was to develop a non-VGA, 3D graphics accelerator, and once this stage was complete, add 2D support to it. The project uses an older model Altera FPGA and does not seem to have been updated in the last three years.

Chapter 3

OpenVGA Outline

This chapter presents an introduction to the architecture of the OpenVGA graphics adapter that was developed for this thesis. OpenVGA has 8 MB of local memory which stores pixel colour information, firmware code, and state information. This memory is accessible to the host computer so it can modify the contents of the framebuffer and firmware. OpenVGA continuously encodes and transmits frames of pixel data to the computer monitor that is connected to it, typically at 60 frames per second.

An outline of the OpenVGA hardware, HDL logic-cores, firmware, and software drivers are covered here. Figure 3.1 introduces OpenVGA hardware and Figure 3.4 is an architecture-overview block diagram that contains the components which will be discussed within this chapter. Many of these components will be further elaborated on in the following chapters.

3.1 Hardware Development

Most of the significant hardware components of OpenVGA are shown in Figure 3.1. Components not shown in this image are the VGA video DAC and the DVI TMDS encoder, as these components are soldered to the reverse side of the PCB (see Figure B.2). The PCB only has two copper layers to meet two of the design goals of OpenVGA, simplicity and low-cost. A full list of all of OpenVGA's electronic components and where to obtain the PCB artwork is found in Appendix B.

Logic cores were developed which interface to many of these hardware components. There are logic cores implementing functionality for PCI, controlling the Video RAM, sending and receiving via the USB UART, and encoding video data for either DVI or VGA displays. These logic cores are introduced later in this chapter in Section 3.2.

Initial hardware development concentrated on adding PCI support to an existing Spartan-3 development board, as shown in Figure 3.2. The PCI logic core needed to be developed first as this allowed the host PC to run testbenches on the logic cores which were developed later. A simple




Figure 3.1: OpenVGA Hardware Features.

PCI stub PCB had to be constructed (the copper-coloured board shown in the photograph). This board contains only two ICs, just simple bus-switches to translate between the 3.3 V signalling of the Spartan-3 development board and the 5 V PCI signalling used in many PCs.

The stub board was then connected to a commercial Spartan-3 Starter Kit development board (the green circuit board shown in the photograph) to develop the PCI-to-Wishbone-bridge logic core (see Sections 3.2.6 and 6.2). An additional advantage of this arrangement was that the connectors between the stub board and the development board allowed digital oscilloscope probes to be attached. This made debugging of the PCI logic core easier since the electrical signals of the PCI bus could be monitored during testing (see Figure 6.6).

Following this, an initial prototype board for OpenVGA was milled in-house, on a LPKFTM PCB milling machine, and is shown in Figure 3.3. This board was an attempt to use Double Data Rate (DDR) SDRAM with OpenVGA to provide greater memory bandwidth¹. Though the DDR SDRAM controller was successfully simulated with the Icarus Verilog simulator, using the Micron-

¹The DDR SDRAM specification requires that signal traces be terminated, since DDR SDRAM was designed to operate at frequencies up to 200 MHz. The OpenVGA DDR version of the PCB did not allow for these resistors since satisfactory placement would be extremely difficult with the two-layer PCB used, since there are no ground-planes. It was decided to attempt DDR support without termination resistors, since the trace-lengths were very short, but DDR SDRAM could not be made to operate.




Figure 3.2: PCI development stub board. The bus-switches, which are visible in the middle of the copper coloured PCB, translate the PCI's 5 V signals to the 3.3 V required by the Spartan-3 FPGA.

provided DDR SDRAM Verilog module, the synthesised design did not work². Due to the problems with DDR SDRAM, the final version of the OpenVGA PCB was designed to use standard, single data-rate SDRAM (labelled as “Video RAM” in Figure 3.1).

3.2 Important Design Elements within the FPGA

Implemented within OpenVGA's Spartan-3 FPGA are many logic cores and the necessary bus and clocking logic. Figure 3.4 shows the significant features of the digital-logic design, and the FPGA's connectivity with the other on-board components of OpenVGA. All logic was described using the Verilog HDL, simulated using Icarus Verilog, and synthesised using tools which Xilinx freely provides for use with their FPGAs.

The design contains a processor core for data processing, initialisation, and mode-managing

²The DDR signals showed significant overshoot and undershoot when monitored on an oscilloscope, and the DDR IC produced a lot of heat.



Figure 3.3: Prototype PCB featuring DDR SDRAM, DVI, and only eight colour VGA (via the DVI analogue outputs).

tasks. OpenVGA can be configured to use either one of the two processors that were developed for this project, TTA16 or RISC16. Additional cores are a small data cache since memory latency is high, a SDRAM controller, a PCI-to-Wishbone bridge, logic to read the SPROM (Serial Programmable ROM), and numerous other support modules.

3.2.1 Logic Core Interconnection

The Wishbone interconnect standard³ was used as the standard interface for communication between logic cores. This standard was designed for System-on-a-Chip (SoC) applications, is very flexible, is commonly used with other open-source hardware projects, and requires very little logic to implement. Appendix C gives a brief introduction to Wishbone, covering the core signals and providing a simple example implementation.

3.2.2 FPGA-Optimised Processor Cores

The OpenVGA initialisation, state-management, and data-processing tasks were to be performed using a processor logic core. Existing processor logic cores were considered but, due to performance,

³Wishbone is called an “interconnect”, instead of a “bus”, standard as it allows for numerous circuit topologies: buses, crossbar switches, point-to-point, and many others.

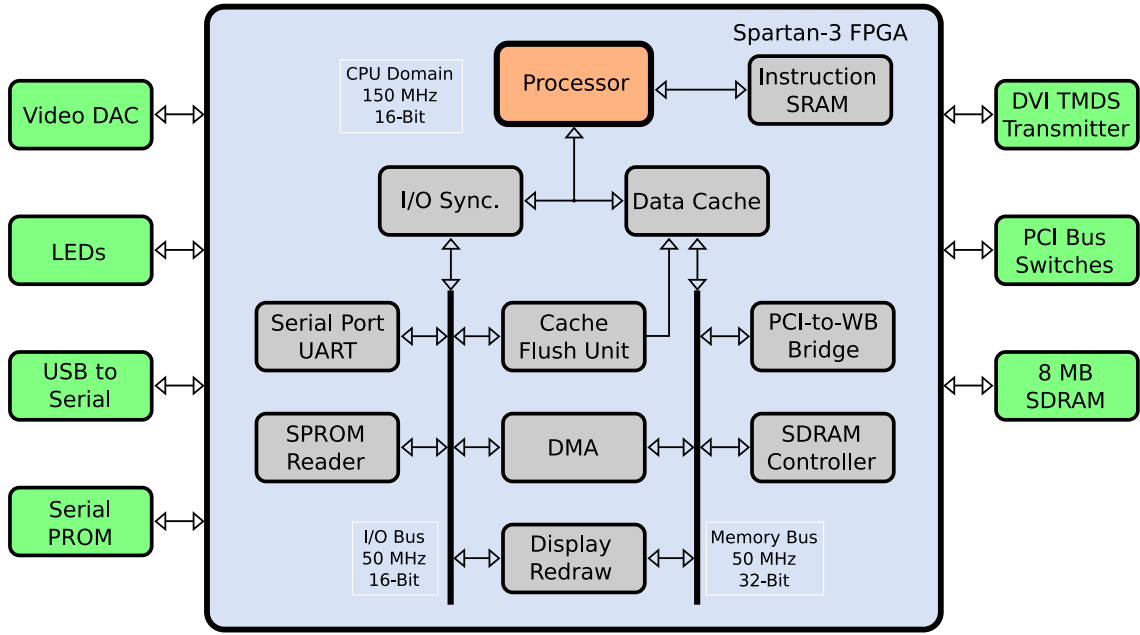


Figure 3.4: OpenVGA Architecture Block Diagram.

capability, and size requirements, two new processors were developed for OpenVGA instead. OpenVGA requires a processor that is be able to address at least 8 MB of memory, so this ruled out existing 8-bit processors. The evaluated 32-bit processors were too large, so a fast 16-bit processor was needed, but nothing suitable was found.

The first processor developed was a novel TTA processor. Significant research has been done with this class of architecture [6, 24], but there are very few commercial processors of this type. The second is a more traditional Reduced Instruction Set Computer (RISC) architecture processor design. It was developed to compare with TTA16, the first OpenVGA processor. Optimised for the same tasks, both processors used the similar functional units. This provides a good platform for the comparison of these two architectures.

TTA16

TTA16 is a 16-bit, Wishbone-compatible, TTA processor that can operate at up to 190 MHz, and has a very small footprint⁴. It has a fixed, 32-bit instruction word and can perform four data transports every clock cycle. This processor is covered in detail in Section 4.5, and Appendix D is a programming guide for TTA16.

⁴These results are for synthesis on a Spartan-3 FPGA, and depend on the build configuration and synthesiser optimisation settings.

RISC16

Due to the difficulty of writing assembly code for TTA16, RISC16 was developed and is of a more traditional design. While not as fast as TTA16, operating at up to 140 MHz, instruction width is only 16-bits so code density is better. The external interface is Wishbone-compatible too, and is therefore identical to TTA16, so that this logic core can be a drop-in replacement. More detail of RISC16 is within Section 4.6, and Appendix E is a simple programming guide.

3.2.3 Processor Data Cache

A low-latency, Wishbone-compatible, data cache was needed for adequate processor memory-access performance. The 8 MB of on board memory is shared amongst multiple logic cores, and the latency is several Wishbone cycles, so any processor access would lead to stalls of many cycles. Both interfaces to this cache are Wishbone, and the number of width memory-address bus is parameterisable.

Other cache features are a fast-hit path with a latency of zero cycles, a slower hit path of one cycle of latency, operating frequency of 150 MHz, two-way set-associativity, a line-size of 64 bytes, and a capacity of 2 kB. A full explanation of the cache's features and design decisions is within Section 5.2).

3.2.4 Clock Domains and Domain Crossing

OpenVGA has multiple external interfaces and several different and asynchronous clocks. OpenVGA has three clock domains, these are the 50 MHz Wishbone bus domain, the 33 MHz PCI Local Bus domain, and the dot-clock domain. While the PCI bus typically operates at a frequency of 33 MHz, though this is not guaranteed, and this clock signal is driven by the system host. The dot-clock of display circuitry depends upon the video mode, as higher resolutions require more pixels be drawn, therefore requiring a higher pixel rate. Currently the dot-clock can be selected as either 25 MHz or 40 MHz (see Section 6.3.1).

As well as asynchronous clock domains, some components of OpenVGA run at integer multiples of the Wishbone domain's 50 MHz clock rate, the SDRAM and RISC16 operate at 100 MHz, and TTA16 operates at 150 MHz. These clocks are synchronous with the Wishbone clock and present fewer problems. A further explanation and solutions are covered in Section 6.1.2.

When logic cores have different clocks, data must be synchronised when passing from one core to the other. When the clocks are asynchronous this becomes even more difficult due to metastability issues. A D-type Flip-Flop (DFF) can enter a metastable state⁵ when setup and/or hold times are violated. Section 6.1 presents the OpenVGA clock architecture and domains, examines the

⁵This metastable state means that the output of the DFF is essentially random, but it oscillates for a significant time before it settles back into a stable state [5].

problems with domain crossing, and presents the solutions used to solve these. Asynchronous First-In, First-Out queues FIFOs allow signals that are multiple bits wide to be synchronised across clock domains and are presented in Section 6.1.3.

3.2.5 Memory Controller

A high performance, small logic foot-print, Wishbone-compatible memory controller was needed for OpenVGA. Tasks the controller has to perform include refreshing the DRAM, support burst reads and writes, and atomic reads and writes. To minimise the size of the controller it has to operate within the Wishbone clock domain so that asynchronous FIFOs are not needed.

The SDRAM data transfers occur at 100 MHz while the controller’s state machine operates at 50 MHz, the Wishbone memory bus frequency. This was achieved by using the DDR input and output primitives within the Spartan-3 I/O Blocks (IOBs). Section 5.1 provides more detail on the controller, including the design of the state machine and data path implementations.

3.2.6 PCI-to-Wishbone Bridge

A small PCI-to-Wishbone-bridge logic core was developed so that the system host can communicate with OpenVGA. This bridge supports the “Plug and Play” standard and Memory-mapped I/O (MMIO), which maps the OpenVGA SDRAM into the host PC’s memory address space. Only the PCI Local Bus features that were needed for OpenVGA are implemented and resulted in a design about one-tenth the size of another available open-source PCI bridge⁶. The design of the bridge, descriptions of its state-machines, and domain crossing issues are explained Section 6.2, and more details of the clock domain issues are covered in Section 6.1.

So that the host OS can access OpenVGA, a simple Linux kernel module was written. Linux kernel modules are written using the C programming language⁷ and OpenVGA is accessed from the GNU/Linux OS by opening the device file which represents the MMIO of OpenVGA. Using read, write, and seek operations, any location within the OpenVGA SDRAM can be read and written.

3.2.7 VGA/LCD Display Controller

This logic core prefetches image data from the memory controller, using the system’s memory Wishbone bus, and generates the timing signals and bit-streams necessary for driving VGA and LCD displays (see Section 6.3). The Wishbone bus can be in a different clock domain to the dot-clock since an asynchronous FIFO is used for the prefetch logic. This has a 2 kB prefetch capability and is discussed in more detail in Section 6.3.2.

⁶There is a very capable, FPGA-tested, PCI bridge available over the Internet from OpenCores site.

⁷A comprehensive guide to developing Linux kernel modules is available from [34].

Within the display controller is a CRT controller which has registers that can be accessed using a Wishbone interface. This allows the timing parameters to be set allowing the display mode to be changed (see Section 6.3.3).

3.2.8 Additional Logic Cores

Other smaller logic cores were developed for OpenVGA. Detailed explanations of these can be found within Chapters 5 (Memory) and 6 (I/O), and a summary of these is:

- USB UART: A Wishbone-compatible interface to a 9600 baud UART which communicates with the on-board FTDI FT232R USB UART IC. More detail can be found in Section 6.4.2.
- LED driver: There are two Wishbone-mapped LEDs that can be set and cleared using this module, covered further in Section 6.4.1.
- DMA controller: A simple Direct Memory Access controller is used to improve the memory write performance of OpenVGA's processor. This logic core is detailed in 5.3 and a programming guide for it is in Appendix D.3.3.
- Serial PROM reader: Spartan-3 FPGAs require configuration upon power-on. The on-board serial PROM contains this configuration data but also has enough capacity to store some extra data. This module can read this data from the SPROM when requested via its Wishbone interface. Section 5.4 presents more information on this module.

Chapter 4

The TTA16 and RISC16 Processor Cores

OpenVGA is designed to use a processor logic core for initialisation, mode setting, and data processing tasks. Two processors were developed for these tasks, the first was TTA16, a novel, 16-bit, transport-triggered processor. The second processor that was developed, RISC16, a 16-bit, RISC processor architecture, was for comparison with TTA16, and for its relative ease of programming in assembly language.

4.1 Processor Architectures

An aim for OpenVGA's processor was to be fast enough to emulate VGA text-mode. After evaluating existing FPGA-based, processor logic cores (see Table 7.2), it was decided to develop a new processor to meet this speed and size criteria. Popular Instruction Set Architectures (ISAs) were considered for the OpenVGA processor. These included Complex Instruction Set Computer (CISC), RISC, and Very Long Instruction Word (VLIW) architectures.

Figure 4.1 gives a breakdown of the tasks that a compiler and processor must perform, and the division of these tasks for differing processor architectures. Modern superscalar CISC and RISC processors use a lot of die area for functionality other than transporting and processing data. This often includes hardware for instruction decoding, out-of-order execution, speculative execution, and branch prediction [31]. TTA processors typically do not have these features [7], with most of the logic gates used for data transports and Functional Units (FUs), resulting in smaller processor designs.

An efficient architecture would operate at a high frequency, have a high FU utilisation rate, and use minimal logic resources. FPGAs have reduced frequency and fewer logic gates than similar

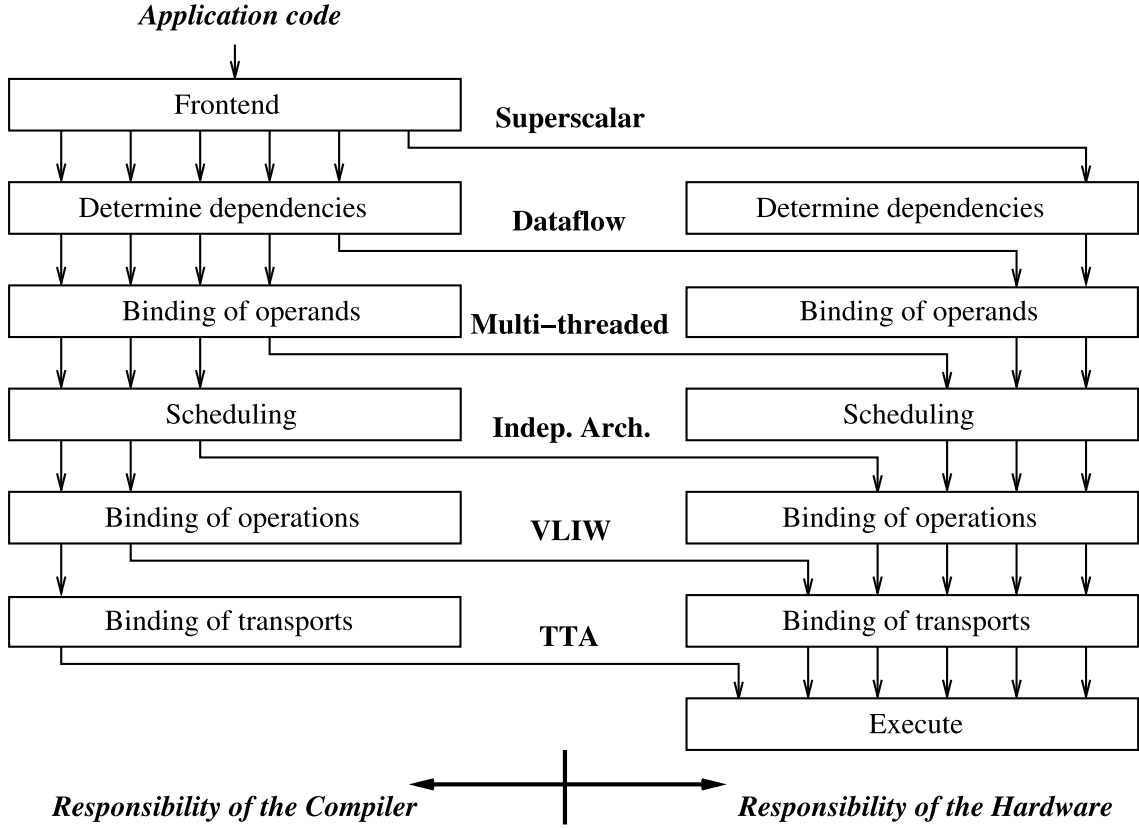


Figure 4.1: The division of responsibilities between hardware and compiler. (Image from [8])

generation ASICs so a smaller, faster processor architecture is arguably even more important. This is why a TTA processor was chosen for OpenVGA. But the complexity of generating efficient firmware code for TTA16 is why a second processor, RISC16, was developed.

TTA processors were investigated (designed, built, programmed, and characterised) at the Technical University of Delft by Henk Corporaal et. al. [8, 7, 6] during the 1990s. Additional TTA research has focused on the automatic generation of the HDL source-code for TTA processor logic cores [22, 24]. The goal has been to develop application-specific processors for certain tasks, like image processing applications [8].

Xilinx Spartan-3 FPGA logic resources, and their latencies, also affect the choice of processor architecture, and which features can be implemented efficiently as well. The Spartan-3 has 18-bit embedded multipliers, fast carry-chain logic, 16-entry and 2 kB RAM primitives, DCMs (Digital Clock Modules), and some additional multiplexing primitives¹. These available FPGA resources influenced the architectures evaluated, and the design choices made.

Both of the processor logic cores developed for OpenVGA, TTA16 and RISC16, feature a data

¹The high combinatorial logic delays of multiplexers synthesised for Spartan-3 FPGAs influenced many design decisions. Even with some hardware support for wide multiplexers these were still quite slow [45].

word-size of 16-bits for two main reasons:

- Functional units are typically half the size of those for 32-bit processors so the resulting processor will be a lot smaller.
- Wider signal paths use more routing resources, and are more difficult for the Xilinx PAR tool to route, usually resulting in lower frequencies.

4.2 Shift Registers as Program Counters

Rather than use a standard, base-2 incrementer for each processor's Program Counter (PC), a Multiple Feedback Shift Register (MFSR) [42] is used as the PC incrementer for both TTA16 and RISC16. The count order of an MFSR is very different to standard base-2 incrementers. To illustrate this different increment sequence, an application of MFSRs is for use as pseudo-random sequence generators. For a maximal-cycle MFSR having n -bits, the cycle length is $2^n - 1$, and can be implemented with a maximum of just one layer of logic gates (see Figure 4.2).

The reason for using MFSRs as PCs is due to the extremely small combinatorial logic delay, just the propagation delay through one FPGA Look-Up Table (LUT) and some routing fabric. This allows the PC increment and branch logic to have very low latency, avoiding the requirement pipeline registers, or becoming the processor's critical path and limiting maximum frequency.




Figure 4.2: An 8-bit MFSR (Multiple Feedback Shift Register) with a cycle size of 255. The total logic required is only eight D-type flip-flops and two XOR logic gates. *Image courtesy of R. Ward [42]*

4.3 Processor Performance

To improve the performance of a processor, the biggest gains are made by making the common operations fast [19]. The primary computational task of OpenVGA's processor is to convert data sent to OpenVGA into the framebuffer's data format. This consists of mainly arithmetic, bit-wise logic, and memory operations. Arithmetic and logical operations take just one clock cycle. Memory operations are far slower but a fast data-cache has been developed (see Section 5.2), reducing the average delay to only about five processor clock cycles.

Amdahl's law defines the speedup obtained from a given enhancement made to a machine [19]. This speedup is based on the performance for the entire task, with and without the enhancement, τ_E and τ_O respectively. For a given enhancement to a CPU, the speedup is the ratio defined as

$$\text{Speedup} = \frac{\tau_O}{\tau_E}$$

Amdahl's Law allows the effect of enhancements to the CPU to be objectively evaluated. For example, consider the question: "How much of a performance effect does doubling the number of processor data-transporters and FUs give?"

Answer: A memory operation for the text-mode conversion algorithm typically takes about 5 clock cycles, due to the benefit of data-caching. If about one third of the instructions contain memory operations (see Section 5.2.1), and all other operations complete in one clock cycle, and about two million memory accesses are required to redraw one frame, then the time in cycles to redraw one frame is

$$\begin{aligned}\tau_O &= 5 \times 2 \times 10^6 + 4 \times 10^6 \\ &= 14 \times 10^6\end{aligned}$$

If the number of transports and functional units are doubled, giving a perfect speedup of two for non-memory instructions, then the total speedup would be

$$\begin{aligned}\text{Speedup} &= \frac{\tau_O}{\tau_E} \\ &= \frac{14 \times 10^6}{12 \times 10^6} \\ &\approx 1.17\end{aligned}$$

The resulting CPU would nearly be twice as large though. The actual speed-up factor will be less than two because of scheduling difficulties. And the Xilinx PAR tool would find the design harder to route, likely reducing processor operating frequency too. So doubling the number of transports would not be an efficient modification in this case since the FPGA logic resources are limited and the speedup minimal.

4.4 Tools & Testing

One of the processors designed and implemented within OpenVGA, TTA16, has an unusual architecture, and it would be a non-trivial task to adapt an existing open source assembler, and

especially a compiler, to this architecture. It would have been extremely tedious to generate large quantities of object code too. Therefore an assembler was needed for both TTA16 and RISC16.

Testing the two processors was initially done with the Icarus Verilog simulator, and later with hardware testbenches. Simple TTA16 and RISC16 programs were written for processor testing to verify correct operation of the FUs. The CPU of the host PC can access all of OpenVGA's memory using a Linux kernel module.

4.4.1 Object Code Generation

Processors execute instructions in the form of object code, which is a binary representation of an instruction. Typically, assemblers, interpreters, and compilers are used to generate object code. There are currently no interpreters or compilers that have TTA16 or RISC16 as an output target since these processors have custom instruction sets. There is an assembler for each of these processors though.

TTA16 Assembler

File:	assemble	
Description:	A general-purpose TTA assembler which loads an XML processor description from a file prior to assembly.	
Related Files:	/data/tta16.xml	
Testing Files:	/src/fw_tta16/*.S	
Author:	Roy Ward, Patrick Suggate	License: GPL

An assembler was created by R. Ward [41] for another project, and additional features were added by him to assist this project. This assembler reads in a processor definition file, which is stored as an XML file, and then assembles the assembly language source file in accordance to the definition in the XML file. This assembler was designed for TTA applications and it supports explicit Instruction Level Parallelism (ILP) (see the TTA16 programming guide in Appendix D.4 for an example of this).

Below is an excerpt from `tta16.xml`, the file defining the TTA16 instruction set. This fragment encodes the available destination registers for transport-0 of TTA16², and has been included to demonstrate the syntax. The XML encoding is very general purpose, and has been used for other TTA processors as well [38].

²This bit-field is called *DST0* in the instruction format diagram, see Figure 4.8

```

<option>
  <token value="\"/>
  <part start="30" width="2">
    <value name="RREG" value="2" ext_field="rs0"/>
  </part>
  <token value="->"/>
  <part start="27" width="3">
    <value name="nil" value="0"/>
    <value name="bra" value="1"/>
    <value name="rad" value="2"/>
    <value name="wad" value="3"/>
    <value name="jb" value="4"/>
    <value name="jnb" value="5"/>
    <value name="jz" value="6"/>
    <value name="jnz" value="7"/>
  </part>
</option>

```

The first line in a TTA16 assembly source file is typically:

```
architecture:tta16
```

This line specifies the name of XML file to use when assembling this file and this line has to precede any assembly language. Comments may precede this line, and since the `m4` macro processor was used as a preprocessor with this assembler, to allow the use of constants and macros for common tasks, macro definitions will often be at the start of TTA16 assembly files too.

A TTA16 instruction encodes the move operations for each of its four transports. Each of these four `move` fields are separated by commas, and an empty field indicates no-operation for that transport within the current instruction.

Arrows (a minus and a greater-than, `->`) are a required part of the assembly syntax and they graphically indicate the direction of data flow. Curly braces surround all instructions, and instructions can be preceded by an optional label, before the braces. More information on TTA16 assembly and programming is in Appendix D, but here is an example of the TTA16 assembly syntax:

```
loop: {r1 ->rad, 2 ->sub, ,r2 }
```

Shown above is a label, `loop:`, two general-purpose registers (`r1` and `r2`), and two special-purpose registers (`rad` and `sub`). The third `move` field is empty, indicating no-operation, and the fourth transport has only one argument, `r2`. This is the source for the data-move, the destination is always the common ALU operand register `com`, and is therefore excluded from the syntax.

RISC16 Assembler

File:	r16asm.py	
Description:	A simple RISC16 assembler.	
Related Files:	/src/r16asm.py, /src/CodeCleaner.py, /src/AsmParse.py, /src/Emit.py	
Testing Files:	/src/fw_risc16/*.S	
Author:	Patrick Suggate	License: GPL

RISC16 uses a custom assembler, written in Python specifically for this processor. Example code written for this assembler is shown in Appendix E.4. RISC architectures are not explicit-ILP processors, and the syntax used with the TTA assembler is somewhat cumbersome for RISC16. Additionally, RISC16 had some instructions that were hard to define within the XML processor definition file (especially the immediate constant field of the `i12` instruction prefix). This justified the creation of yet another assembler.

The general form of a valid line of assembly code that is accepted by this assembler is:

`[label:] mnemonic [arguments]`

The `label:` field is used to label the destination of branches, and to label constants so they can be referenced within the code. The `mnemonic` field is a keyword identifying the desired instruction, which is translated into an opcode (condensed form of operation-code) in the final phase of the assembler. Fields surrounded by square braces indicate that the field can be optional. All constants are required to have a label, and most instruction require arguments.

An overview of the algorithm the assembler uses to generate an output file is:

1. Parse command-line options for the input file name, the optional output file name, and the assembler options.
2. Read in assembly file, line-by-line. Line numbers are stored, with each line, at this point. This is needed so useful error messages can be generated later.
3. Remove comments and extra white-spaces, throwing away lines containing no code or no assembler directives.
4. Tokenise each of the remaining lines according to the simple format discussed above.
5. The parser takes the stream of tokens, builds a table of constants, and then a table of branch destination labels. The next parse step is evaluating the instruction arguments, which involves substituting any labels, in the arguments, for its corresponding constant or address value. The final parse step is to evaluate any argument that represents a numerical value, the set arithmetic and bit-wise logical operations used in the C language is accepted.

6. Since both TTA16 and RISC16 use an MFSR instead of a standard PC incrementer, the generated object code needs to be reordered according to the MFSR count order. Since the count order is pseudo-random, it will span the entire instruction RAM block(s), so before reordering the instruction stream, it is padded with `nop` (no-operation) instructions so that it contains the same number of instructions as the block RAMs can hold.
7. The parsed and reordered instruction mnemonics and arguments are matched against valid instruction formats, and upon match, the hexadecimal representations of the object code is then emitted to the output file.

4.4.2 Processor Testing and Testbenches

Initially, each FU was developed and tested using the FOSS (Free and Open-Source Software) simulator Icarus Verilog³. Processors can be very difficult to debug, especially RISC16 with its many hazard detection flags and states, so it was important that the components were well tested first. Another FOSS Electronic Design Automation (EDA) application, GtkWave⁴ was used to display the timing diagrams of the simulated signals, which were saved to a VCD output file by Icarus.

When a processor had been developed to the point where it simulated as expected, larger quantities of test code were written, with a disassembler coded within the processors HDL source displaying the instructions (with input and output values) being executed. The MMIO functionality and correct data cache operation could then be verified.

Once a design was believed to operate correctly, hardware testbenches were synthesised and uploaded to OpenVGA, at first just running code for a trivial task like flashing some LEDs. Figure 4.3 shows the synthesiser output for one of the TTA16 hardware testbenches. As can be seen, TTA16 is very fast for a processor implemented on Spartan-3⁵.

The Verilog synthesiser used was ISE WebPackTM 9.1 developed by Xilinx, since it is free of charge, though proprietary, and available for both the Windows and GNU/Linux operating systems. Being free, it keeps with the goal of OpenVGA being a low-cost platform for development.

³Icarus Verilog is available on the Internet from <http://www.icarus.com/eda/verilog/>

⁴GtkWave is available on the Internet from <http://gtkwave.sourceforge.net/>

⁵It is often estimated that an FPGA has about one-tenth the performance of an ASIC made with the same manufacturing process, in this case 90 nm.

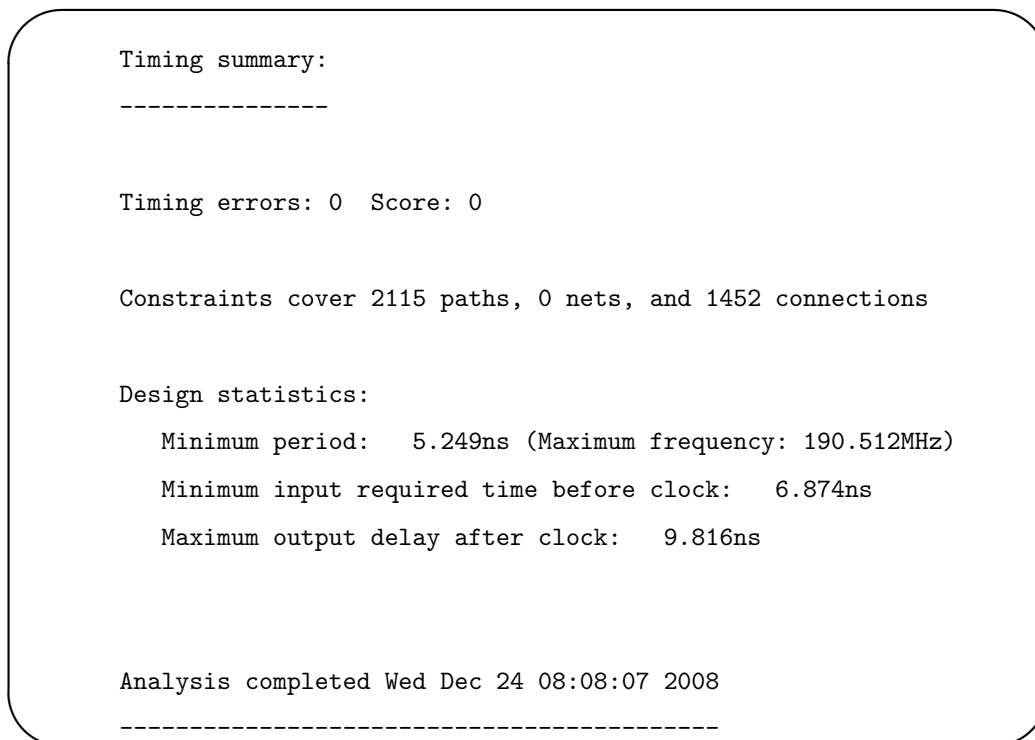


Figure 4.3: Xilinx Place and Route (PAR) report showing the synthesised TTA16 hardware test-bench timing and routing information.

4.5 TTA16

Module:	tta16	
Description:	A 16-bit, TTA processor logic core designed for use with FPGAs.	
Related Files:	/rtl/cpu/tta16/tta16.v, /rtl/cpu/tta16/tta16.xml, /rtl/cpu/tta16/tta_stream4to4.v, /rtl/cpu/tta16/tta_stream4to8.v, /rtl/cpu/tta16/tta_stream8to8.v, /rtl/cpu/fastbits.v	
Testing Files:	/sim/cpu/tta16_tb.v	
Author:	Patrick Suggate	License: GPL

TTA16, is a small, fast, 16-bit, transport-triggered, processor logic core designed for OpenVGA initialisation, state management, and the data processing tasks for emulating VGA functionality. TTA16 features a three-stage pipeline, four data transports, explicit ILP, 16 general purpose registers, up to 32-bit memory address support, logic and arithmetic operations, a multiplier, a fast local instruction store, and a fast Wishbone bus interface to access OpenVGA Wishbone bus peripherals, including the system memory.

TTA16 is a partially-connected TTA, each FU can only get data from, and place data on, a subset of the processor transports. The FU and data-transport connectivity of TTA16 is shown in

Figure 4.5. Partially connected TTAs allow the use of smaller multiplexers for connecting FUs to transports, resulting in smaller and faster designs [2].

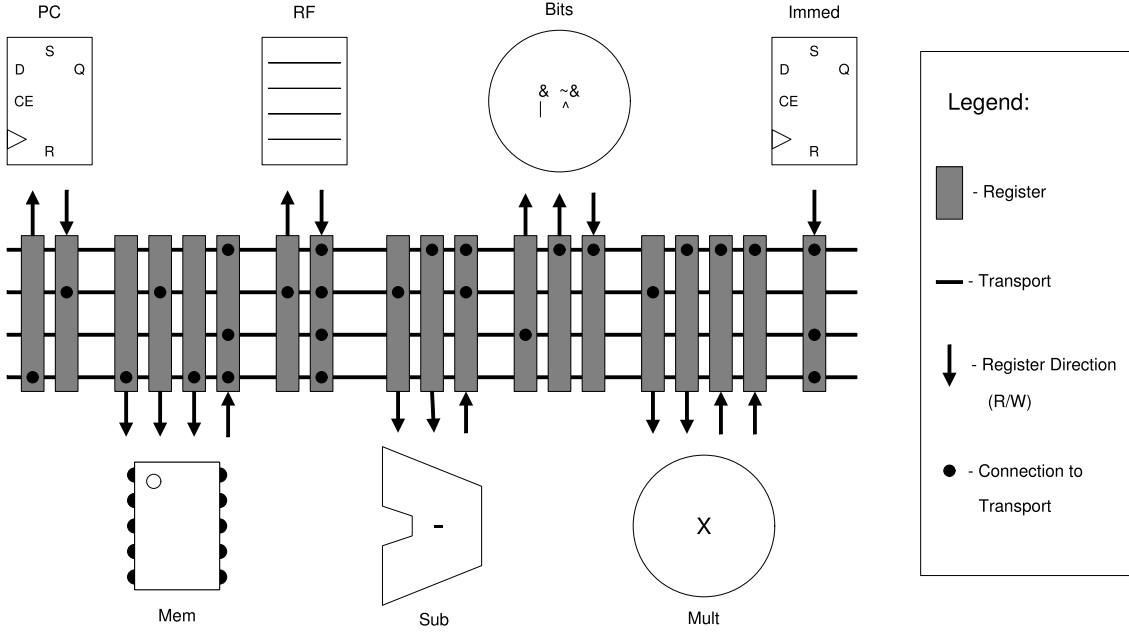


Figure 4.4: TTA16 transport-based diagram which shows the connections between transports and FUs.

Even though TTA16 is a very capable processor, when synthesised for the Spartan-3 architecture TTA16 uses only approximately 220 logic slices⁶ which is very small for the feature set. Additionally, the maximum clock rate at which it will run is up to 190 MHz (see Figure 4.3), very fast for a processor implemented within the Spartan-3 FPGA family. A comparison with other Spartan-3 based processors is shown in Table 7.2.

4.5.1 Introduction to TTA Processors

A TTA processor is different from more common processors, like RISC and CISC based processors such as ARM, SPARC, x86, Power, and MIPS. TTA instructions are simply a list of data moves for each transport that the processor has to perform. For comparison, instructions of more traditional architectures typically specify the desired operation and the data to use for that operation.

To perform useful work, processors need FUs (Functional Units) that carry out some operation, like branching, memory load and store, and subtraction and other calculations. Work is performed within a TTA processor by transporting values to special purpose registers, some of which can trigger side-effects, like computations. There are three main classes of registers

⁶This is when XST is set to optimise for speed, when optimising for size it uses even less.

- **Trigger Registers:** These trigger the start of operations within FUs. The operations which are initiated depend on the FU, they can be an operation like a memory write, or a subtract, or whatever computations the FU supports.
- **Operand Registers:** These supply additional values for an operation, a subtract requires both a subtrahend and a minuend, for example. Unlike trigger registers, writing to an operand register has no side-effects. The number of operand registers required by a FU varies, and can even be zero.
- **Result Registers:** These are the results of computations by FUs, they can have a latency of zero or more cycles, and FUs can have more than one result register, like the multiplier (see Section 4.5.4) for example.

Figure 4.1 demonstrates how to perform a calculation in a TTA processor, in this case an addition. RISC-type assembly code is given for comparison. The simple addition example looks to be a clear win to the RISC processor, even if it is assumed that a TTA can be made to operate at a 50% higher clock rate (see Table 7.2).

Also note that the RISC instruction reads two values from the Register File (RF, `r0` and `r1`), and writes back one (`r0`) This requires the RF to have a bandwidth of three operations per clock cycle, or else it will take longer than a single cycle to complete, or maybe affect/restrict other instructions.

RISC-type instructions	TTA-type instructions
<code>add r0, r0, r1</code>	<code>{r0 -> add }</code>
	<code>{r1 -> add_t }</code>
	<code>{sum -> r0 }</code>

Table 4.1: A simple addition operation to demonstrate the TTA concept. The addition is only initiated once a write is made to the addition-trigger register (`addt`).

For the second example, lets say the task is to sum four numbers, stored in `r0-3`, and the example TTA now has an ILP of two⁷. As Figure 4.2 shows, the new TTA code now fares a little better, it can even be considered to win if it can operate at a 50% percent higher clock rate. The TTA still requires less register-file bandwidth as well, five accesses compared to nine of the RISC.

Due to TTA16's explicit ILP, four data moves occur per instruction. TTA16 is not fully connected though, each execution stream cannot access all of the processors registers, just a smaller

⁷An ILP of four [6] is common but much higher than this was experimented with by Henk Corporaal[6]

RISC-type instructions	TTA-type instructions
add r0, r0, r1	{r0 -> add, r1 -> add _t }
add r0, r0, r2	{r2 -> add, sum -> add _t }
add r0, r0, r3	{r3 -> add, sum -> add _t }
	{ sum -> r0 }

Table 4.2: TTA Accumulate Example. Accumulate four numbers and store the result in `r0`. This TTA has an ILP of two, meaning two data moves are performed every clock cycle.

subset. This approach makes the transports simpler, and therefore faster, and if the connections are chosen well, the net effect is a performance increase [2] due to allowing a higher clock rate.

Even with four moves per clock, TTA16 still required less RF bandwidth than RISC16 (a general feature of TTAs [21]), so the RF is only half the size of RISC16 (see Section 4.6.6), 16 Spartan-3 slices (32 LUTs) versus 32 of RISC16. Within OpenVGA, the total RF bandwidth of TTA16 is identical to RISC16 since a two port RF operating at 150 MHz has the same bandwidth as three ports at the 100 MHz of RISC16⁸.

4.5.2 Internal Tri-States vs. Multiplexers

Multiple FUs are connected to the four transports of TTA16, but since only one FU result register can place its value onto a transport during any given clock cycle, there needs to be hardware that selects the register to be placed on the transport. Typically, multiplexers and internal tri-states are used, but the Spartan-3 family does not include internal tri-states, unlike previous FPGA families from Xilinx [45]⁹.

The Spartan-3 does support internal pull-up networks though [45]. Internal pull-ups are implemented using the general-purpose, four-input LUTs, but using them as internal pull-ups often requires adding one extra layer of LUTs. This extra layer of LUTs, and the routing cost, lead to a large performance penalty as shown in Table 4.3. Multiplexers are a clear win in this situation, so they were used for TTA16.

⁸The operating frequencies of TTA16 and RISC16 are lower than their absolute maximum values due to routing factors, the data cache frequency, and the Wishbone bus clocks.

⁹Internal tri-states were implemented within the Virtex II Family via primitives called TBUFs. This was often of limited value since the number was relatively scarce and the locations fixed. The Spartan-3 does not have any TBUFs.

Routing Method	Speed (MHz)	Size (Slices)
Internal Pull-ups	147	245
Multiplexers	190	203

Table 4.3: Synthesis results comparing multiplexers and internal pull-ups.

4.5.3 The TTA16 Pipeline

TTA16 is pipelined so that the work of executing a single instruction is spread over multiple clock cycles. Instructions are first fetched, data are then moved via transports, and then the specified functional units are triggered.




Figure 4.5: TTA16 pipeline overview showing a three-stage pipeline.

A new instruction is fetched every clock cycle, except during a memory read/write interlock. There are up to three instructions “in-flight” within the pipeline (see Figure 4.5.3). The purpose of pipelining is to allow more instructions to be issued per unit time, but each instruction will probably take longer to execute than a the same instruction in non-pipelined, but otherwise similar, processor¹⁰. It is a form of ILP since there are three instructions being processed simultaneously within the TTA16 pipeline, though each in a different stage.

4.5.4 Functional Units

The basic design of a TTA processor is several transports connecting multiple FUs together. This is essentially true for all processors but the concept of using transports to move data between FUs

¹⁰This is due to two main reasons, the first being that the pipeline registers each add a small combinatorial component. The second reason being that the total work required to execute one instruction will not be uniformly divided amongst each pipeline stage, a critical path will limit the maximum frequency, other paths will be slightly quicker.

is very explicit with TTA processors.

Program Counter and Flow Control

Functional Unit Name:	Program Counter
Trigger register(s) and aliases:	bra, jb, jnb, jz, jnz
Operand register(s) and aliases:	N/A
Result Register(s):	pc

There is only one trigger register but the aliases determine the condition that is required to be met in order to follow the branch. Branches not followed are treated as NOPs.

All branch destinations are determined by absolute addresses, there is no support for relative branches. This is due to the Program Counter (PC) using an MFSR as its increment operator [42]. MFSR propagation delays are lower since there is no carry logic and the logic depth is one.

The PC is only 10-bits wide, and the immediate field of a TTA16 instruction is up to 11-bits wide, so loading the PC with the value from the immediate field of the instruction allows all locations in instruction memory to be reached. Registers can also be used as sources for branch instructions, since the PC is treated as just another FU register, any value can be written to it.

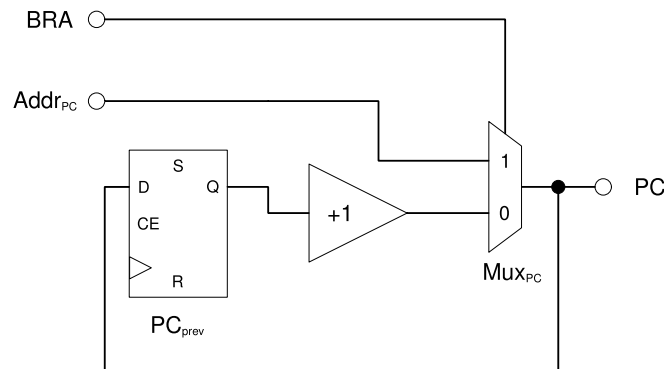


Figure 4.6: The PC increment and branch unit.

The incrementer (labelled “+1”) is actually implemented as an MFSR to reduce PC calculation latency [42].

There is an instruction fetch delay (due to pipelining) of one clock cycle. This delay causes an extra cycle of latency between the branch instruction being executed and the instruction at the branch address fetched and ready to execute. The gap of one cycle, between the branch being executed and the instruction fetched from the branch destination, there are a couple of options as to what the CPU can do:

- It can either stall (or issue a NOP).

- It can execute the instruction following the branch anyway, since it has already been fetched.

The instruction(s) immediately following a branch, which are executed, are called branch delay slot instructions. Processors that support the execution of instructions in branch delay slots include SPARC [32], MIPS [3], PA-RISC, ETRAX CRIS, and SuperH [15]. This is also the approach that TTA16 uses. This means that the instruction following a conditional branch is always executed, irregardless of whether the branch is taken or not. It is often possible to fill this slot with a useful instruction (see Figure D.4, but if it is not possible to do something useful, this slot must be filled with a NOP.

Register File

Functional Unit Name:	Register File
Trigger register(s) and aliases:	r0-r15
Operand register(s) and aliases:	N/A
Result Register(s):	r0-r15

These are 16, 16-bit, general-purpose registers, and function in the same way as general-purpose registers in RISC architectures. Up to two of these register can be read per clock cycle, and up to one write. All transports can read from the RF, but only transport three can write results back to the RF. (See Appendix D to obtain more details on the quirks and restrictions on RF use.)

The TTA16 Register File (RF) was implemented using Spartan-3 distributed RAM resources. There is hardware support for using this distributed RAM in dual-port mode allowing a read and a write to occur within each processor clock cycle (though reads are actually asynchronous with this primitive, but they were pipelined anyway to increase performance).

Bitwise Operations

Functional Unit Name:	Bitwise Unit
Trigger register(s) and aliases:	and, nand, or, xor, cmp
Operand register(s) and aliases:	com
Result Register(s):	bits, zf

The bitwise FU can perform one of four logic operations on the 16-bit input values, and these are selected by the two-bit mode input, which is set depending on the alias used. These modes are shown below:

Logical Operation	Mode Bits
AND	00
NAND	01
OR	10
XOR	11

The CMP (CoMPare) alias performs an XOR within this FU, but also simultaneously triggers a subtract operation. The intended side-effect is that both the CPU flags, the ZF and BF, are set at once, which can be useful for conditional branches.

The bitwise FU also asserts/deasserts the ZF, depending on whether the result of a bitwise operation contains all zeroes, or not. To calculate the value of the ZF, the carry-chain is used (see Figure 4.7), saving an additional cycle by not needing a separate zero test unit.

The actual output of the bitwise FU is the inverse of the desired result, due to the carry-chain's MUX requiring a "1" input to propagate the carry signal. But there is no performance penalty when inverting the bitwise result to get the desired value though, since the TTA transport multiplexers perform this without penalty. This is because these multiplexers are implemented using the same arbitrary function LUT4s (four-input Look-Up Tables) as used in this FU.

The latency once synthesised is about 4.5 ns, depending on the optimisations used and how complex the overall design is. The total logic cost, including the zero flag is only 17 LEs, less than nine slices. The Verilog source for this FU uses a Xilinx specific primitive, MUXCY, which means that this FU will only synthesise when targeting a subset of Xilinx FPGA architectures. The MUXCY primitive had to be explicitly instantiated since the Xilinx optimiser failed to pack the design using this primitive automatically.

The carry chain within the Spartan-3 FPGA runs vertically so all elements of the bitwise FU are placed vertically adjacent, but the Xilinx Synthesis Tool (XST) handles this automatically, explicit use of the MUXCY primitive ensures this.

Subtractor

Functional Unit Name:	Subtractor
Trigger register(s) and aliases:	sub, sbb, cmp
Operand register(s) and aliases:	com
Result Register(s):	diff, bf

The Spartan-3 has fast carry chain logic which is designed to increase the speed of arithmetic operations (typically additions, subtractions, and multiplications) that are implemented using general-purpose logic resources. This carry chain also supports the borrow functionality needed

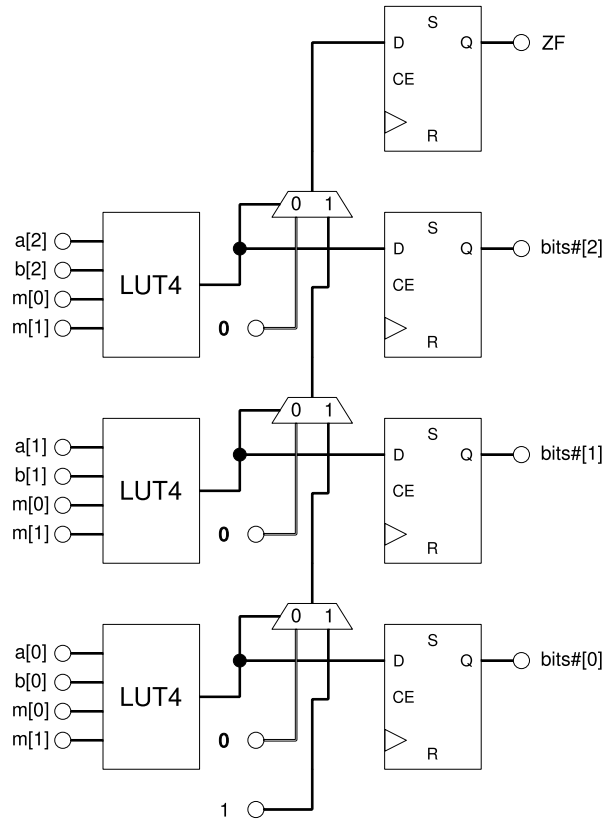


Figure 4.7: 3-Bit, Bitwise Operations Unit.

This shows how a 3-bit, bitwise operations FU implementation would look when synthesised for the Spartan-3 architecture. The inputs are $a[2:0]$, $b[2:0]$, and the mode $m[1:0]$. The bitwise operations FU takes advantage of the fast carry chain logic to compute the value of the ZF (Zero Flag) simultaneously with the bitwise result, and so implementing the ZF uses only one additional DFF (See the accompanying text for a full description).

for subtraction. Consequently, subtract operations take about 5 ns for a 16-bit addition in a synthesised design, but only an extra 0.064 ns for each extra bit width of the input numbers, as long as PAR does a good job (see Figure 4.7 for an example of the fast carry/borrow chain).

The borrow chain implementation uses the MUXCY primitive (like with the bitwise FU shown in Section 4.5.4) but also uses an XORCY¹¹ primitive and a different connection scheme. Fortunately the XST places these automatically for additions and subtractions, since this is a common use-case, so implementation of subtractors only takes as many LEs as there are bits in the widest input number.

¹¹The CY suffix of XORCY denoting that its intended use is to be part of the carry chain, like with the MUXCY primitive, but they can be freely instantiated for any purpose the designer can think of. They have a more restrictive connection scheme than LUT4s though, so often they are only useful as carry/borrow logic.

The subtractor FU also supports the subtract-with-borrow operation which allows subtract operations with numbers longer than 16-bits to be implemented quite efficiently. This is achieved by having a processor Borrow Flag (BF) which is set by any subtract operation, but only used for SBB and conditional branch operations. This BF requires an additional LE which brings the total for TTA16's subtract FU to 17 LEs.

As explained above, the CMP alias simultaneously triggers an XOR in the bitwise FU and a SUB within this FU.

Multiplier

Functional Unit Name:	Multiplier
Trigger register(s) and aliases:	mul
Operand register(s) and aliases:	com
Result Register(s):	plo, phi

Within every Spartan-3 FPGA there are embedded multipliers capable of operating at about 200 MHz [45]. These embedded multipliers take two 18-bit (or less), signed integers and calculate a 36-bit, signed product. Both the upper (PHI) and lower (PLO) halves of the 36-bit product can be selected, so supporting multiplications on numbers wider than 18-bits is fairly easy and efficient¹².

Within each embedded multiplier there exists a pipeline register, use of which is optional, so a registered multiplier can be implemented using no additional general-purpose logic elements.

Due to the way Xilinx chose to implement these multipliers, there exists a combinatorial delay component on either side of the pipeline register. When synthesising TTA16, this combinatorial component causes the path from the multiplier output to transport register to be the critical path. This is the path which limits TTA16 performance to about 190 MHz.

There exists a synthesis flag for TTA16 to optionally use an extra pipeline stage to hide this combinatorial component, but this also increases the size of TTA16 which, due to the extra routing costs, does not actually increase clock rate either, and adds an extra cycle of latency penalty on multiply operations too, so this option was left disabled.

¹²It is still quite complicated since the processor does not have an adder for adding partial products. Negations and the subtractor has to be used.

Wishbone Bus Interface

Functional Unit Name:	Load-Store Unit
Trigger register(s) and aliases:	rad, wad
Operand register(s) and aliases:	mem, rseg, wseg
Result Register(s):	mem

To support large addresses, with a native CPU data size of 16-bits, a memory segment model is used. Each segment consists of 2^{16} (65536), 16-bit words, or 131072 bytes. There can be up to 2^{16} of these segments to give a total addressable memory limit of 8 GB.

The registers **rseg** and **wseg** are (respectively) the segments for memory read and write accesses, and these are set via writes to the special function registers FU. These segment values are concatenated with the lower 16-bit addresses, **rad** and **wad**, for reads and writes respectively, to produce the full 32-bit, word-aligned address.

The contents of the **mem** register are written to the Wishbone bus when WAD is triggered, and the contents of **mem** are set after the Wishbone bus read transaction, triggered by a write to **rad**, completes. Both the registers called **mem** are actually different registers too, so a **mem** to **mem** transport is needed when performing a memory copy.

Memory access, for example, has an undeterminable latency, it depends on Wishbone bus activity and whether the SDRAM is refreshing. When a Wishbone access is triggered, TTA16 asserts the interlock signal to stall the CPU until the Wishbone bus operation completes. Due to this stalling, Wishbone access is the slowest CPU operation, but interlocking ensures correctness.

Special Function Registers

Functional Unit Name:	System Registers
Trigger register(s) and aliases:	msr
Operand register(s) and aliases:	com
Result Register(s):	N/A

There are currently five special registers implemented, two are the segment registers, **wseg** and **rseg**, as detailed above. The other three are used for modifying the contents of the instruction memory.

MSR	Register
Index	Name
0	RSEG
1	WSEG
2	IADR
4	IDAT _{LO}
5	IDAT _{HI}

IADR is the 10-bit address of the instruction to overwrite, and since instructions are 32-bits wide, it takes two 16-bit writes, the LSW of the instruction is written to IDAT_{LO}, while the MSW is written to IDAT_{HI}.

The primary purpose is to allow new pages of instructions to be fetched from main memory, though it would be possible to use this mechanism for self-modifying code, like setting the value of a constant before entering in a tight-loop.

4.5.5 Instruction Format

TTA16 has just one instruction format (see Figure 4.8), a list of data moves between registers and a some immediate data fields. The immediate is used for specifying constants used in a program, like a count value or an address offset, or for specifying a register index when performing a RF read or write.




Figure 4.8: TTA16 instruction format. More details are available in the TTA16 programming guide (see Appendix D)

For performance reasons, TTA16 features overlapping instruction RF index bit-fields. The RF read-index bit-fields, for the current instruction, are stored in the previous instruction so that the processor will have fetched them already when the current instruction is fetched. The reason this design decision was taken is because it saves one pipeline stage for the targeted processor frequency, which was > 150 MHz. One less pipeline stage means fewer pipeline registers resulting in a smaller processor.

From a programming point of view, this increases complexity slightly, but is not a significant problem and is supported by the assembler. Example code is listed in Appendix D.

4.5.6 Instruction Memory

With a 32-bit instruction being issued nearly every clock cycle, instruction fetch bandwidth must be high, at 150 MHz, about 600 MB/s. Also, since branches need to have a penalty as small as possible, instruction fetch latency must be minimised too. Modern high-performance processors use instruction caches to solve these problems but there exist some additional issues which is why a Xilinx Block RAM was used as instruction memory.

Below is a list of arguments against a traditional instruction cache, and in favour of a BRAM storing instructions.

- The added logic of an instruction cache, more specifically its sense logic, running at the CPU frequency, will make timing closure even more difficult.
- The pseudo-random count order of the MFSR means poor cache behaviour, along with poor memory addressing behaviour in general¹³.
- The processor needs code to run at power-on, a boot-loader, and this will need to be stored in a BRAM, so using two BRAMs as instruction memory saves one BRAM for a boot-loader ROM. (It would be quite complex to do this using the BRAM built into the instruction cache.)
- Since a cache requires sense logic (see Section 5.2.8), total latency is higher, either increasing the cost of branches, or reducing maximum frequency, relative to a simple fetch from a BRAM.

The obvious disadvantage of using BRAMs is that the number of 32-bit instructions that can be stored in the 4 kB of BRAMs is only 1024 . There is a mechanism for fetching instructions from memory and storing them in the BRAMs. New instructions can be fetched from the SDRAM and written into the local BRAMs by writing to the appropriate MSRs. More information is available in Section D.2.1 of the TTA16 programming guide.

¹³Solutions to the cache coherency problem of MFSR PCs were examined, the most promising being a hybrid PC, part standard counter, part MFSR, but was not implemented. It could be a future area of research.

4.6 RISC16

Module:	risc16	
Description:	A 16-bit, RISC architecture, processor logic core designed for use with FPGAs.	
Related Files:	/rtl/cpu/risc16/risc16.v, /rtl/cpu/risc16/risc16.xml, /rtl/cpu/risc16/risc16.v, /rtl/cpu/risc16/defines.v, /rtl/cpu/risc16/fetch.v, /rtl/cpu/risc16/branch.v, /rtl/cpu/risc16/decode.v, /rtl/cpu/risc16/execute.v, /rtl/cpu/risc16/memory.v, /rtl/cpu/risc16/risc_alu.v, /rtl/cpu/risc16/risc_rf.v, /rtl/cpu/fastbits.v	
Testing Files:	/sim/cpu/tta16_tb.v	
Author:	Patrick Suggate	License: GPL

RISC16 has been designed to be easier to program in assembly language than TTA16. It also has a similar instruction set to other RISC processors [16, 15] which have C compiler ports available. It should therefore be possible to port the LCC [16], or maybe GNU Compiler Collection (GCC), C compiler to this processor, but this was beyond the scope of this work. The trade-off for this easier-to-program approach is a larger and slower processor than TTA16. It is still faster and smaller than other similar processors which were evaluated though (see Table 7.2).

4.6.1 General RISC16 Characteristics

The RISC16 processor architecture has the following features, which are typical of many RISC architectures:

- RISC16 has a fixed-width instruction word format¹⁴ and therefore simpler instruction encoding when compared to CISC designs. Early RISCs used 32-bit data and instruction words, but some later RISC processors have used other instruction word widths. Like RISC16, some have used 16-bit wide instruction words [16, 15, 46].
- Simple memory addressing modes. RISC16 only supports memory access using load and store instructions, as is typical with RISC architectures [19, 9].
- RISC16 has 16, 16-bit, general-purpose registers. A large number of general-purpose registers is typical of RISC processors. These registers can all be used interchangeably, because they

¹⁴Both ARM and MIPS have had 16-bit instruction word support for processors targeted at embedded applications. Typically, they mix 16-bit and 32-bit instructions within the same program and claim a 40% reduction in code size [19]

function identically¹⁵.

- RISC architectures are typically heavily pipelined and RISC16 features a five-stage pipeline. The RISC16 pipeline is considered a classic RISC pipeline [37].

4.6.2 Design Choices

A five-stage pipeline was chosen because a greater number of pipeline stages usually allows a processor to run at higher clock rates [19, 31]. A disadvantage to highly pipelined processors is often a greater silicon area usage due to the extra control logic and pipelining registers. The extra register usage penalty was minimal with RISC16 since every logic element within the Spartan-3 family contains a DFF [45], so there was little area increase due to extra registers alone.

The greatest increase in logic usage, as a result of more pipeline stages, was due to the data forwarding mechanism. Data forwarding is typically more complex, therefore using more logic gates, when there are more pipeline stages. More specifically, when there are more pipeline stages between the register-read stage and the register write-back stages, more logic is needed for either interlocking and/or data-forwarding. There are processors with fewer pipeline stages than RISC16, so their forwarding mechanisms are simpler [35], or not even present [16].

Another significant design choice with the RISC16 processor was the omission of an `ADD` instruction. This decision allows the Arithmetic Logic Unit (ALU) to be faster, resulting in shorter processor cycle times, but does not cause a significant programming penalty. This is because even in the worst cases, an addition operation can be completed by issuing two subtract instructions instead (see Table 4.4). Additionally, many algorithms can be modified to use subtract operations instead of additions. A simple example is that many iteration operations, like the common `for-loop`, can be rewritten to decrement the counter instead¹⁶ so that there is no extra penalty.

Traditional RISC	Subtract-Only RISC
<code>add r0, r0, r1</code>	<code>sub r1, #0, r1</code> <code>sub r0, r0, r1</code>

Table 4.4: RISC16 does not have an addition instruction, but the same operation can be completed with two subtract instructions.

¹⁵The Intel 8086, for example, had special registers for accessing memory, like a dedicated stack pointer, a dedicated stack base-pointer, several memory segment registers, and a couple of memory index registers that some instructions could only be used with [28].

¹⁶This trick is a common optimisation technique with some processors anyway, like the 8086, since the compare with zero operation is typically faster than the compare to count-limit operation.

To have both an ADD and a SUB instruction would have required using a different instruction encoding resulting in a more complex, therefore slower, instruction decoder¹⁷. Additionally: either more logic would have been required to include an adder, and the necessary multiplexers; or would have required using a single ADD/SUB functional unit, which due to how they are implemented within a Spartan-3 FPGA, would have added about 1 ns to the cycle period.

Unconditional branching to a constant address typically requires two instructions, one to set the upper 12-bits (the `i12` instruction), and then the actual branch instruction, since this only contains a 4-bit immediate constant. Since only a small percentage of program flow control instructions have this form [28], it does not cause a significant performance penalty.

4.6.3 The RISC16 Pipeline

RISC processors are typically pipelined, meaning that the work of executing an instruction is spread over multiple clock cycles. Typically, new instructions are fetched and enter the pipeline every clock cycle. Pipelining is also form of ILP, as multiple instruction are executed at simultaneously, although each is in a different stage of the execution pipeline.

The classic RISC pipeline [37] has five stages¹⁸, as does the pipeline of RISC16 (see Figure 4.9). These five stages are [19]¹⁹:

- F: instruction Fetch. Fetches an instruction, pointed to by the PC, from the instruction memory.
- D: instruction Decode. Decodes the instruction and fetches the operands.
- X: eXecute. Performs an ALU operation on the operands.
- A: memory Access. Loads/stores data from/to the system memory.
- W: Writeback. Updates the RF.

The design target for performance was an 8 ns cycle time, corresponding to 125 MHz²⁰, and this limit was determined by the two slowest paths:

- The ALU subtract operation: hindered by a long borrow-chain propagation delay.

¹⁷The current instruction function field is encoded with just a 3-bit value, adding both an addition and an addition-with-carry would have required a 4-bit field.

¹⁸Other RISC processors that had the classic, five-stage, RISC pipeline were the original MIPS, SPARC, Motorola 88000, and later versions of the DLX[37]

¹⁹Different sources give different names for these stages, but they are actually the same.

²⁰The goal is 125 MHz since designs get slower as more logic is added when using automatic placement when synthesising using the Xilinx PAR tool. The final OpenVGA design uses a large percentage of the device so conservative constraints for individual components allows the final timing constraints to be more easily met.




Figure 4.9: RISC16 has five pipeline stages.

- The calculation of the next PC value: which is either an increment operation performed on the current PC value, or if a branch instruction has been issued, can depend on the evaluation of conditions codes to choose a branch location.

With the typical propagation delay for a single layer of Spartan-3 logic taking 3-4 ns, depending on the situation, including expected routing and registering delays, a design using just one layer of logic will easily meet the design target. This was the goal for most of the design since it would allow plenty of timing “slack”. The purpose of so much slack was to allow the Xilinx automatic Place And Route (PAR) tool to have more flexibility when placing and routing the logic. This is to allow the optimiser to expend more effort achieving the timing constraints for the two critical paths.

4.6.4 RISC16 Instruction Set

RISC16 supports five different instruction formats. Writing instructions with these different formats are covered in the RISC16 programming guide in Appendix E. This section describes the bit-fields of these formats, which are shown in Figure 4.10.

The three most-significant bits (**OP**, for **OP**eration-code) encode either the instruction format, or the instruction operation if the instruction is of the three-operand **rri** format. The set of **rri** instructions are load-word, store-word, and subtract with immediate (**lw**, **sw**, and **subi** respectively).

For ALU operations, except multiply, the **SF** bit determines whether the condition code flags are updated by the current instruction. The condition code flags are evaluated for conditional jumps. With load and store operations, the **SF** flag selects one of two segment registers to use for

	15	12	8	4	0
rri	OP	SF	RD	RS	IMM4
rr	0x0	SF	RD	RS	FN CR
ri	0x1	SF	RD	FN CR	IMM4
i12	0x7	X	IMM12		
bx	0x6	CND	IMM10		

Figure 4.10: RISC16 instruction formats.

the upper word of the memory address. These segment register values are set using **msr** instruction. The multiply instruction uses the **SF** bit to determine whether the upper or lower half of the 32-bit product will be stored.

The **FN** bit-field selects either the ALU operation, or an unconditional branch. This bit-field is only present within **rr** or **ri** format instructions. The position of this bit-field is different for each format as well. The related **CR** bit selects whether the current instruction is to write back to the register file. Some instructions like bit-test (**test**) are non-destructive, so this bit is zero, but for most instructions, it is set.

The **RD** and **RS** bit-fields choose destination and source registers to be used for the current instruction. When used with the **rr** and **ri** instruction formats, **RD** is used as both a source and destination register. For **rri** instructions **RD** is only used as a destination. Register **RS** is always selects a source register when present.

The **rri** and **ri** instruction formats have a 4-bit signed immediate constant field, **IMM4**. The 4-bit is sign extended, when used without the **i12** instruction, and can therefore have a value from -8 to 7. When an immediate constant of more than 4-bits is required, the immediate override instruction prefix must be used (format **i12**). This sets the upper 12-bits for just the following instruction. The **IMM12** bit-field becomes the upper 12-bits of the immediate constant of the following instruction.

There are 1024 instructions within the RISC16 instruction memory, so a 10-bit absolute address can reach all locations. This is encoded as **IMM10** within the **bx** instruction format. The **CND** bit-field selects one of eight branch conditions. The processor condition code flags are evaluated and if the branch condition code matches the processor flags, the branch is taken.

4.6.5 Instruction Memory

A 2 kB BRAM is used for instruction memory, and this has space for 1024 instructions. The rationale for using a simple BRAM as instruction memory for RISC16 was the same as with TTA16 (see Section 4.5.6). The key difference between RISC16 and TTA16 instruction memory is that firstly, RISC16 instructions are only 16-bits wide, so just one BRAM can store as many instructions as TTA16 can with two. Secondly, TTA16 has a mechanism for fetching instructions from memory and storing them in the BRAM, but it has yet to be implemented within RISC16, this is future work. This is because RISC16 was created primarily to compare with TTA16 and was not needed during testing and evaluating RISC16.

4.6.6 Functional Units

A minimal set of FUs was chosen so that the processor is as small and fast as possible. The units chosen allow most typical CPU operations to be performed, but some operations can be slower, and requiring more instructions, than typical RISC processors. Most of the FUs are exactly the same as used with TTA16 (see Section 4.5.4), so only the FUs that differ from TTA16 will be covered in detail here.

Register File

This is a 16 entry, 16-bit wide, RF implemented using what Xilinx calls distributed RAMs. Two of these RAMs can be connected together to yield a dual-port RAM. Since a RISC CPU typically has two RF read ports and a write port, RISC16 required two parallel banks of dual-port memory to support two reads and a write every clock cycle. The write-ports of both banks are written to simultaneously, and with the same data, so that they remain coherent. The result of having two banks of distributed RAM is that the RF is quite a large component of RISC16, 4 LUTs per bit-width, resulting in 64 LUTs.

Machine Special Registers

These are currently just the memory segment registers of RISC16, they contain less functionality than their TTA16 equivalent (see Section 4.5.4), though it would be straight-forward enough to make them the same.

Subtract Unit

The only difference compared with the TTA16 subtractor is that there is one extra output flag, indicating whether the difference is negative or positive. This is called the Negative Flag (NF), and indicates a negative number when asserted. The subtract unit only updates the processor's

condition codes when the instruction's **SF** bit is set. The **NF** is used by the branch unit when evaluating condition codes for jumps.

Branch Unit

The branch unit is similar to that used with TTA16 except that there is one extra condition code, the **NF**. There are also eight conditional jump instructions (see Section E.2) compared to the four that TTA16 has.

RISC16 does not use branch delay slots (see Section 4.5.4), the processor stalls for two cycles while the branch destination instruction is fetched. Since it is not always possible to do something useful with the instruction within the branch delay slot, eliminating it improves code density slightly, which was a goal of RISC16.

Multiplier

The RISC16 multiply is also implemented using a Spartan-3 embedded multiplier. This allows a multiply to take just one processor cycle, which is less than many simple RISC processors [43, 3].

4.6.7 Data Forwarding

In a pipelined processor, a data dependency occurs when an instruction depends on a result, which is still within the pipeline, calculated by a preceding instruction²¹. Since the result is still within the pipeline, and has not been written back to the RF yet, the processor needs to either wait until the writeback has been completed, or have a mechanism for transferring the result back into the ALU. The latter approach is known as data forwarding (or data bypassing). If the processor does not handle data dependencies correctly, data hazards will result, causing erroneous computations.

Consider the RISC16 assembly fragment in Figure 4.11, the register **r0** is used as a destination in the first instruction and a source in the second. RISC16 has a five-stage pipeline so the result of the first instruction (**r0**) is present on the outputs of the ALU when the second instruction is ready to be evaluated within the ALU. This creates a data dependency as the result of the first instruction is needed before it has been written back to the register file.

```
sub  r0, r1, -3
sub  r3, r0, -2
```

Figure 4.11: Data Dependency.

Figure 4.12 shows the bypassing mechanisms used within RISC16. Results from preceding

²¹For heavily pipelined processors, data dependencies can be caused by instructions dispatched many cycles earlier than the current instruction.

instructions, that are still within the pipeline, can be forwarded back to the ALU's input registers. This is to avoid hazards caused by data dependencies. The RISC16 control logic is used to select the data source for the ALU input registers (ID_a and ID_b). This is achieved by examining destination register fields from previous instructions with the register sources for the current instruction.

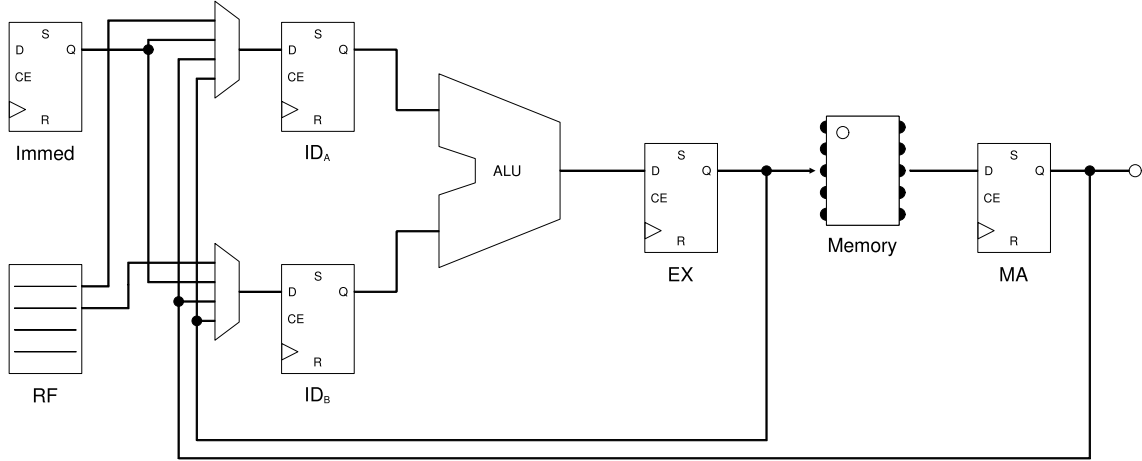


Figure 4.12: This shows the bypassing scheme used by RISC16.

Due to the cost (logic and latency) of multiplexers, data forwarding was only implemented for data hazards caused by instructions with a gap of one or two cycles. For data hazards caused by immediately following instructions, the processor stalls for one cycle, and then forwards the data.

The exception to this is a memory load. If the result of a memory load is used immediately after the load instruction there is no bypassing issues since the processor stalls for a minimum of three cycles automatically, therefore no bypassing is needed.

The data hazard prevention logic consisted of bypassing and stalling logic, with the data hazard stall logic needing only an extra comparator, and the bypassing logic using a two 16-bit wide multiplexers and some comparison logic. The complete cost in logic was about 30 logic slices, and there was very little effect on processor cycle times. This was because neither the comparison logic or the multiplexers were within critical timing paths. Without this logic, data hazards would have to be avoided by the programmer, or by a compiler, and the code produced would potentially be slower and less compact too.

4.7 Processor Logic Cores Summary

The objectives of the TTA16 and RISC16 designs were high data-conversion throughput, low logic usage, and a large address space. These were achieved, with the results of TTA16 and RISC16 processor logic cores, when synthesised for a Spartan-3, are shown in Table 4.5. TTA16 also has the

highest performance of all evaluated processors. The processors evaluated are listed in Table 7.2 and TTA16 has the highest operating frequency and the second smallest size.

Property	TTA16	RISC16
Speed (MHz)	190	140
Size (Slices)	200	320
Has Assembler?	Yes	Yes
Has C Compiler?	No	No
Wishbone	Yes	Yes
Bus Width (bits)	16	16
Instruction Width (bits)	32	16
Instruction SRAM Size (kB)	4	2
Data Forwarding	No	Yes
Data-Hazard Detection	No	Yes

Table 4.5: Processor properties.

TTA16 has very good performance but it unfortunately comes at a cost. It is very time consuming to write small and fast assembly routines for TTA16. This is mostly due to the explicit ILP, the limited connectivity between the transports and the FUs, the interleaved instructions, and that this particular TTA processor performs no hardware checks for data dependency hazards, this must be done by the programmer or unexpected results will occur.

TTA16 has a three-stage pipeline and no data-dependency checking hardware. Data dependencies must be manually checked because there are typically three instructions being executed simultaneously. The results of previous instructions take two cycles to become available to the data-transport. Transporting a value from a FU result register before two clock cycles have elapsed will simply move an earlier value than the one currently being calculated. Consequently, efficient TTA16 assembly code is hard to write and also difficult to read and gain an understanding of (see Section D.4).

Also worth noting is that the TTA16 instruction word is 32-bits wide, twice as wide as the instruction word of RISC16. Even highly optimised TTA16 assembly is typically much larger than the equivalent RISC16 assembly routines. The memory copy routines for each processor are listed in their own programming guides (Appendix D for TTA16 and Appendix E for RISC16). A similar number of instructions are required for each processor, giving TTA16 a performance advantage due to the higher clock frequency. But RISC16 instructions are half the size so RISC16 significant has a code density advantage.

Chapter 5

Local Memory and Caching

OpenVGA features two on-board memory ICs, a SDRAM and a Serial PROM, and to improve processor-to-memory performance, there is a data cache logic core too. The SDRAM IC has a capacity of 8 MB and is used for storing frame buffer data and as a general-purpose local memory. The memory and controller were successfully tested at 120 MHz, giving a peak transfer rate of 240 MB/s, but the read latency is a minimum of eight Wishbone clock cycles. Memory latency can be considerably higher than this because the memory is shared by multiple devices. There is therefore the potential for congestion. And since DRAMs need to be periodically refreshed, the memory is also not available during these refresh cycles either.

The Serial PROM stores the Spartan-3 configuration information which is read by the FPGA at power on¹. Since not all of the capacity is used by the configuration information, OpenVGA can store processor firmware and other data in this ROM too. At initialisation, the processor (either TTA16 or RISC16) transfers this extra data to the local memory, the SDRAM. All processor accesses to local memory pass through the processor's data cache. This data cache significantly improves processor performance due to decreased memory access latencies. To further improve performance, the data cache operates at the processor frequency, up to 150 MHz, and has a fast-hit calculation path with a latency of zero clock cycles.

The two principal demands on system memory throughput are due to the processor - converting display data between formats, and the display redraw logic. The redraw logic requires a lot of memory bandwidth, but has a 2 kB prefetch buffer to maximise throughput. This allows the use of large burst sizes to reduce the cycles spent initiating a transfer (see Section 6.3.2).

The processor's memory accesses are all single-word (atomic) transfers. Without a cache the full memory access latency penalty would apply to each read or write, which due to congestion averages at about 45 processor clock cycles in 640x480, 16-bit colour mode. By using a cache, and a Direct Memory Access (DMA) controller, the average memory access latency was reduced

¹Xilinx refers to these types of ICs as Platform Flash

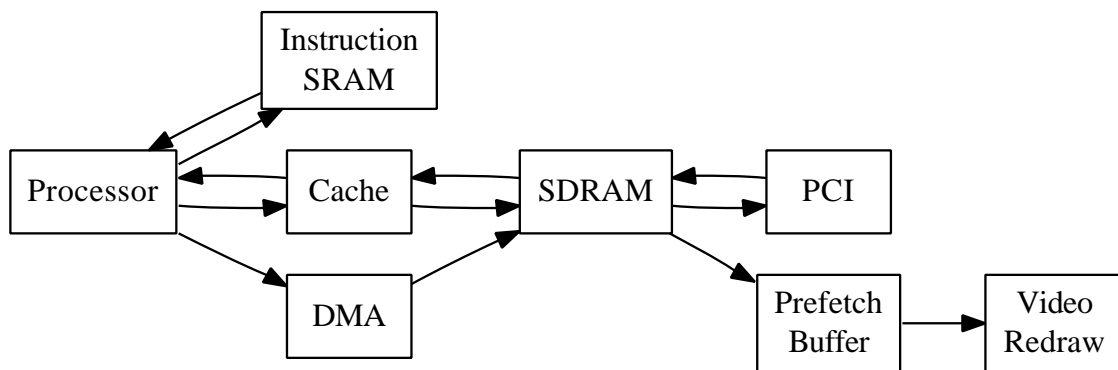


Figure 5.1: OpenVGA memory overview presented as a directed graph.

to about 5 processor cycles (see Section 5.2).

5.1 SDRAM Controller

Module:	wb_sdrdram_ctrl
Description:	A parameterisable, Wishbone-compatible, SDRAM controller which supports burst and atomic reads and writes.
Related Files:	/rtl/mem/wb_sdrdram_ctrl.v, /rtl/mem/ddr_datapath.v
Testing Files:	/sim/mem/wb_sdrdram_ctrl.tb.v
Author:	Patrick Suggate
License:	GPL

Single Data Rate (SDR) SDRAM was chosen over DDR1/2 SDRAM because it is less difficult to implement, especially when considering that OpenVGA uses a two-layer PCB (see Section 3.1). OpenVGA uses a standard Micron SDRAM IC, part number: MT48LC4M16A2-75 (which is an 8 MB, 16-bit data bus, 133 MHz capable device[29]).

The theoretical peak data throughput is significantly less for SDR vs. DDR, fewer FPGA I/O pins are needed and PCB complexity is lower. Additionally, signal integrity issues were assumed to limit throughput to below the maximum rate of the specification anyway, and they did. During OpenVGA's SDRAM testing, maximum error-free throughput of the SDRAM was at 120 MHz.

5.1.1 Controller Overview

The controller's major functional components are: the state machine logic for initialising and issuing commands to the SDRAM; the data path that sends data to, and latches data from, the SDRAM, via tri-state drivers; the Wishbone bus interface; and timers to ensure that the SDRAM is regularly refreshed and that SDRAM timing specifications are not violated. The SDRAM controller

is a Wishbone slave device, it is a target of Wishbone transactions but it does not initiate them.

State Machine and Other Control Logic

The SDRAM control logic clock is driven by the Wishbone bus clock so asynchronous FIFOs are not needed to cross clock domains. The controller state machine (see Figure 5.1.1) and other control logic is more complex than the data-path logic, therefore requiring more layers of logic, so the control clock has half the rate of the data clock. Since the Wishbone bus clock is 50 MHz, the control clock is 50 MHz, but the data clock is 100 MHz, though higher frequencies were used during testing.

To achieve the maximum theoretical transfer rate, 200 MB/s at 100 MHz, the data burst size is two cycles (four bytes), so read and write commands are only required to be issued at the Wishbone clock rate of 50 MHz to keep the data path saturated.

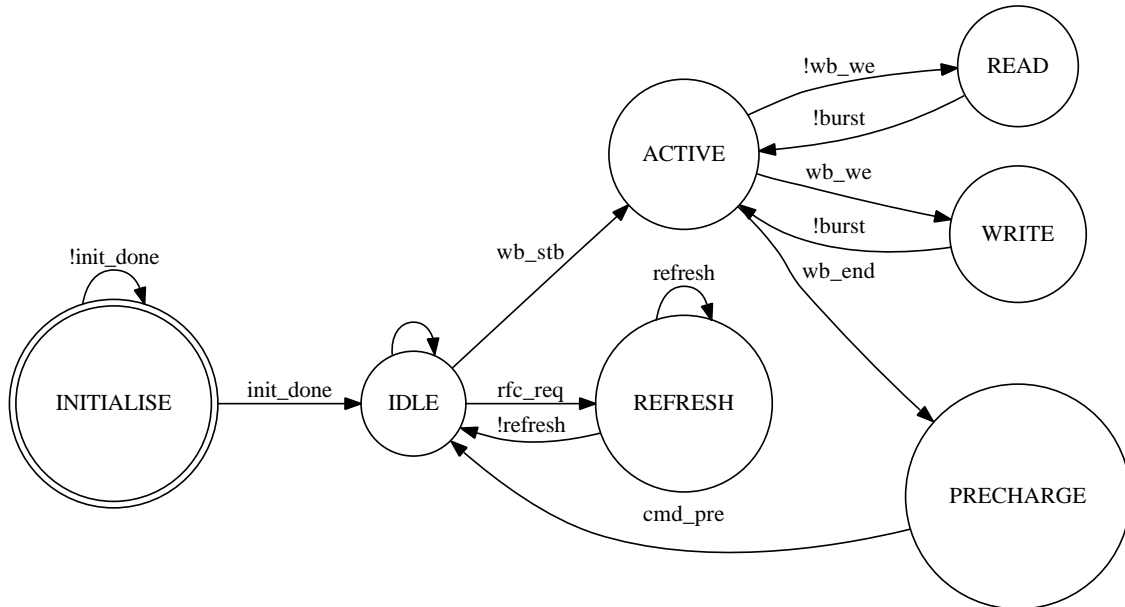


Figure 5.2: Wishbone clock domain state machine.

The following is a brief description of the operation of the SDRAM controller state machine and the effect each state has on the SDRAM and the Wishbone bus. A complete datasheet which details all of the SDRAM commands, and the meanings of the symbols used below, is available from Micron [29].

- **INITIALISE:** Upon reset or power-on this is the state entered. The controller remains in this state until the initialisation sequence is complete:
 1. Wait 100 μ s after clocks stabilise, issuing either COMMAND INHIBITs or NOPs.

2. Perform a PRECHARGE ALL and issue NOPs or DESELECT until t_{RP} is met.
 3. Issue AUTO REFRESH followed by NOPs until t_{RFC} is met. Do this twice.
 4. LOAD MODE REGISTER with the operating parameters, $CL = 2$, burst size = 2 .
Issue NOPs until t_{MRD} is met.
 5. Initialisation is now complete, change state to IDLE.
- IDLE: Accepts memory requests from the Wishbone bus, which causes a change in state to ACTIVE, or enters the REFRESH state if requested by the refresh timer.
 - REFRESH: Issues an AUTO REFRESH command and waits until t_{RFC} is met before returning to the IDLE state. AUTO REFRESHes are issued every $15.625 \mu s$ on average (this is because the device requires 4096 refresh commands, one for each row, to be issued every 64 ms [29]).

If the controller is busy, when the refresh request signal asserts, it stays asserted until serviced. The refresh counter will keep track of up to eight missed refreshes to help ensure that the average refresh rate requirement is met.

- ACTIVE: Issue an ACTIVE command, which activates the SDRAM bank and row selected by the upper address bits provided by the Wishbone bus. If the Wishbone bus request is for a memory write then the next state is WRITE, if it is a read request, then the next state is READ. If the read or write has completed then the next state is PRECHARGE.
- WRITE: The lower Wishbone bus address bits select the SDRAM column, and 32-bit data from the Wishbone bus is written to the device, 16-bits at a time, and can be masked using the Wishbone data select signals. If the Wishbone bus transaction is a burst operation, stay in the WRITE state until either: the transaction completes; or a column boundary is reached, and then return back to the ACTIVE state.
- READ: The lower Wishbone bus address bits select the SDRAM column, and two 16-bit words are read from the SDRAM and assembled into a 32-bit word and placed onto the Wishbone bus. If the Wishbone bus transaction is a burst operation, stay in the READ state until either: the transaction completes; or a column boundary is reached, and then return back to the ACTIVE state.
- PRECHARGE: Once a read or write Wishbone transaction completes, the PRECHARGE command is issued. This command deactivates the selected bank and row. This is done since the SDRAM only allows one row in each bank to be activated at once, and the row must be precharged after an elapsed period as well. The simplest solution is to deactivate a row after each Wishbone bus transaction is complete.

Datapath

The memory controller datapath transfers data during both read and write transactions. Read data is latched from the IOBs and transferred onto the Wishbone Bus, and write data is transferred from the Wishbone bus to the IOBs. The datapath also implements byte masking so that single bytes (or none at all) can be transferred, rather than requiring the size all data transfers be multiples of 16-bits.

SDRAM, and therefore the SDRAM controller, uses the same data pins for both read and write data. The Spartan-3 FPGA has sophisticated IOBs which support bidirectional data transfer (see Figure 5.3). When the IOB T (Tri-state) signal is de-asserted the IOBs drive the data lines, when asserted the IOBs can receive data.




Figure 5.3: A simplified schematic of the DDR output primitive within a Spartan-3 IOB[45].

The Wishbone interface is 32-bits wide while the SDRAM has only 16 data lines. This is because the datapath makes use of the Spartan-3 IOB's DDR driver primitive. This allows transferring the lower half of the 32-bit data during the first half of the clock cycle and the upper half during the second half of a clock cycle².

5.1.2 Testing and Synthesis

Testbenches were created at various stages of development, a simple Verilog testbench was used initially. Micron provides a Verilog behavioural description file of an SDRAM IC, and this was

²Since SDRAM is not a DDR device, the SDRAM system clock still has to be driven at the same frequency as the data rate, in this case 100 MHz. Using the IOB DDR primitive simply allows the controller to run at half the frequency of the datapath, and remain synchronised to the Wishbone bus clock. Another benefit it that this trick also makes timing constraints easier to acheive.

used with simulations.

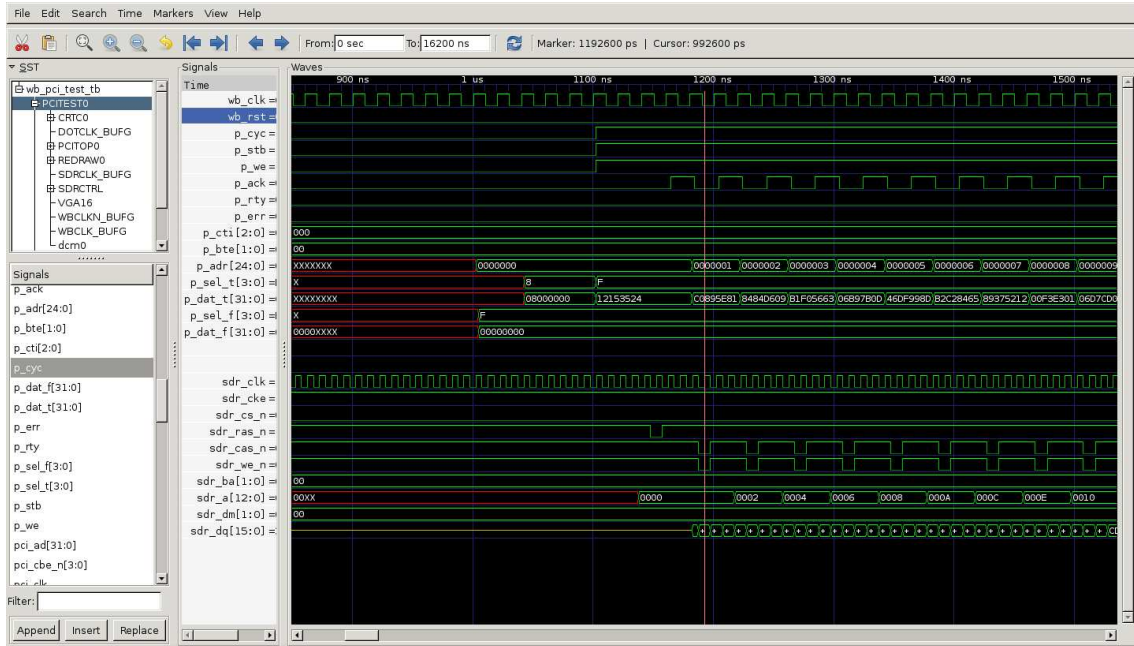
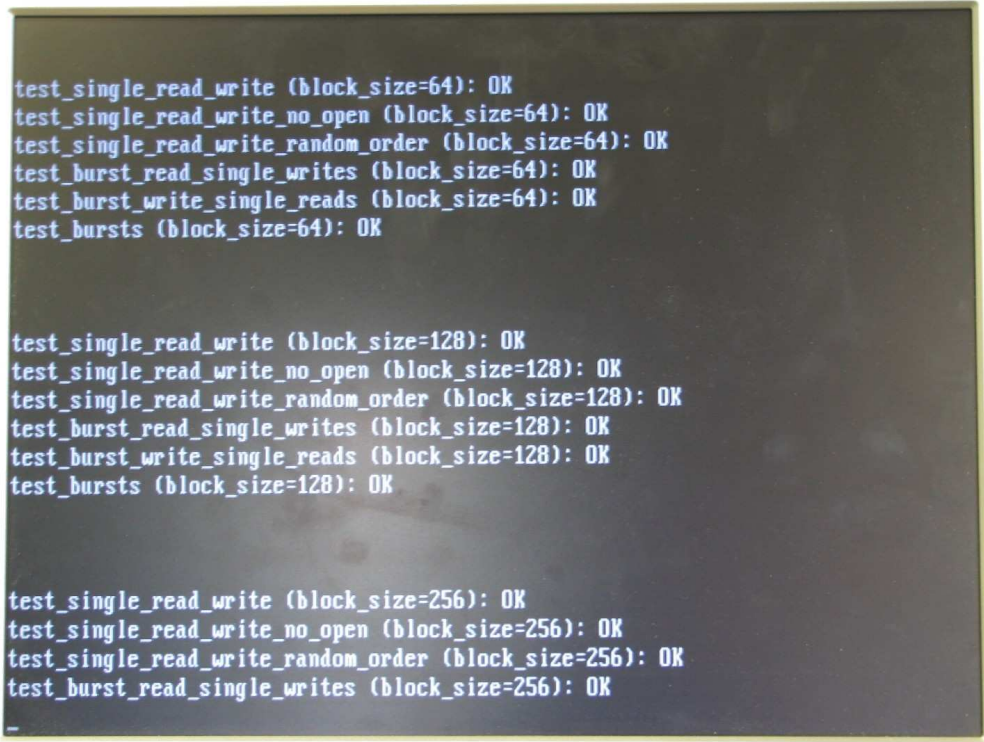


Figure 5.4 is a GtkWave³ screen capture showing the simulation waveforms of a write-to-SDRAM transaction. The transaction is initiated by the PCI bridge, transferred over the Wishbone bus to the SDRAM controller. The SDRAM controller then writes this data to the SDRAM device. This Wishbone bus transaction is a non-burst transfer (as determined by the Wishbone signals CTI and BTE).

The testing procedure consisted of writing blocks of data, of various block sizes, and write orders, and with different transfer modes, and then reading back and checking the data. The design passed memory tests without data corruption at up to 120 MHz (see Figure 5.5).



```
test_single_read_write (block_size=64): OK
test_single_read_write_no_open (block_size=64): OK
test_single_read_write_random_order (block_size=64): OK
test_burst_read_single_writes (block_size=64): OK
test_burst_write_single_reads (block_size=64): OK
test_bursts (block_size=64): OK

test_single_read_write (block_size=128): OK
test_single_read_write_no_open (block_size=128): OK
test_single_read_write_random_order (block_size=128): OK
test_burst_read_single_writes (block_size=128): OK
test_burst_write_single_reads (block_size=128): OK
test_bursts (block_size=128): OK

test_single_read_write (block_size=256): OK
test_single_read_write_no_open (block_size=256): OK
test_single_read_write_random_order (block_size=256): OK
test_burst_read_single_writes (block_size=256): OK
```

Figure 5.5: OpenVGA passing memory tests while running at 120 MHz.

Known Bugs

There exists just one known bug with the SDRAM controller and this has been narrowed down to a problem with write transactions. The system occasionally hangs when performing large PCI burst-write transfers to OpenVGA. From the tests performed, the following is known about this bug:

- This bug is intermittent. There seems to be a small probability of it occurring for any PCI write transactions with a burst length of more than 256 bytes.
- It occurred at all of the frequencies and memory timings tested, from 25 MHz to 120 MHz, and with column address strobe latencies of 2, 2.5, and 3.
- Only occurs when the SDRAM controller logic core is used, other memory cores have not demonstrated this bug.
- It has only occurred during write operations, all reads seem unaffected.
- Changing the SDRAM refresh period seems to have no effect on the frequency of these hangs.

The system freeze does not occur when a Wishbone SRAM logic core is used instead, indicating the problem is with the SDRAM controller, but all simulations have failed to reproduce this bug, and hardware testing to date yields too little information to locate the source of this bug. The current workaround is to use many small burst transfers, rather than a few large ones.

5.2 Data Cache

Module:	wb_simple_cache	
Description:	A Wishbone-compatible, 150 MHz, Spartan-3 optimised, 2 kB data cache which also features a fast-hit lookup path.	
Related Files:	/rtl/cache/wb_simple_cache.v, /rtl/cache/fetch_wb.v, /rtl/cache/tag_ram.v, /rtl/cache/cache_bram.v	
Testing Files:	/sim/mem/wb_sdram_ctrl.tb.v	
Author:	Patrick Suggate	License: GPL

A data cache for the OpenVGA processor was developed to significantly improve memory access performance. A large quantity of data needs to be processed when updating the contents of the frame buffer, such as when performing ASCII text to pixel conversion. The average latency of a single SDRAM memory access is high, about 45 processor cycles per word read or written (see Section 5.2.1), and the processor has to stall while waiting for the memory controller.

So that the processor's memory access operations have a high throughput, low average access latency is required. By using a data cache, average memory access latency, and therefore performance, is improved by a factor of about eight for the ASCII to pixel conversion algorithm (see Section 5.2.1).

The data cache has been designed to achieve a good compromise of the following characteristics:

- Logic resources used
- Clock frequency
- Design complexity
- Hit-time calculation
- Miss penalty
- Miss rate

For the sample workload, converting ASCII character data into pixel data, the cache miss rate is extremely low (0.7%), the clock rate high (≈ 150 MHz), the average latency is only 5.1 clock

cycles, and only about 80 Spartan-3 logic slices⁴, and a Block RAM, are used.

5.2.1 Cache Justification

How much of an improvement in average Clock cycles required Per Instruction (CPI) does a data cache give? First, assuming OpenVGA is operating at screen resolution of 640x480, in 16-bit colour mode (2 bytes/pixel), and with a 60 Hz redraw, the data rate (r_{RD}) is given by

$$r_{RD} = 640 * 480 * 2 * 60 = 36864000 \text{ Bs}^{-1}$$

The SDRAM controller operates at 50 MHz, the word size is 32-bits, the maximum burst transfer size is 512 bytes, and the memory controller's fixed access penalty (minimum memory access time, t_{MA}) at 50 MHz is another eight cycles for reads, or four cycles for writes. The parameter t_{MA} is subsequently assumed to be six cycles per access for simplicity, though reads are more frequent⁵.

The total redraw transaction rate (r_{RT}) is given by

$$r_{RT} = \frac{r_{RD}}{512} = \frac{36864000}{512} = 72000 \text{ Hz.}$$

If a transaction requires t_{RD} cycles to complete, the sum of the fixed and per-word transfer penalties, the total Wishbone bus cycles for 60 redraws, (T_{WB}) is given by

$$\begin{aligned} T_{WB} &= r_{RT} * t_{RD} \\ &= 72000 * (128 + 8) = 9792000 \end{aligned}$$

Wishbone bus congestion (c_{WB}), for the Wishbone bus frequency (f_{WB}) of 50 MHz, is given by

$$\begin{aligned} c_{WB} &= \frac{T_{WB}}{f_{WB}} \\ &= \frac{9792000}{5 \times 10^7} \\ &\approx 20\% \end{aligned}$$

This result, c_{WB} , is approximate because it does not have prefetch-overshoot and refresh cycles factored in. Average memory access latency is affected by (c_{WB}). There is a probability of 0.2

⁴Slice is the term Xilinx uses for its general purpose logic resources, and within the Spartan-3 a slice contains a four-input Look-Up Table (LUT) and two DFFs

⁵Using six as the average penalty results in worse statistics for the cache, so this is not an attempt to make the cache results appear more impressive. Also, the DMA controller can be used for data writes with ASCII text to pixel algorithm since it is well-behaved. This reduces memory bus congestion.

that a memory access will have to wait for a redraw prefetch transaction to complete. The average wait is $\frac{1}{2}t_{\text{RD}}$ cycles.

Average memory access latency (l_{WB}) in Wishbone bus cycles is given by

$$\begin{aligned} l_{\text{WB}} &= \frac{1}{2}t_{\text{RD}} \times 0.2 + t_{\text{MA}} \times 0.8 \\ &= 68 \times 0.2 + 6 \times 0.8 \\ &= 18.4 \end{aligned}$$

Additional parameters affecting the average memory access latency for the processor (l_{CPU}) are

- m_{CPU} : RISC16 has a CPU multiplier of two. This doubles the average memory latency seen by RISC16.
- i_{mem} : RISC16 has a memory instruction latency cost of four additional cycles, due to the interlocking and flushing mechanism used.
- t_{sync} : An additional synchronisation cost of four processor clock cycles because the processor and memory operate at different frequencies.

Giving l_{CPU} (in processor cycles)

$$\begin{aligned} l_{\text{CPU}} &= l_{\text{WB}} \times m_{\text{CPU}} + i_{\text{mem}} + t_{\text{sync}} \\ &= 18.4 \times 2 + 4 + 4 \\ &\approx 45 \end{aligned}$$

Profiling the text-mode code shows that about 900,00 memory accesses (see Figure 5.6) which shows the total accesses for 50 such conversions) are required for the conversion of a screen full of ASCII text data to pixel data, (640x400 resolution, 16-bit colour).

If the desired display update rate, for the ASCII to pixel transformation, is 20 Hz, for smooth scrolling, and the cost for each of the memory accesses required is 45 CPU cycles, then 720 million CPU clock cycles will be required per second. The RISC16 processor will not be fast enough, it only operates at 100 MHz. If a fast cache is used, which reduces average memory access time to about 5.1 cycles, then the 100 MHz processor will be about fast enough.

5.2.2 Implementation Considerations

The Spartan-3 FPGA architecture has 2 kB BRAMs, which are built-in primitives, and are suitable for implementing small caches. Each BRAM has two independent ports and supports several

different data widths. Due to the locations of the BRAMs within the Spartan-3, aligned in two vertical columns on either side of the FPGA, half in each column, two or more BRAMs can be connected in parallel to give greater capacity, data transfer rates, and widths.

To design an efficient cache, the performance effects of set associativity, write policy, capacity, data widths, line size, and replacement policy were examined. The processor within OpenVGA has the primary purpose of converting data into a format suitable for displaying, and for this particular application, memory access consists of reading blocks of data, applying a transformation, and writing the output to the framebuffer.

This sequential, block-reading/writing behaviour was speculated to heavily benefit from a cache, since there is a high probability the next datum to be read/written will be from one of the blocks currently cached. Furthermore, such sequential behaviour should mean that small cache sizes will not reduce performance much.

Factors responsible for performance when designing caches are:

- Reducing the miss penalty.
- Reducing the miss rate.
- Reducing hit calculation time.

A cache that excels at all of these may well be too large to fit within the 200k-gate Spartan-3 FPGA used. To design a cache that gives good performance without being too large, some objective measure of cache performance is needed.

Since designing and implementing a fast cache can also be very time-consuming, a parameterisable cache simulator was written in C++ (see Appendix ??). This allows cache design parameters to be chosen before the cache is constructed. Simply modifying the simulator parameters and (re)running a test allows a different cache design to be evaluated. By methodically changing each of the parameters, a cache has been chosen that will perform well for the simulated task.

If an unrealistic simulation model is used for evaluating cache performance, due to poor understanding of the problem, then the chosen parameters may not be suitable for final cache design. The simulation model used for evaluating cache performance was the ASCII text to pixel conversion algorithm, this was implemented in C++ too (see Appendix A. The C++ implementation of this algorithm closely matches the desired characteristics of the assembly routines that need to be written for the RISC16 or TTA16 processors.

5.2.3 Set Associativity

The set-associativity cache parameter represents how many cache look-up indices are hashed from an incoming address. These indices are used to retrieve cache-line tags for comparison with the

```
Memory Statistics:
Reads: 19280100
Writes: 25684152
Total: 44964252

Total CPU memory access cycles: 2113319844 (Av: 47.0/op)
```

Figure 5.6: Text-mode simulator output without using a cache.

```
Cache Properties:
Cache Size:          2048
Tags per Bank:       16
Set associativity:    2-way
Line (Block) Size:   64
Write Policy:        thru

Memory Statistics:
Reads: 7812096
Writes: 28882131
Total: 36694227

Cache Statistics:
Accesses: 67364279
Hits: 67120151 (99.6%)
Fast Hits: 61714821 (91.9%)
Misses: 244128 (evicts = 0)
Miss rate(per 1000): 3.6

Total CPU memory access cycles: 358827429 (Av: 5.3/op)
```

Figure 5.7: Text-mode simulator output when using a cache.

incoming address. A cache with higher set-associativity has a higher hit-rate, but also has larger and more complicated cache sense logic.

A cache look-up consists of:

- Hashing the incoming address to generate an index into the tag memory.
- Retrieving the tag from the tag memory.
- Comparing the tag-field of the incoming address to the retrieved tag.
- Signalling a hit and retrieving the data from cache upon a match, or else signalling that a miss has occurred.

With a direct-mapped cache, the incoming address can be hashed to only one cache memory location, so there is only one tag to check. If the incoming address is hashed to produce two different indices into the cache memory, which data (if any) chosen depends on if either retrieved tag matches. Such a cache design is said to be a two-way set associative cache. With two tag banks, these two tag comparisons typically happen in parallel[19].

A fully associative cache requires that every tag be checked for a match. This requires a lot of parallel comparison logic and tag memory bandwidth, and is quite rare[19]. Caches can be designed to be direct-mapped, or fully associative, or anywhere in between these two extremes.

Fully associative caches typically have the highest hit rates[31], if all other parameters held constant, since data can be stored anywhere within the cache and the cache is free to choose any line to evict. Direct-mapped caches have the lowest hit performance, but also the lowest implementation cost (complexity, logic, and latency)[19, 31].

A two-way set associative cache was chosen since this is typically a good trade-off for a small caches. It requires only one extra n -bit comparator and a few extra logic gates, compared to a direct-mapped cache, and results in lower miss-rates[31]. A four-way set-associative cache would require two more comparators than a two-way cache, and extra tag memories, and have a higher latency on a Spartan-3 FPGA⁶.

5.2.4 Cache Size

The results from simulating the text-mode redraw algorithm (see Table 5.1) show that performance is extremely good with even relatively small cache sizes because the total size of the read-data is quite small, 2 kB of font data, 4 kB of text-mode frame buffer data, and some stack reads.

Based off these results, 2 kB, the size of a Spartan-3 BRAM, appears to be a good choice. There is little advantage, FPGA resource usage-wise, for going smaller. With sizes greater than

⁶The extra tag memories would be needed because the Spartan-3 FPGA family has distributed RAMs with 16 entries. To store the 32 tag entries, as determined later in this section, would require a exactly two banks of distributed RAM for either a direct-mapped or two-way set associative cache, but moving to a four-way set associative cache would require two more banks of distributed memory, if simply to provide the required number of memory read ports.

Cache Size	Lines (total)	Misses (total)	Miss Rate (per 1000)	Average Latency
1 kB	16	485065	25.2	5.9
2 kB	32	14255	0.7	5.1
4 kB	64	2633	0.1	5.1
8 kB	128	97	0.0	5.1
16 kB	256	97	0.0	5.1

Table 5.1: Performance effects of different cache sizes, with a two-way set associative, 64 byte line size, write-through cache.

2 kB there is no significant average latency improvement, and the extra complexity and resource usage of larger caches is of minimal benefit in this application.

5.2.5 Cache Line Size

Shorter lines for a given size of cache require more entries within the cache lookup table. Longer cache-lines mean that more data has to be fetched from main memory, but the hit rate for subsequent sequential accesses will be higher.

Line Size (bytes)	Lines (total)	Miss Rate (per 1000)	Average Latency
8	256	1.9	5.2
16	128	1.2	5.2
32	64	0.8	5.1
64	32	0.7	5.1
128	16	1.1	5.2
256	8	3.4	5.4

Table 5.2: Performance effects of different cache line sizes, with a two-way set associative, 2 kB, write-through cache.

The 2 kB data cache demonstrates excellent performance with all tested line-sizes (see Table 5.2), but the number of tags to store doubles as the line-size halves, for a given cache size. Since a single Spartan-3 distributed RAM has 16 entries, and there are two tag banks within a two-way set associative cache, this means that a line size of 64 bytes is the most efficient hardware implementation, as well as having a slight performance edge. A line size of 64 bytes was chosen for the final cache.

5.2.6 Replacement Policy

To determine which cache line to evict on a data miss, three cache data replacement policies were evaluated:

- Least Recently Used (LRU)
- First In, First Out (FIFO, also called round-robin)
- Random

Results from the cache simulator (see Figure 5.4) agreed with published results for simulations done with an Alpha 21264 processor[19] (see Table 5.3).

Size	LRU	Random	FIFO
16 kB	114.1	117.3	115.5
64 kB	103.4	104.3	103.9
256 kB	92.2	92.1	92.5

Table 5.3: Different cache eviction policies, with a two-way set associative cache, and their effect upon cache miss rate, per 1000 instructions, with a DEC Alpha 21264 processor running the SPEC2000 benchmarks[19].

Replacement Policy	Miss Rate (per 1000)	Average Latency
Least Recently Used	0.8	5.1
Random	0.7	5.1
First-In, First-Out	0.8	5.1

Table 5.4: Different cache replacement policies and the effect on miss-rates and average latencies when simulated using the text-mode redraw algorithm and a 2 kB, two-way set associative, 64 byte line size, write-through, data cache.

Random replacement was therefore chosen since it is the easiest to implement. The random replacement selector is a single bit of a seven bit LFSR, as LFSRs have been shown to generate good approximations of random sequences[12].

5.2.7 Write Policy

The two approaches with writes to the cache are to attempt to write it straight away (the write through policy), or store incoming data in the cache so that it can be written to the memory in

bursts. The second approach, the write back policy, was less efficient in this case (see Table 5.5), and is more complex to implement, so the data cache uses write through instead.

Since the majority of data written with the cache simulator, and simulating the text to pixel conversion algorithm, uses writes to the DMA controller, not direct memory writes, the performance of each approach will not have a significant effect in this case anyway.

Line Size (bytes)	Write Back		Write Through	
	Misses/1000	Cycles/Op	Misses/1000	Cycles/Op
8	75.0	8.7	1.9	5.2
16	37.9	7.1	1.2	5.2
32	19.3	6.3	0.8	5.1
64	10.9	6.0	0.7	5.1
128	9.5	6.2	1.1	5.2
256	20.2	8.6	3.4	5.4

Table 5.5: Different cache write policies and line-sizes, with a 2 kB, two-way set associative cache, and their effect upon average memory latency, and cache miss rate per 1000 instructions.

5.2.8 Pipeline Length

The data cache interface to the processor needs to run at the processor frequency to minimise latency but, since hit calculation (performed by the sense logic) requires cache tag lookup and then address comparison, there is a large combinatorial delay, so this logic was pipelined.

The previous two tags used for a cache look-up operation are stored so if the same cache tag index is subsequently used, the look-up operation will take one less clock cycle. These are called “Fast Hits” in the cache simulator output, and this was about 90% of all hits with the cache simulations, so the average latency saving from this optimisation is about 0.9 cycles per memory access. The previous-tag registers have an additional purpose as well. Since tag retrieval takes several nanoseconds, they are used as pipeline registers so the cache can operate at a higher clock frequency.

Upon hit, the total latency is zero cycles for a fast hit and one cycle for a normal cache hit. The miss latency is determined by the state of the SDRAM as well as the overhead imposed by the cache, plus the signal synchronisation, in each direction, since the memory access crosses from the processor domain to the Wishbone bus domain and back. Minimum total latency for a read is 15 cycles plus an additional one quarter of the line size (and the processor design adds another five), but average latency for a miss will be a lot higher since the video redraw logic is competing for SDRAM access too.

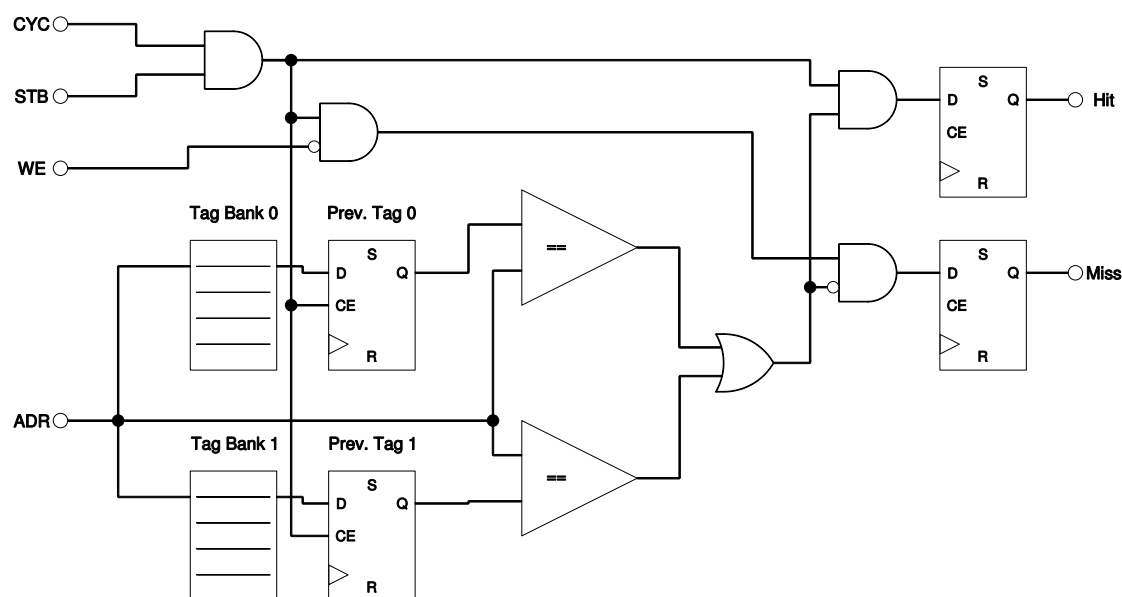


Figure 5.8: Simplified schematic showing the cache sense logic.

Figure 5.8 is a simplification the data cache sense logic. The Wishbone signals used to initiate a cache look-up are shown on the left (CYC, STB, WE, and ADR). The Miss signal is used to initiate a fetch from memory and is only asserted by cache-read misses, never for cache-writes. The Hit signal is used to complete a cache look-up transaction.

5.2.9 Testing

The first part of testing the cache was building a software cache simulator. This allowed evaluating changes to the design can be made quickly and the performance implications simulated and evaluated. A text-mode conversion algorithm was written in as well. This was used with the cache simulator to provide a realistic test scenario, and the CRT simulator was used to verify correctness.

To hardware-test the synthesised cache, software writing to OpenVGA using PCI-to-Wishbone bridge logic core provided a workload for the cache. All PCI bridge transactions had to pass through the cache to access the SDRAM, and the same memory-test applications that were written for the SDRAM controller (see Section 5.5) could be used.

5.3 DMA Controller

Module:	wb_dma	
Description:	A DMA controller for writing data in bursts to a memory controller.	
Related Files:	/rtl/lib/wb_dma.v	
Testing Files:	/sim/lib/wb_dma_tb.v	
Author:	Patrick Suggate	License: GPL

The DMA allows a peripheral to queue up write data so that it can be written to the memory controller as a burst transfer, improving total system memory throughput. The DMA consists of a FIFO, using a 2 kB BRAM, to buffer data until it is ready to be written. If the DMA has been filled with 2 kB of data, it can optionally be re-written multiple times to any location in memory. This is useful for clearing the OpenVGA display memory for example. Information on how to access and program the DMA controller is contained in Section D.3.3 of the TTA16 programming guide.

5.4 Serial PROM

Module:	wb_sprom	
Description:	A Wishbone-compatible, serial PROM read-back module designed to work with Xilinx serial PROMs.	
Related Files:	/rtl/lib/wb_sprom.v	
Testing Files:	/sim/lib/wb_sprom_tb.v	
Author:	Patrick Suggate	License: GPL

Since the FPGA state is stored internally, using SRAM, it needs to be reloaded every time the device is powered on. Xilinx sell a Serial Flash PROM to accomplish this task. After the initial FPGA state is loaded from this PROM, the user can then access the PROM and read out additional data.

When building the binary image file for the PROM contents, extra data can be appended to the end of the PROM image file. This requires some padding bytes and an ID pattern preceding the appended data. This can then be read back once the FPGA performs its own initialisation.

The serial PROM is needed to store the VGABIOS ROM image and additional processor instructions, which is transferred to SDRAM upon start-up. The VGABIOS ROM image is 32 kB of 16-bit, x86 code which resides at address 0x0C8000 in the OpenVGA host computer's memory address space[39]. The purpose of this ROM is too implement standard routines for communicating with the VGA hardware.

5.4.1 Serial PROM Read-Back Considerations

Since this is typically done only on power-on, since sequential access without direct control of the address counter is very slow, the PROM read logic should be small so to minimise logic usage for hardware that is used just once.

The simplest way to read data from this serial PROM is to clock the data out into a shift register, searching first for the ID pattern to locate the beginning of user data, then continue clocking while transferring the user data onto the Wishbone bus and to the SDRAM.

Chapter 6

I/O Logic Cores and Data Synchronisation

This chapter describes the logic cores used to communicate with the PCI Local Bus, the video DAC, the DVI TMDS encoder, the USB UART, and the LEDs. Also contained is the clocking and data synchronising issues that had to be solved so that the previously mentioned logic cores can exchange data using the OpenVGA Wishbone bus.

6.1 Clock Architecture

All OpenVGA clock signals are derived from either the 33 MHz PCI clock, or the 50 MHz on-board crystal oscillator (see Figure 3.1). The display dot-clock, while derived from the 50 MHz oscillator, has to be treated as a separate clock domain since its frequency can be set to either 25 MHz or 40 MHz, depending on the video mode.

6.1.1 Spartan-3 Clocking Resources

The Spartan-3 contains low-skew clock lines which can be connected to all the logic resources within the FPGA. To drive these lines there are a couple of logic primitives, BUFGMUX and DCM, that can be used.

Global Clocks

Module:	BUFGMUX, BUFG	
Description:	Simulation modules which emulate the functionality of the Xilinx Spartan-3 primitives with the same name.	
Related Files:	/sim/xilinx/BUFGMUX.v, /sim/xilinx/BUFG.v	
Testing Files:	/sim/xilinx/BUFGMUX_tb.v, /sim/xilinx/BUFG_tb.v	
Author:	Patrick Suggate	License: GPL

The 200k-gate Xilinx Spartan-3 FPGA has eight dedicated, low-skew, global clock lines [45]. These are routed throughout the FPGA and designed primarily for use as clock and reset signals. Skew on the clock or reset signals can cause metastability so it is essential that skew on these lines is low, or at least known. XST automatically calculates the skew of a synthesised design, and generates a timing report for the clock signals. These reports also include warnings when a design uses clock signals in a way that may cause unpredictable skew.

Clock signals can be explicitly routed into global clock lines using the Xilinx **BUFGMUX** primitive, though this is normally performed automatically by XST. A **BUFGMUX** is a two input multiplexer and the output is a global clock line. This primitive allows one of two different clock signal inputs to be selected as the source. This primitive contains logic to avoid clock glitches so the input source clock can be dynamically switched.

Digital Clock Managers

Module:	DCM	
Description:	Simulation modules which emulate the functionality of the Xilinx Spartan-3 primitives with the same name.	
Related Files:	/sim/xilinx/DCM.v	
Testing Files:	/sim/xilinx/DCM_tb.v	
Author:	Patrick Suggate	License: GPL

The Spartan-3 family of FPGAs typically contain four DCM (Digital Clock Manager) primitives, except for the smallest device which only has two [45]. DCMs use DLLs (Delay-Locked Loops) to condition clock input signals, and optionally phase-shift and derive new clocks as well. Conditioning the input clock signals ensures a clean clock output, corrects the duty-cycle (optional), and allows XST to more accurately assess and control clock-skew. Additionally, the input clock period can be multiplied and divided to generate new clocks, and this feature was used to generate processor, SDRAM, and VGA clocks from the on-board 50 MHz oscillator.




Figure 6.1: BUFGMUX and DCM example showing the OpenVGA pixel clock generation.

6.1.2 Clock Domains

The operating frequency of the PCI Local Bus is nominally 33 MHz [36], but this is not certain. There are PCI power saving modes in which the bus frequency can be lowered, possibly even to zero¹, for example. The video DAC and TMDS encoder ICs operate at the dot-clock frequency, 25 MHz for the 640x480 resolution VGA. Since these clocks differ, and also differ from the 50 MHz Wishbone bus clock, data crossing these separate domains need to be synchronised.

The Spartan-3 FPGA has clocking resources which support clock generation and conditioning, and multiple clock signals [1]. There is also logic primitives suitable for constructing asynchronous FIFOs, a logic core used to transfer data between clock domains.

Clock Name	Clock Frequency (MHz)	Domain Name
Wishbone Clock	50	Wishbone
CPU Clock (RISC16)	100	
CPU Clock (TTA16)	150	
SDRAM Controller Clock	50	
SDRAM Datapath Clock	100	
Video Clock 0 (640x480)	25	Dot Clock
Video Clock 1 (800x600)	40	
PCI Local Bus Clock	33	PCI

Table 6.1: OpenVGA's clock domains and frequencies.

¹In addition, off-the-shelf motherboard's PCI system clock can differ significantly from the specification as the PCI clock is often derived from the Front Side Bus (FSB) clock, which can be adjusted by the user, or differ due to manufacturing tolerances.[36]

OpenVGA has multiple, asynchronous clock domains (see Section 6.1.2, and data needs to be transferred between clock domains. To prevent problems due to metastability², data crossing asynchronous clock domains needs to be synchronised. A single-bit synchroniser is shown in Figure 6.2, and for signals many bits wide, asynchronous FIFOs (see Section 6.1.3) are used for synchronisation.

Synchronous Clock Domains

Module:	wb_sync	
Description:	Synchronises Wishbone transactions between two Wishbone buses with synchronous, but different frequency, clocks.	
Related Files:	/rtl/lib/wb_sync.v	
Testing Files:	/sim/lib/wb_sync_tb.v	
Author:	Patrick Suggate	License: GPL

Data is transported between separate OpenVGA functional units, internal to the Spartan-3, via a Wishbone compatible bus (see Appendix C). The Wishbone bus clock is 50 MHz and is the DCM conditioned output generated by the 50 MHz on-board oscillator.

The CPU, SDRAM datapath, and the pixel clock, all share a common root clock, the 50 MHz Wishbone bus clock. Though different frequencies are used for each, the CPU (see Chapter 4) and SDRAM controller (see Section 5.1) clocks are synchronous with the Wishbone bus clock, and with frequencies that are integer multiples of the Wishbone bus clock frequency (see Table 6.1.2).

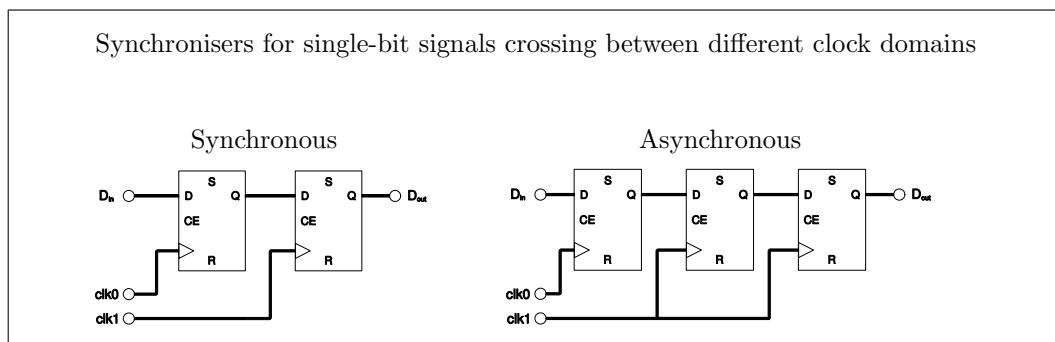


Figure 6.2: A synchronisers for crossing clock domains.

Transitions between these synchronous clock domains requires two DFFs (see Figure 6.2) for each signal that crosses a domain boundary. Additionally, signals multiple bits wide retain the same relative phase when crossing synchronous domains using simple synchronisers, which is not the case for crossing asynchronous domains using simple synchronisers[10].

²A full discussion of metastability is beyond the scope of this work but a D flip-flop can enter a metastable state if setup and/or hold times are violated. Its data output can become essentially random.

There are three additional problems when crossing synchronous clock domains:

1. A signal crossing from the higher-frequency domain, to the slower domain, needs to be asserted for at least one cycle of the slower clock so it is registered correctly.
2. When a signal crosses from the slower clock domain, to the higher frequency clock domain, the higher frequency domain needs to correctly handle that the signal might remain asserted for more than one cycle of the faster clock, but should only trigger once.
3. At the boundary of two synchronous signals, the setup and hold times must be met for both D flip-flops, which means that the smallest period of the two timing constraints has to be applied to both DFFs.

To solve the first two of these issues, an acknowledge (ACK) signal was used when crossing synchronous clock domains. This indicates that a signal has been correctly registered. The signal source then de-asserts the signal at the end of the cycle which the ACK was received. For the final problem, XST automatically applies the correct constraint but the designer still needs to be aware that the tightest of the two constraints will apply to any combinatorial logic.

Asynchronous Clock Domains

To cross asynchronous clock domains there are two common methods: a simple two flip-flop synchroniser (see Figure 6.2), which is useful for single bit flags like an interrupt signal; and asynchronous FIFOs, which are useful for multiple bit-width signals, like buses, where a phase relationship has to be maintained for all signals, across the entire width [10].

6.1.3 Asynchronous FIFOs

Module:	afifo16, afifo2k	
Description:	Parameterisable, asynchronous FIFOs, with either 16 or 2048 entries, which synchronise and queue data which has to cross asynchronous clock domains.	
Related Files:	/rtl/lib/fifo/afifo16.v, /rtl/lib/fifo/afifo2k.v, /rtl/lib/counter/bin2gray.v	
Testing Files:	/sim/lib/wb_sync_tb.v	
Author:	Patrick Suggate	License: GPL

An asynchronous FIFO takes write data from one clock domain and stores it internally within a queue, implemented using SRAM, and the data can then be read out to the other clock domain. The FIFO generates control signals to indicate its state (empty, full, etc.) so that data is not

written to a full FIFO, or requested from an empty FIFO. To avoid metastability, even these control signals are valid in just one domain. Typically the empty signal is valid in the FIFO's read clock domain, and the FIFO's full signal is typically valid within the FIFO's write clock domain.

The penalty due to having to use an asynchronous FIFO is both a latency penalty and extra logic usage (approximately 60 Spartan-3 logic slices per 32-bit wide FIFO, see Section 6.1.3). Asynchronous FIFOs are only used for the PCI and the video redraw circuits since each are within separate asynchronous clock domains.

There are two logical divisions of the design for an asynchronous FIFO: the logic for adding data to the FIFO, and removing data from the FIFO; and the logic for comparing the read and write pointer values across clock domains to allow the setting of FIFO state flags.

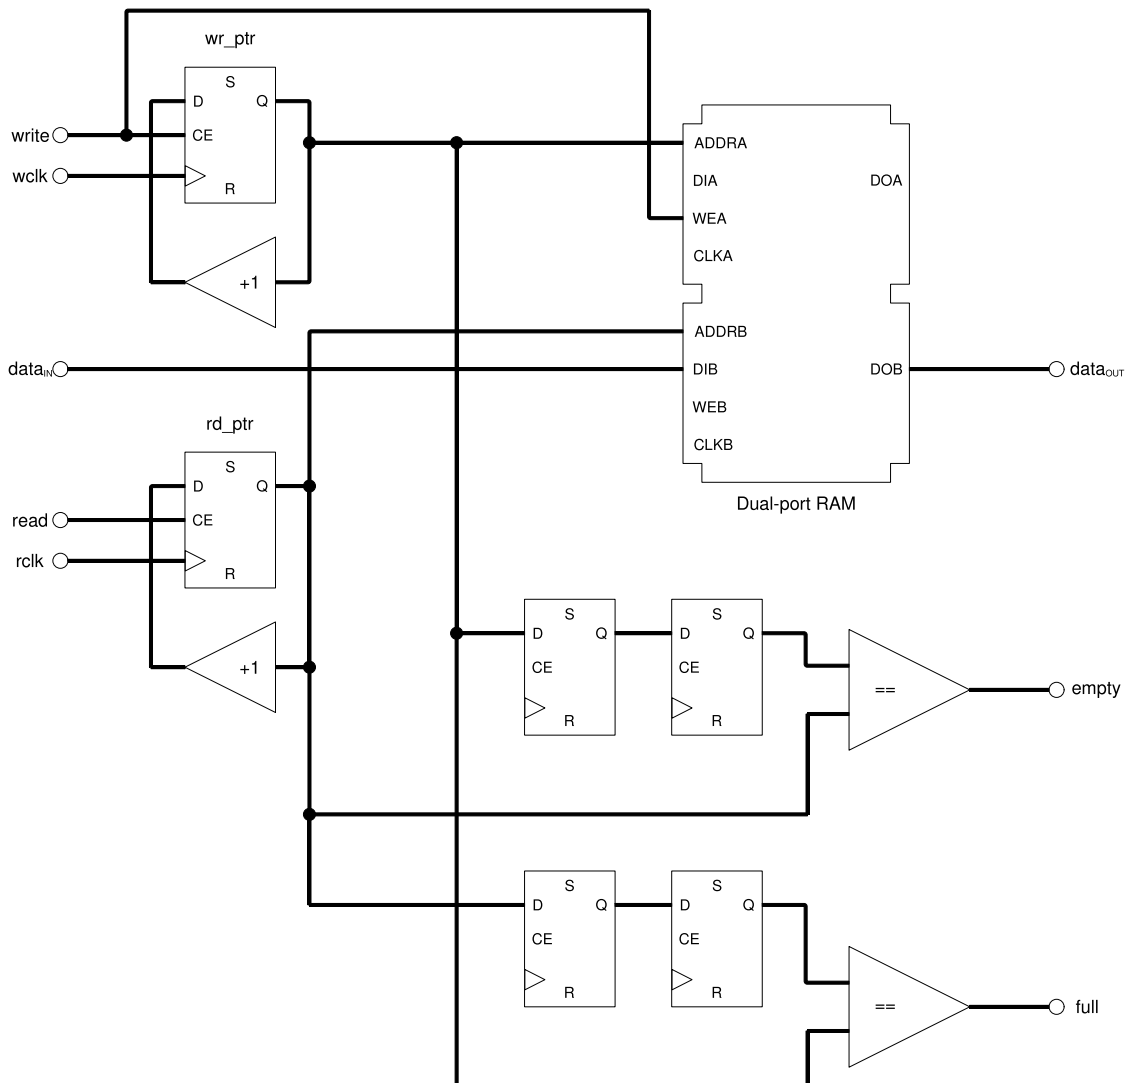


Figure 6.3: Simplified schematic showing the basic components of an asynchronous FIFO.

Decimal	Binary	Gray Code
0	0000	0000
1	0001	0001
2	0010	0011
3	0011	0010
4	0100	0110
5	0101	0111
6	0110	0101
7	0111	0100
8	1000	1100
9	1001	1101
10	1010	1111
11	1011	1110
12	1100	1010
13	1101	1011
14	1110	1001
15	1111	1000

Table 6.2: Gray codes with the standard binary and decimal equivalents.

The logic for adding (and removing) data is straightforward, simply increment the write pointer every time data is added, and increment the read pointer whenever data is removed. Figure 6.3 is a simplified diagram of a FIFO, the Gray-encoders are not shown, and neither is a faster `full` and `empty` assertion path. These are covered in detail in [5].

Comparing these pointers across clock domains, as is needed to set the full and empty flags, is more difficult due to the possibility of metastability. Standard base-2, ripple-carry incrementers can have multiple, even all, bits change during an increment operation. This could potentially cause every bit to become non-deterministic within the post-synchronised, pointer representation.

The first step for the control signal generation is to convert the pointers into a unit-distance representation³, like Gray encoding, and this was used here (see Table 6.2. Even with metastability the desired pointer value has a maximum difference of one from any generated post-synchronised pointer. The only two values that the post-synchronised pointer can have is the previous value or the desired value, and this will become the correct value the following clock cycle anyway [10, 5].

³A unit distance encoding means that only one bit changes between successive values in a sequence, but arithmetic in this encoding is more complex which is why a standard representation is used for the pointers.

Control Signals

Either the **full** or **empty** signals are asserted when the two Gray-encoded counters match. Depending whether the match was caused by the FIFO going full, or going empty, determines which of the two flags is asserted.

The count order of the two Most Significant Bits (MSBs) of a Gray code counter are identical to the two MSBs of the standard binary encoding. By sampling the two MSBs of the Gray codes we can determine which quadrant the value is in. Using this quadrant detection, and by comparing the current pointer values with the previous pointer values, it can be determined whether the FIFO was going full, or going empty, immediately before the pointers matched.

Half-Full Detection

In some applications, as with the video redraw circuit, it is useful to prefetch multiple words, since this makes more efficient use of the Wishbone bus. Before a prefetch is issued the prefetch control logic needs to know whether there is space in the FIFO for the entire data transfer. An extra FIFO control signal, calculated by comparing pointer quadrants, was added for this purpose.

By comparing the two MSBs of the FIFO read and write pointers, since these quadrant bits have the same value as the standard binary encoding, the FIFO can calculate the difference in quadrants between the two pointers. If the quadrant bits are equal then both pointers are within the same quadrant so the FIFO is less than 25%, or more than 75% full. If there are more than two quadrants between the write and the read pointer, then the FIFO can also more than three quarters full. The result is a control signal that detects whether the FIFO is between 25% and 75% full (or not), and has been labelled **halfish**.

Performance and Size

The size of a synthesised asynchronous FIFO, with a data width of 32-bits, and with the target a Xilinx Spartan-3 FPGA was 59 slices⁴. With high optimisation settings, the maximum speed at which the FIFO clocks is 144 MHz, with the data write clock domain being the slowest (clock period of 6.9 ns). The read clock domain achieved a clock period 3.8 ns, but since the logic for both domains is very similar, this difference was speculated to be due to how XST optimised the design.

⁴The 200k gate Spartan-3 device that was used for OpenVGA contains 1920 logic slices

6.2 PCI-to-Wishbone Bridge

Module:	pci_mem_top	
Description:	A <i>Plug and Play</i> compliant PCI-to-Wishbone-bridge logic core.	
Related Files:	/rtl/pci/pci_mem_top.v, /rtl/pci/pci_mem.v, /rtl/pci/cfgspace.v	
Testing Files:	/rtl/pci/pci_mem_top_tb.v	
Author:	Patrick Suggate	License: GPL

OpenVGA connects to the host computer using the Peripheral Component Interconnect (PCI) Local Bus. All data exchanged between OpenVGA and the host is through the PCI-to-Wishbone bridge logic core. PCI was developed by Intel in 1991 as a replacement for the aging Industry Standard Architecture (ISA) Bus [36, 33]. PCI features data widths of 32 or 64 bits and has nominal operating frequencies of 33 or 66 MHz (but can be reduced by power saving modes).




Figure 6.4: PCI-to-Wishbone bridge block diagram.

6.2.1 PCI-to-Wishbone Bridge Logic Core

This logic core maps OpenVGA's local memory into the host-system's address space, and this bridge allows data to pass between the host and OpenVGA. Figure 6.4 represents a simplified diagram of this logic core. The OpenVGA PCI bus interface is in a separate clock domain to the system Wishbone bus. PCI typically operates at 33 MHz, with the clock signal provided by the host system, and the Wishbone bus domain operates at 50 MHz. This clock is generated using an on-board crystal oscillator.

The PCI bridge logic core consists of two state machines, two asynchronous FIFOs, and some additional control and pipelining logic. Asynchronous FIFOs were used so that data can be transferred from one clock domain to another (see Figure 6.4). When synchronising signals across clock domains, these FIFOs typically add a latency penalty of two clock cycles. Domain crossing

therefore adds two cycles of latency for a PCI write and four cycles of latency for a PCI read request.

PCI-Domain State Machine

This state machine queues PCI read and write requests, and write data, in the PCI-to-Wishbone asynchronous FIFO. PCI read transactions require that data be fetched from the Wishbone bus domain and placed in the Wishbone-to-PCI FIFO. The contents of this FIFO can then be read from the PCI domain, with the data being placed on the PCI bus. The states of this state-machine, and the signals that cause transitions between these states, are shown in Figure 6.5.

To prevent the PCI bridge from being overwhelmed by many back-to-back PCI write transactions, the bridge will generate **TARGET ABORT** responses when the bridge is already busy. Multiple PCI read requests cause no problems though. This is because they must wait for data retrieved from the Wishbone domain before the transaction is completed.

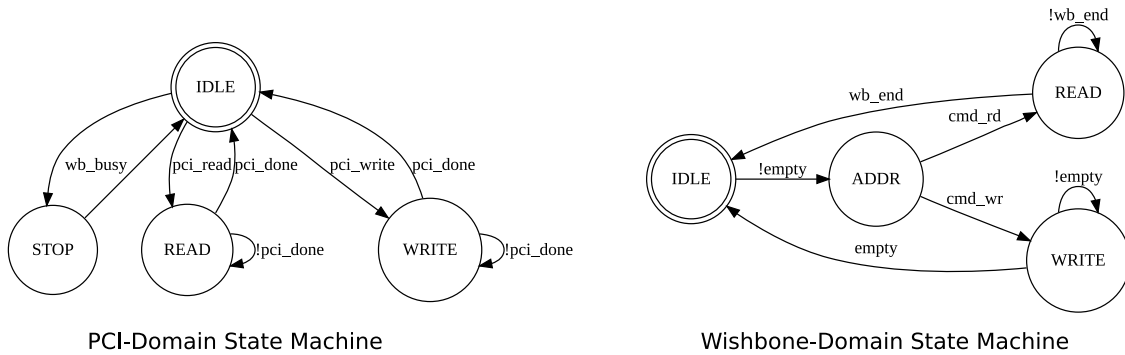


Figure 6.5: The PCI bridge state machines.

Wishbone-Domain State Machine

The PCI-to-Wishbone FIFO's **empty** signal de-asserting indicates that there is a PCI request pending. The Wishbone-side state machine decodes this request and issues a corresponding Wishbone transaction. Reads result in data being transferred to the PCI domain using the Wishbone-to-PCI asynchronous FIFO, writes use the PCI-to-Wishbone FIFO to transfer data to the Wishbone domain.

Summary of this state's machines signals:

- **empty**: Indicates the PCI-to-Wishbone FIFO is empty. When this signal de-asserts, there are pending PCI transactions to process, causing the state-machine to transition to the **ADDR** state. When this signal asserts in the **WRITE** state all data have been transferred, ending the write transaction.

- **cmd_rd**: When in the **ADDR** state, this signal causes a transition to the **READ** state, causing a Wishbone read transaction to be issued.
- **cmd_wr**: When in the **ADDR** state, this signal causes a transition to the **WRITE** state, causing a Wishbone write transaction to be issued.
- **wb_end**: Indicates the end of a Wishbone transaction, triggering a return to the **IDLE** state when in the **READ** state.

PCI Configuration Space

Upon system start-up, each PCI function and device has its configuration space registers read to determine its capabilities. The host (BIOS or operating system) can choose to write to configuration space registers to map the PCI device's registers and memories into the system's memory or I/O address spaces⁵. Figure 6.6 shows a configuration space **WRITE** access generated by a *Plug and Play* BIOS.

It is this configuration space information that informs a *Plug and Play* BIOS as to whether a device is a VGA-compatible video controller, a network interface card, or any other category of device listed in [33]. To be VGA-compatible, OpenVGA can be synthesised with VGA configuration space settings.

6.2.2 PCI Testing and Summary

The PCI specifications expect a device to respond within 25 clock cycles [33], and since four PCI cycles will be consumed for data to cross clock domains and resynchronise, this means that the OpenVGA has only 630 ns to complete the request and signal **TARGET READY**⁶.

⁵I/O space is a legacy feature from early IBM-compatible x86 architectures. The PCI specification encourages the use of memory address space instead[33, 36]

⁶There is a PCI command **TARGET ABORT** which is effectively “try again later”. This command is issued whenever there is still write data within the PCI to Wishbone asynchronous FIFO. The rest of the time all PCI accesses are routed to the SDRAM controller via the Wishbone bus and are expected to complete within this 630 ns.

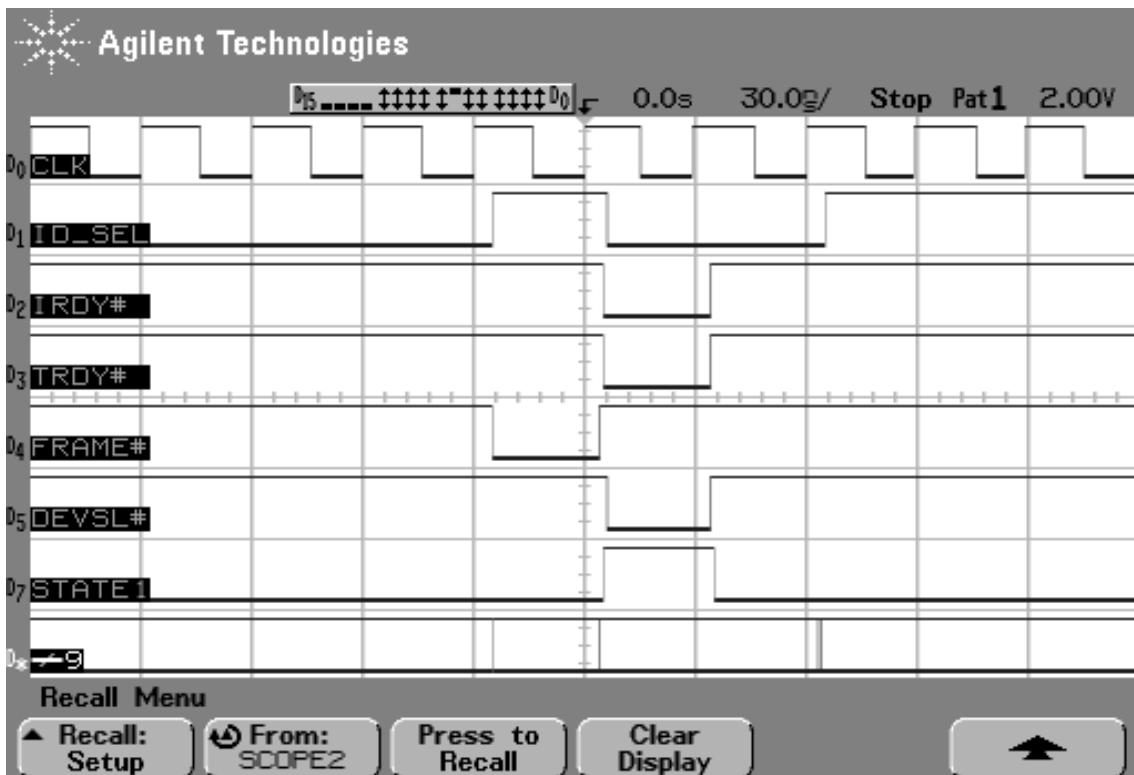


Figure 6.6: Screen capture of a PCI configuration space write transaction.

6.3 VGA and DVI Controller

Module:	wb_video_top
Description:	This logic core reads 16-bit pixel data from the framebuffer and converts it to 24-bit colour, outputting this to the video DAC and TMDS encoder, as well as generating all of the other signals necessary for redrawing the display.
Related Files:	/rtl/video/wb_video_top.v, /rtl/video/wb_vga_ctrl.v, /rtl/video/wb_redraw.v, /rtl/video/wb_crtc.v, /rtl/video/crtc.v, /rtl/video/vga16.v, /rtl/lib/fifo/afifo2k.v, /rtl/video/dvi_ctrl.v
Testing Files:	/sim/video/vga_tb.v, /sim/video/crtc_tb.v, /sim/xilinx/DCM.v, /sim/xilinx/BUFGMUX.v
Author:	Patrick Suggate
License:	GPL

OpenVGA outputs image data, made up of multiple picture elements (called “pixels”), to a standard computer display, typically a VGA connected CRT monitor or a VGA or DVI connected LCD monitor. The VGA standard is analogue, so a video DAC is used, the Philips TDA8777 IC – since the Spartan-3 has no analogue output pins, and DVI is a serial digital interface – encoded

using a TMDS transmitter IC, via the Texas Instruments TFP410.

VGA and DVI display formats are 2D (2-Dimensional) and multicolour, and both the VGA and DVI use similar control and synchronisation pulses⁷. The display data, in graphical modes, is laid out in memory as a flattened 2D array of pixel colours. To draw a complete screen of data, this pixel colour information is sent, one pixel at a time, to the display, using either a VGA or DVI connection. Display update rate is typically 60 Hz, meaning that the display is completely redrawn 60 times per second.

The block of memory that stores the pixels for a 2D display is called a frame-buffer. The size of this memory depends on the screen resolution and the colour depth (and can be from 1-bit to 32-bits per pixel).

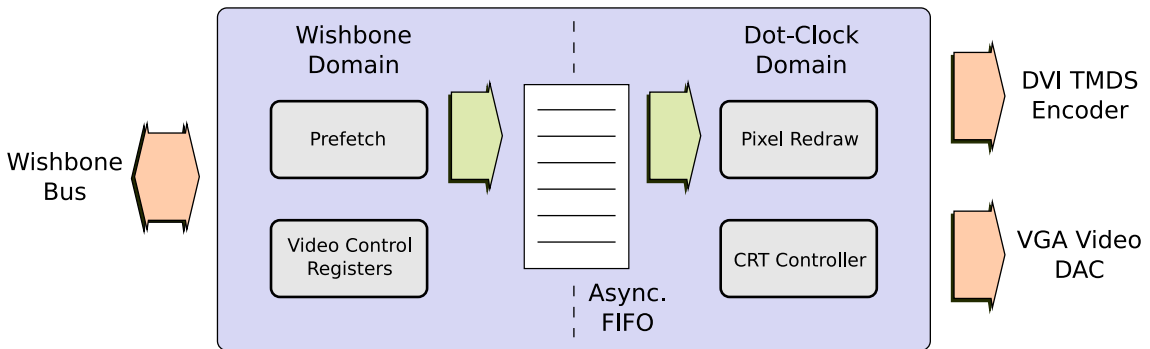


Figure 6.7: Video controller block diagram.

6.3.1 Display Modes

Text Mode

The original IBM VGA had hardware accelerated text modes that converted ASCII characters, and an additional attribute byte, into pixel data as the screen was being redrawn. This allowed a very small block of memory to represent the pixels of an entire screen of text data. Additionally, the ASCII to pixel conversion is quite slow on a general purpose CPU, and would have been prohibitively so in 1981 [14], which is why the original IBM Monochrome Display Adapter (MDA), and subsequent display adapters, used dedicated logic for the task.

OpenVGA does not use legacy dedicated hardware for this ASCII to pixel conversion as the on-board processor is capable of performing this task⁸. OpenVGA's firmware will contain the code necessary for the ASCII to pixel transformation. Appendix A contains the text-mode routines

⁷This was intentional to accelerate the adoption of the newer DVI specification[17].

⁸The original IBM XT ran at only 4.77 MHz, and required multiple clock cycles to execute a single instruction, and had quite a slow system bus[39]. TTA16 can operate at up to 190 MHz, executes an ALU operation every clock cycle, and has a fast connection to the framebuffer.

written in C. Assembly for both RISC16 and TTA16 was written, based upon the C code, and used to evaluate each processor.

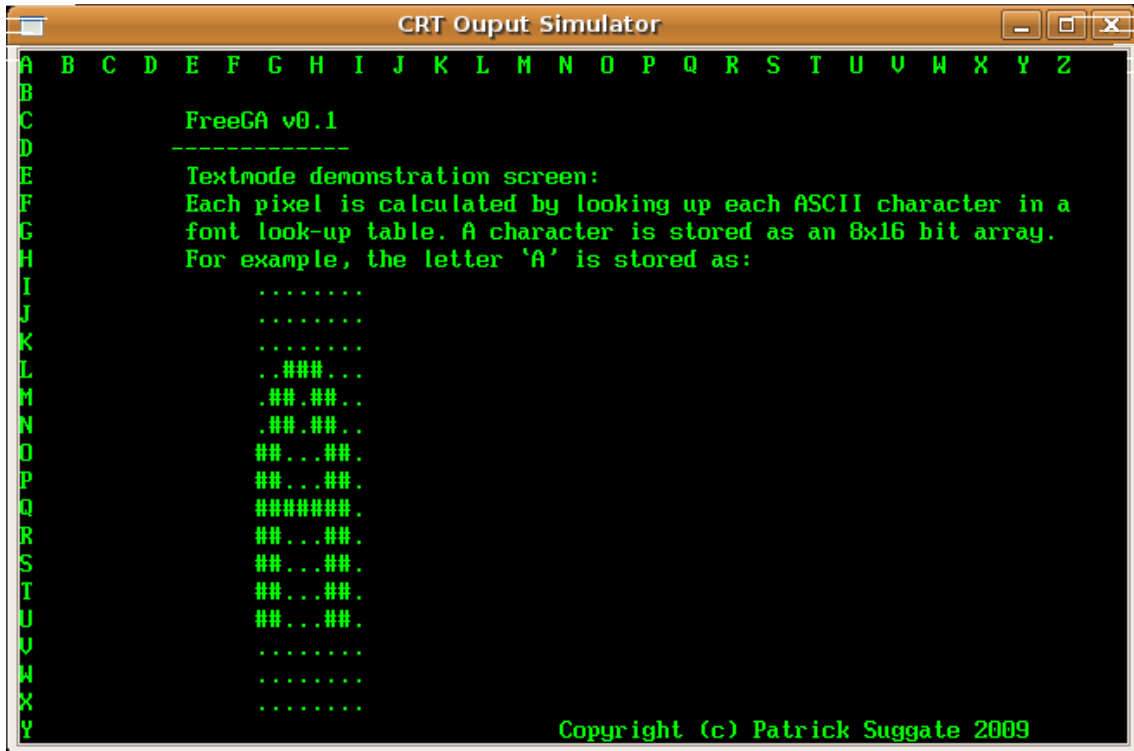


Figure 6.8: Emulated text-mode display data produced by a simulation.

To display the output of simulations of text mode, a Python script was written that takes in the generated pixel data and displays it in a window. The results produced by both the Icarus Verilog test-harnesses and the C code can then be verified. A screen capture of the generated text-mode output is shown in Figure 6.8.

Graphic Modes

The current version of OpenVGA has only two frequencies available as the clock source for the display timing, 25 MHz and 40 MHz. Table 6.3 shows common modes which use these dot clock frequencies.

6.3.2 Pixel Data Prefetch

The OpenVGA display data is stored in the SDRAM, which is in the Wishbone clock domain, and needs to be prefetched and queued within the dot-clock domain for redrawing the display. The data needs to be prefetched since the SDRAM is shared with other peripherals but a steady stream

	640x400 (VGA text)	640x480 (VGA)	800x600 (SVGA)	Unit
<i>General Timings</i>				
Dot Clock	25	25	40	MHz
H-sync	31.5	31.5	37.9	kHz
V-sync	70	60	60	Hz
Redraw data rate	36.9	34.6	57.6	MB/s
<i>Horizontal Timings</i>				
Front porch	16	16	40	clock cycles
Back porch	48	48	88	clock cycles
H-sync duration	96	96	128	clock cycles
H-total	800	800	1056	clock cycles
<i>Vertical Timings</i>				
Front porch	12	10	1	clock cycles
Back porch	35	33	23	clock cycles
V-sync period	2	2	4	clock cycles
V-total	449	525	628	clock cycles

Table 6.3: OpenVGA video modes which use a dot-clock of either 25 MHz or 40 MHz.

of data is needed to redraw the display without artifacts⁹

A 2 kB asynchronous FIFO achieves both these requirements, synchronises data which has to cross clock domains, and queues up a large quantity of pixel data so redraw can occur uninterrupted. The FIFO's half-full signal is used to trigger prefetches, which are always 256 byte burst read transfers from the SDRAM controller, as this is the SDRAM column size with the 8 MB SDRAM that was used.

Video Clocks

The signals to drive a display device, like a monitor, exist in their own clock domain as well. This is because different video modes typically use different pixel clocks to fetch and display pixel data, therefore has to be decoupled from the system's internal bus clock. For example, the display mode 640x480 @60 Hz uses a 25.175 MHz clock, 800x600 @60 Hz uses a 40 MHz clock, and 1024x768 @60 Hz uses a 65 MHz clock. To transfer data from the clock domain of the system's Wishbone bus to the video redraw circuit's clock domain, another asynchronous FIFO is needed, but only one since data is transferred in just one direction. This asynchronous FIFO was implemented using a

⁹The required data rate to redraw the display, at a resolution of 640x480, and 16-bit colour, is 50 MB/s. Pixel data has to be ready at each edge of the dot-clock when redrawing or else that pixel is skipped.

BRAM and also functions as a prefetch queue (see Section 6.3.2).

6.3.3 CRT Controller

The CRT Controller (Cathode Ray Tube Controller, called this for historic reasons) generates the horizontal and vertical refresh signals which determine the display resolution. To redraw the display, pixels are drawn from the top left to the bottom right of the display[14] (see Figure 2.2). The display is redrawn row-by-row with a horizontal synchronisation pulse marking the end of every row. Once all rows have been redrawn, a vertical synchronisation pulse marks the end of the screenful of data.

The basic design of the CRTC is two counters, a column counter and a row counter. The column counter increments at each positive edge of the dot-clock until the column limit is reached and then restarts at zero. Every time the column counter resets to zero, the row counter increments until it reaches its row limit, and then resets to zero.

6.3.4 VGA Output

OpenVGA's output to a VGA monitor are three analogue colour components (red, green, and blue), the horizontal synchronisation signal (subsequently referred to as hsync), and the vertical synchronisation signal (subsequently referred to as vsync).

The analogue signals are generated by an external video DAC since the Spartan-3 has no analogue outputs. The video DAC is connected to the Spartan-3 in 24-bit colour mode, and since OpenVGA's framebuffer is stored as 16-bit colour data, to reduce memory bandwidth, this data has its Least Significant Bits (LSBs) padded with zeros.

6.3.5 DVI Output

The DVI transmitter IC is connected to the Spartan-3 in 12-bit DDR mode, which means half the 24-bit colour is transferred on each clock edge. DVI support is incomplete and the DVI output functionality has been only partially simulated but not tested on OpenVGA hardware due to time constraints.

6.4 Miscellaneous Peripherals

6.4.1 LEDs

Module:	wb_leds	
Description:	MMIO module for driving on-board LEDs.	
Related Files:	/rtl/lib/wb_leds.v	
Testing Files:	/sim/lib/wb_leds_tb.v	
Author:	Patrick Suggate	License: GPL

OpenVGA has two LEDs for diagnostic and debugging output. These are memory-mapped to the address 0x480000 . The following listing is a TTA16 assembly routine from the OpenVGA firmware that sets the LEDs.

```
-----  
; set_leds - The two LSBs of 'r0' determines the LED outputs.  
set_leds:    {  
              ,  
              ,1    } ; 1  
              ; Set segment to the memory-mapped LEDs.  
              {  
              ,com    ->msr    ,0x48    } ; 2  
              {\r0    ->wad    ,\r0    ->mem    ,1    } ; 3  
              ; Restore SS and return.  
              {\r15    ->bra    ,com    ->msr    ,\r12    } ; 4  
              {  
              ,  
              ,  
              } ; 5
```

6.4.2 USB Serial Port

Module:	wb_serial_port	
Description:	A Wishbone interface to a standard UART.	
Related Files:	/rtl/lib/wb_serial_port.v, /rtl/lib/uart/*	
Testing Files:	/sim/lib/wb_serial_port_tb.v	
Author:	Jeung Joon Lee, Patrick Suggate	License: Modified BSD, GPL

A USB port was also included as part of the OpenVGA hardware. This is connected to the FPGA via a FTDI serial port IC. The Wishbone module allows the CPU to send and receive data via the serial port, with the goal of being able to dump OpenVGA's state, to aid with debugging OpenVGA.

The UART modules used as part of `wb_serial_port` is an open source module, with a very permissive license, and was obtained from OpenCores.org, a repository of open source HDL modules. The baud rate is configurable in the Verilog code and currently operates at 9600 baud, 8-bits data, one stop-bit, and no parity-bit.

Chapter 7

Summary, Conclusions, and Future Work

Concluding this thesis is a summary of OpenVGA, the useful logic cores that were developed, how these cores compared to those of other projects, and possibilities for future work following on from this project. Future work, even if completed by others, can be directly contributed back to the OpenVGA project. This is because OpenVGA is an open-source project and licensed under the GPL.

7.1 Summary of the OpenVGA Project

OpenVGA development has reached the point where it can function as a secondary graphics adapter. The OpenVGA local memory, an 8 MB SDRAM IC, is mapped into the host system's address space, and its contents are accessible via the PCI Local Bus. The contents of a portion of this memory, the framebuffer, can be displayed on a VGA monitor, either a CRT or LCD display, that is connected to OpenVGA.

The design of OpenVGA is also very modular, and many of its components interconnect using standard protocols. This applies to both the connections to and from the PCB, connections between components on the PCB, and connections between the logic cores within the Spartan-3 FPGA.

OpenVGA also contains a processor, implemented within the FPGA, with associated firmware for initialisation tasks. Additional functionality can easily be added by writing additional firmware. Firmware can be included during synthesis, stored in the on-board serial PROM, or uploaded through the PCI Local Bus. OpenVGA was designed so that this on-board processor can be used to emulate a subset of VGA functionality. This would allow it to function as a primary graphics adapter once complete, but this is future work.

A Linux kernel module has been developed that allows software to access OpenVGA's local memory. Data can be written to, and read from, OpenVGA using this kernel module. Any data written to the framebuffer region of memory will be displayed on a VGA monitor connected to OpenVGA.

Software has been produced for this project too. There are utilities for converting data between the different formats required, such as the tool to convert the files generated by the assembler into a format that can be included within the processor's Verilog source code. Other tools developed are an assembler for RISC16, a CRT simulator for developing the display-controller logic core, and a cache simulator that was used to analyse and select cache design parameters, and code to simulate text-mode.

The text-mode simulator was written to provide a realistic workload for the processor when choosing cache parameters. This piece of code would be a good starting point for developing firmware to support VGA's alphanumeric display-modes.

7.2 OpenVGA Logic Cores

The complete set of OpenVGA digital logic circuits requires only a small quantity of logic and it can be embedded into larger SoC projects with little penalty. The complete design uses less than 1100 slices of the Spartan-3 FPGA that it was synthesised for. The Spartan-3 used does not contain a lot of programmable logic compared to many other FPGAs, it is the second smallest device within the Spartan-3 product range [45]. OpenVGA used only 60% of the available FPGA logic. According to Xilinx, the XC3S200 is supposed to contain the equivalent of about 200,000 logic gates, and modern ASICs can contain hundreds of millions of logic gates.

7.2.1 Summary of Important OpenVGA Logic Cores

Many of the logic cores that have been developed for OpenVGA have been heavily tested and optimised. All of the logic cores are written in the industry standard Verilog HDL. Except for the asynchronous FIFOs, the external interfaces of these logic cores are compatible with the Wishbone interconnect standard. These factors mean that many of the following logic cores will likely be of use to other projects.

The logic cores presented in Table 7.1 are the sizes and speeds of the cores when synthesised alone, with high optimisation settings, optimised for speed, not area, and with parameters set to the configuration used within OpenVGA. These numbers are approximate and depend a lot on many other factors. Speeds will be lower when included in large designs, unless manually floor-planned. And changing the optimiser setting to area, not speed, should result in lower resource usage in more resource-limited designs.

Logic Core	Speed (MHz)	Size (LEs)	Other logic elements used
TTA16	190	200	2 BRAMs, 1 multiplier
RISC16	140	320	1 BRAM, 1 multiplier
Data Cache	150	80	1 BRAM
PCI to Wishbone Bridge	110	320	
SDRAM Controller	120	100	
Display Controller	140	200	
Async. FIFO, 32-bit, 16-entry	>150	60	

Table 7.1: Summary of significant OpenVGA logic cores.

TTA16: A 16-bit TTA Processor

A TTA processor logic core with a 32-bit instruction width, 16-bit data width, 4 kB of local instruction SRAM, three pipeline stages, and four data transports was developed for OpenVGA. When synthesised for the Spartan-3, operating frequency was up to 150 MHz, and using about 200 logic slices.

Unusual amongst processors, TTA16 does not support exceptions due to the extra complexity this would have added to the design. This processor also features interleaved instructions and no data-hazard interlocks, so it is difficult to write assembly code for (see the programming guide in Appendix D). The result is a processor that compares favourably with other FPGA-based processors (see Table 7.2).

RISC16: A 16-bit RISC Processor

This is a RISC processor with a 16-bit instruction width, 16-bit data width, 2 kB of local instruction SRAM, five pipeline stages, and data forwarding. It was developed for OpenVGA because TTA16 was such a difficult processor to write code for and the result is a processor that is far easier to program (a programming guide for RISC16 has been included as Appendix E). Consequently, it shares many of the same functional units and design techniques as TTA16, and therefore has value for comparison with TTA16.

This logic core has the same external interfaces as TTA16 and these logic cores be used interchangeably within a design, as long as the correct code is provided for each processor as the instruction sets differ. Compared to TTA16, RISC16 operates at a lower frequency >100 MHz, but not higher than 140 MHz, and uses more logic slices, 320 compared to 200, than TTA16, but is easier to program. Creating a compiler would be easier as well, though beyond the scope of this work.

Data Cache

The data cache can operate at more than 150 MHz and uses about 80 logic slices when synthesised for a Spartan-3 FPGA (see Table 7.1). The data cache significantly improved average memory latency with the text-mode conversion algorithm, from 47 down to 5 clock cycles. This cache is 2 kB in size, a line size of 64 bytes, and is a 2-way set-associative design. The cache sense logic features a fast-hit path, with a latency of zero clock cycles, which is used when calculating a hit within the same cache-line as the previous request. And a slower-hit path, with a latency of one cycle, if the request is for a different cache line.

An additional feature is that this cache is a dual synchronous-clock design. It has two Wishbone interfaces, one that is connected to a processor and supports a data width of 16-bits. The other Wishbone interface is to be connected to the memory bus, which has a data width of 32-bits. The processor's Wishbone interface can be driven at a frequency which is an integer multiple of the memory bus clock frequency, as long as the two clock signals are synchronous, and both satisfy their required timing constraints.

Parameterisable SDRAM Memory Controller

This memory controller design is simple, fast, and supports burst data transfers. The maximum frequency at which the controller could operate, error free, was 120 MHz, which is 240 MB/s with a memory IC that has a 16-bit data bus, on a two layer PCB with a Spartan-3.

The bit widths of the address, internal, and external data buses are parameterisable. This memory controller can therefore be configured for many applications. The configuration used within OpenVGA was a 32-bit internal data bus width, an external 16-bit data bus width, and a 21-bit address width (the size and speed for this configuration are shown in Table 7.1).

As an example, the controller could be parameterised to support an external bus of 32-bits. This could be a design which used two 16-bit SDRAMs, or four 8-bit SDRAMs. An external data width of 32-bits would require that the internal Wishbone interconnect be 64-bits wide, since the external data bus operates at twice the frequency of the internal bus.

PCI-to-Wishbone Bridge

The PCI-to-Wishbone bridge allows data to be transferred between clock domains and bus protocols. The PCI Local Bus is a standard for transferring data between components on system boards within a PC. The Wishbone interconnect is intended for data transfer between logic cores within a single IC. This bridge features asynchronous FIFOs so that PCI and Wishbone can operate within separate clock domains. A FIFO is needed for each data-transfer direction, so two were used.

Even though the original PCI Local Bus specification is now considered obsolete, and succeeded by PCI Express [4], PCI is still popular and has a considerable installed base. The PCI Local Bus

is still found in many computers, even those currently being produced. Due to the size of the installed base, there are still reasons to use it, and therefore this logic core, in new designs.

Video Controller

This is a logic core that generates a data stream and the timing signals for driving VGA and DVI monitors (though DVI is untested and this is future work). This Wishbone-compatible core is sufficiently general-purpose and flexible to be useful for any application which connects to a VGA or DVI display. The core features a 2 kB prefetch queue, implemented using an asynchronous FIFO, so the dot-clock and Wishbone clock can be in separate clock domains. The data prefetch also makes the display redraw less susceptible to memory bus congestion.

Asynchronous FIFOs

Two asynchronous FIFOs were developed, one with 16 entries, and the other with a capacity of 2048 bytes. These FIFOs are optimised for the Spartan-3 architecture, to take advantage of the Spartan-3 RAM primitives for efficient implementation. The width of the data bus is parameterisable for the 16-entry FIFO. A width of 32-bits is used for the PCI-to-Wishbone bridge, and Table 7.1 lists the speed and size in this configuration.

Asynchronous FIFOs are important when multiple bit-width data needs to be transferred across clock domains. The design of these FIFOs ensures that metastability problems are avoided [5].

7.2.2 Processor Comparison and Summary

Two processors were developed for OpenVGA, a traditional RISC design, RISC16, and a processor with a novel transport-triggered architecture, TTA16. Relative to RISC16, the TTA16 processor uses fewer logic resources, operates at higher clock frequencies, and can perform more operations per clock cycle (see Section 4.5).

RISC16 has a 16-bit instruction word, instead of the 32-bit wide instructions of TTA16, leading to far better code density. It is also a simpler architecture to program in assembly language, data-forwarding and interlocking help here, and RISC16 is similar to processors with GCC ports. But due to the added complexity of RISC16 it is both larger and slower than TTA16.

A summary of the processors that were compared with TTA16 and RISC16 are listed in Table 7.2. All of the operating speeds and numbers of logic slices used are approximate and depend heavily on the parameters used, optimisation settings, and the other components within a design. This table just indicates what can be expected in best-case scenarios.

In terms of size and performance, both of OpenVGA's processor logic cores compare favourably to other FPGA-based processors, as Table 7.2 shows. TTA16 is the standout here in terms of performance, and RISC16 is fast as well. All processors listed have assemblers, but TTA16 and

CPU Name	Speed (MHz)	Size (Slices)	Compiler	Open- Source	Notes
MicroBlaze [43]	115	~1000	Yes	No	Proprietary
PicoBlaze [44]	90	90	No	Yes	Limited address space
ZPU [18]	90	~400	GCC	Yes	
OR1k [13]	30	2000	GCC	Yes	
GR0040 [16]	30	200	LCC	Yes	Restrictive license
RISC16	140	320	No	Yes	Lcc port possible
TTA16	190	200	No	Yes	Fast but difficult to program and poor code density.

Table 7.2: Evaluated processors and a comparison with TTA16 and RISC16.

RISC16 lack C compilers. ZPU, MicroBlaze, and OR1k are all 32-bit processors, so based on this table, ZPU is clearly the best of these. It would have been a good candidate for OpenVGA, but unfortunately this processor only recently reached a usable state [18].

7.2.3 Known Bugs

There exists a known bug within the SDRAM controller logic core and this is covered in Section 5.1.2. Fortunately, this is the only known bug with OpenVGA, and there is a software workaround. The size of burst transfers can be limited, which seems to solve the problem. Fixing this bug is future work.

7.3 Future Work

Open-source hardware is an area which is still in an early state of development. There is no free and open implementation of a VGA-compatible graphics adapter. An obvious direction for OpenVGA to take is enough VGA support to allow it to function as a PC's primary graphics adapter. OpenVGA was designed as a step towards VGA functionality and features a processor suitable for this task.

Another avenue of future work discussed here is software drivers so that OpenVGA can be used with the Graphical User Interfaces (GUIs) of modern operating systems, like GNU/Linux, Windows, and Mac OSX. Other areas for future work include improvements to the hardware, like PCI Express support, using a more modern FPGA, and greater memory bandwidth. These changes will require changes and additions to the current set of logic cores as well.

7.3.1 VGA Compatibility

OpenVGA does not currently have a VGA-compatible graphics adapter, this is due to the lack of the firmware for VGA emulation. Chapter 2 covers the components of a VGA which will have to be emulated. The quantity of firmware code required is substantial and was beyond the scope of this project. A C compiler would simplify this task, and could be another area of research.

7.3.2 Software

Due to PCI Plug and Play support, OpenVGA is detected and initialised correctly by the host system, but it is not currently taken advantage of by any operating systems since only a simple Linux kernel module has been written, and just for testing purposes. A driver for X11 or Microsoft Windows would allow OpenVGA to be a supported display device for these operating systems. Writing an X11 driver should be straight-forward, since the OpenGraphics project has an open-source driver available.

If OpenVGA is used for other uses, such as hardware computation, it will most likely require additional custom software and drivers, though the current kernel module may be enough for some applications, since it allows OpenVGA to be written to, and read from, like a file.

TTA processors that have multiple data-ports are very time consuming to program in assembly code. TTA16 would be a far more useful processor if there was a C compiler for it. There are open-source C compilers for other TTA processors [24, 6], these are GCC ports, and maybe one of these could be modified for TTA16. Generating optimised code for TTA processors is very difficult though.

7.3.3 Hardware Improvements

OpenVGA hardware works well and is low-cost but future modifications could significantly improve performance and features. By changing to a four-layer PCB, adding a PCI-Express PHY, an on-board clock-generator IC, a larger FPGA with more I/O pins, and larger, wider, and faster memory IC would allow OpenVGA to support high-resolution video modes, and allow it to be more useful for hardware computation tasks. A current area of research, for example, is real-time ray-tracing[40, 38].

Xilinx have recently introduced a new low-cost FPGA product line, the Spartan-6 family, replacing Spartan-3 generation FPGAs. These are faster and have more logic resources than their previous generation devices. FPGAs from other vendors should be evaluated as well, like the Cyclone III from Altera, or a Lattice ECP3¹. Another benefit of these newer FPGAs is they have hardware support for DDR2 and DDR3 memory, addressing a current OpenVGA weakness.

¹A disadvantage with Altera and Lattice is that these vendors do not have a free tool-chain that supports the GNU/Linux operating system, only currently Microsoft Windows.

7.3.4 Logic Core Improvements

There are two options for improving the current logic cores. The first approach would be changes that expand the functionality of OpenVGA, and the second class of improvements would optimise the current set of logic cores.

New Features

To complete VGA-compatibility, the PCI-to-Wishbone bridge would need modifying so that it identifies itself as a VGA device. This is straight-forward, and is covered here [33].

It is also common for graphics adapters to contain logic for hardware acceleration of 2D and 3D tasks. Some support for this could be done through firmware with the existing processors, new functional units could be added to these processors², or new logic cores could be added that provide this functionality.

Optimisations

While individual logic cores have been heavily optimised for speed and size, they could be modified to make more efficient use of the buses which connect them together. For example, the PCI-to-Wishbone bridge does not use the burst-transfer functionality of the Wishbone interconnect standard, nor does it support PCI burst read transfers. This is because the logic required to support this is quite complex. These changes will lower bus congestion, thereby increasing system performance.

Another area that can be optimised is the cache. Currently, when the processor performs a memory write it stalls until the completion of the transaction, which can be many cycles. A write-back cache design reduces the frequency of this problem. Also, when a memory read is issued by the processor, and if there is a cache miss, the processor stalls until the entire cache-line has been fetched. It would be possible to implement a system where the desired word is fetched first, and then allowing the processor to continue operation.

A change to the SDRAM controller, adding write-data FIFOs, would allow multiple memory writes to be queued, and completed at a time suitable to the memory controller. This modification would probably add 60 logic slices to the size of the controller though.

7.4 Conclusion

The design and construction of a low-cost, open-source graphics adapter has been presented in this thesis. Open-source hardware is a relatively new area of research. As long as the capabilities and

²Automatic synthesis of TTA processors is a possibility too [22], this would allow processors to be developed quickly for specific uses.

price of programmable logic keeps improving, open-source hardware will likely be an expanding domain too. This project contributes OpenVGA to this area. Both as a stand-alone project, and as a collection of free and open logic cores.

Bibliography

- [1] *Using Digital Clock Managers (DCMs) in Spartan-3 FPGAs*, 2003.
- [2] M. Arnold and H. Corporaal. Data Transport Reduction in Move Processors. In *Proceedings of the third annual conference of the Advanced School for Computing and Imaging*, pages 68–75, 1997.
- [3] R.L. Britton. *MIPS assembly language programming*. Prentice Hall, 2004.
- [4] R. Budruk. *PCI express system architecture*. Addison-Wesley Professional, 2003.
- [5] E. Clifford and C.P. Alfke. Simulation and Synthesis Techniques for Asynchronous FIFO Design with Asynchronous. Pointer Comparisons. *Inc, www.sunburst-design.com/papers*, 2002.
- [6] H. Corporaal. MOVE32INT Architecture and Programmer’s Reference Manual. *Delft University of Technology, EE Faculty, the Netherlands*, 1993.
- [7] H. Corporaal. Transport triggered architectures examined for general purpose applications. In *Sixth Workshop on Computer Systems*, pages 55–71, 1994.
- [8] H. Corporaal. TTAs: Missing the ILP complexity wall. *Journal of Systems Architecture*, 45(12/13):949–973, 1999.
- [9] Harvey G. Cragon. *Computer Architecture and Implementation*. Cambridge University Press, 2000.
- [10] C.E. Cummings. Simulation and synthesis techniques for asynchronous FIFO design. *Synopsys Users Group, San Jose, CA*, 2002.
- [11] Advanced Micro Devices. Ati radeonTMhd 4890 – overview. Technical report, Advanced Micro Devices, 2009.
- [12] C. Dufaza and G. Cambon. LFSR based deterministic and pseudo-random test pattern generator structures. *2nd ETC, Munich*, pages 27–34, 1991.

- [13] Unneback M. Erlandsson, M. and R. D’Addio. *OpenRISC 1000: Overview*. OpenCores, 2009.
- [14] Richard F. Ferraro. *Programmer’s Guide to the EGA, VGA, and Super VGA Cards*. Addison Wesley, 3rd edition, 1994.
- [15] Yukiho Fujisawa. *32-Bit SuperH Microcomputers*. IOS Press, 2005.
- [16] J. Gray. Designing a simple fpga-optimized risc cpu and system-on-a-chip. *DesignCon’2001*, 1:128, 2001.
- [17] Digital Display Working Group. *Digital Visual Interface, DVI*. VESA Inc., 1999.
- [18] O. Harboe. *ZPU - the worlds smallest 32 bit CPU with GCC toolchain: Overview*. Zylín AS, 2009.
- [19] J.L. Hennessy and D.A. Patterson. *PCI Hardware & Software: Architecture & Design*. RTC Books, 2003.
- [20] R. Herveille. *WISHBONE System-on-Chip (SoC) Interconnection Architecture for Portable IP Cores*, 2002.
- [21] J. Hoogerbrugge and H. Corporaal. Register file port requirements of transport triggered architectures. In *Proceedings of the 27th annual international symposium on Microarchitecture*, pages 191–195. ACM New York, NY, USA, 1994.
- [22] J. Hoogerbrugge and H. Corporaal. Automatic synthesis of transport triggered processors. In *Proc. First Ann. Conf. Advanced School for Computing and Imaging, Heijen, The Netherlands*, 1995.
- [23] NXP Inc. *NXP CBT16211 Datasheet*. NXP Inc., 2001.
- [24] P. Jaaskelainen, V. Guzman, A. Cilio, and J. Takala. Codesign toolset for application-specific instruction-set processors. In *Proc. SPIE-Multimedia on Mobile Devices*, 2007.
- [25] Khronos Group, Inc. *OpenGL Overview*, 2009.
- [26] N. Knutsson. An FPGA-based 3D Graphics System. Master’s thesis, Linköping Institute of Technology, 2005.
- [27] J. Lerner and J. Tirole. Some simple economics of open source. *Journal of Industrial Economics*, pages 197–234, 2002.
- [28] G. McFarland. *Microprocessor design*. McGraw-Hill, 2003.
- [29] Micron. *Synchronous DRAM, 64mb: x4, x8, x16 sdram features edition*, 2008.

- [30] Timothy Miller. The open graphics project, 2009.
- [31] B. Parhami. *Computer Architecture: From Microprocessors to Supercomputers (Oxford Series in Electrical and Computer Engineering)*. Oxford University Press, Inc. New York, NY, USA, 2005.
- [32] Richard P. Paul. *SPARC Architecture, Assembly Language Programming, and C*. Prentice Hall, 2nd edition, 2000.
- [33] PCI SIG. *PCI Local Bus Specification Revision 2.2*, 1998.
- [34] P.J. Salzman and O. Pomerantz. The Linux Kernel Module Programming Guide. *Url: <http://jamesthornton.com/linux/lkmpg>*, 2001.
- [35] David Seal. *ARM Architecture Reference Manual*. Addison-Wesley, 2nd edition, 2001.
- [36] Tom Shanley and Don Anderson. *PCI System Architecture*. Mindshare, Inc., 4th edition, 1999.
- [37] JE Smith, S. Weiss, C.R. Inc, and C. Falls. PowerPC 601 and Alpha 21064: A tale of two RISCs. *Computer*, 27(6):46–58, 1994.
- [38] Ward R. Suggate, P. and T. Molteno. Massively Parallel MIMD Processing in Hardware, 2007.
- [39] G. Suttty and S. Blair. *Advanced Programmer’s Guide to SuperVGAs*. Addison Wesley, 2nd edition, 1994.
- [40] I. Wald. *Realtime ray tracing and interactive global illumination*. PhD thesis, Universitätsbibliothek, 2004.
- [41] Roy Ward. *Assembly Code Generation from XML Processor Descriptions*, 2007.
- [42] Roy W. Ward and T.C.A. Molteno. Counter Representations in Microprocessors. In *ENZ-Con06 Conference Precedings*. ENZCon, 2006.
- [43] Xilinx Inc. *Microblaze processor reference guide*, 2008.
- [44] Xilinx Inc. *PicoBlaze 8-Bit Embedded Microcontroller for Spartan-3 Generation, Virtex-II, & Virtex-II Pro FPGAs*, 2008.
- [45] Xilinx, Inc. *Spartan-3 FPGA Family: Complete Datasheet*, 2008.
- [46] Joseph Yiu. *Computer Architecture and Implementation*. Newnes, 2007.

Appendix A

Source Code Layout and Overview

OpenVGA is an open-source project, licensed under the GPL. All of the source code is available on the Internet. The URL is <http://openvga.sourceforge.net/> and also contains the snapshot of the project at the point of the completion of this thesis.

A.1 Top-Level

Also included in the top-level are the standard LICENSE and README text files which are included within most open-source projects.

rtl/	Verilog RTL (Register Transfer Level) description of the design. There is a Makefile to synthesise the design using XST and targeting the Spartan-3 FPGA family.
sim/	Verilog RTL (Register Transfer Level) description files for the test-harnesses of the logic-cores in this project.
src/	OpenVGA source code for kernel driver, the TTA assembler, the RISC16 assembler, the cache and text-mode simulators, and a partial VGABIOS implementation.
tools/	Scripting files used for tasks such as assigning data to BRAMs, converting fonts, and a CRT simulator.

A.2 HDL Hierarchy

The Verilog source code is consists of both logic-core source files and hardware testbenches. Simulation test-harness code is in a separate **sim** directory.

rtl/	The top-level OpenVGA modules, testbenches, and makefile.
rtl/cpu/	Both the RISC16 and TTA16 processors, and hardware testbenches.
rtl/lib/	Library containing copies of modules not specific to the OpenVGA project. This includes LFSRs, multiplexors, FIFOs.
rtl/misc/	Miscellaneous modules, and testbenches, that are too OpenVGA specific for the “lib” directory.
rtl/pci/	The files for the PCI-to-Wishbone Bridge logic core and a hardware testbench.
rtl/sdram/	SDRAM controller logic core modules and hardware testbench.
rtl/video/	The OpenVGA video logic core. There are CRTC, prefetch buffer, redraw logic, and hardware testbench modules.

A.3 Simulation Test-Harness Hierarchy

This directory contains the Verilog simulation files for the logic cores within the **sim** directory.

sim/	The top-level OpenVGA, module test-harnesses, and Makefile.
sim/cpu/	RISC16 and TTA16 simulation test-harnesses for use with Verilog simulators.
sim/lib/	Library containing test-harnesses for modules not specific to the OpenVGA project. This includes test-harnesses for LFSRs, multiplexers, FIFOs. Additionally, many Xilinx built-in primitives were emulated and these modules and test-harnesses are within this folder as well.
sim/misc/	Miscellaneous test-harnesses for modules that are too OpenVGA specific for the “lib” directory.
sim/pci/	Test-harnesses used while developing the PCI logic core.
sim/sdram/	SDRAM controller test-harnesses.
sim/video/	The OpenVGA CRTC, prefetch buffer, redraw logic, and testbenches.

A.4 Tools

There were many scripts and other tools created for this project. They are typically used for converting data between the many formats used, like font data, creating Verilog include files, and other post-processing tasks.

Other tools include `peek`, `poke`, `set`, `test`, and other OpenVGA test applications that were used for developing and testing the PCI bridge and SDRAM controller. These were run on the host computer that OpenVGA was attached to.

Appendix B

OpenVGA Components and PCB

This appendix provides more detail of the OpenVGA hardware. The CADSoft EagleTM schematic, board, and parts-list files are available on the Internet from <http://openvga.sourceforge.net/>. OpenVGA uses of a two-layer PCB for cost and simplicity and the PCB form-factor is a full-height, short, PCI expansion board with rear I/O panel connectors (see Figure B.1). Two-layer PCBs restrict the choice of electronic components due to routing and signal integrity issues. Data-sheets of high-pinout BGA ICs often recommend six-layer PCBs or more, for example [45].




Figure B.1: OpenVGA PCB artwork, both sides.

The FPGA has mostly user assignable inputs and outputs so this allowed flexible placement of most of the other ICs. The video transmitters could be near the back edge of the PCB and SDRAM could be placed so as to minimise trace lengths.

The components were laid out with the objective of having the majority of the signal routing on the top-layer (the component-side) so that the bottom-layer (the solder-side) could be used for routing the power-supply and ground nets. The VGA and DVI transmitter ICs were mounted on the solder-side as well, and this was to simplify routing. The decoupling capacitors and termination

resistors were mounted on both sides of the PCB since they were required to be as close as possible to their related ICs.

B.1 Off-Board Connections

OpenVGA has connectors for PCI, DVI, VGA (through the DVI connector, using an adapter), JTAG, and USB (see Figure 3.1. Only the JTAG connections are connected directly to the FPGA, the rest use encoder or voltage-translation ICs and are detailed below.

B.2 Integrated Circuits

Table B.1 is a list of the ICs chosen for OpenVGA. More complete descriptions are included in later sections of this appendix. Figure 3.1 shows the ICs which are on the top (component layer) of the PCB, the two video encoders, the TMDS encoder and the video DAC are shown in Figure B.2.

Part#	Description
XC3S200-4PQG208C	Xilinx Spartan-3 FPGA, 200k Gates
MT48LC4M16A2TG-75	Micron 8 MB SDRAM
TDA8777HL/14/C1,15	Philips video DAC, 10-bits/colour component
TFP410PAP	Texas Instruments TMDS encoder
LM1086IS-ADJ	National Semiconductor SMT, adjustable, voltage regulators
CBT16211DDG	Fairchild Semiconductor 24 I/O, Bus Switches
XCF04SVOG20C	Xilinx 4 Mb, Platform Flash, Serial PROM
FT232RL	Future Technology Ddevices International USB UART
SG-8002JA-MPT	Epson Toyocom Corporation 50 MHz, SMD oscillator

Table B.1: OpenVGA integrated circuits.

B.2.1 The Xilinx Spartan-3 FPGA and SPROM

The core of OpenVGA is a FPGA which contains the logic necessary for implementing this display device. The FPGA is a 200 k-Gate, Xilinx Spartan-3, XC3S200 FPGA, which is from the Xilinx low-cost product range. It has a quantity of programmable logic that Xilinx considers as approximately equivalent to an IC with about 200,000 gates.

The FPGA can be programmed using a JTAG boundary-scan chain or on power-on using a Xilinx SPROM. The state of the Spartan-3 is lost on power-off [45] so a SPROM (Serial Programmable Read-Only Memory) is used to restore the state at power-on. The SPROM is programmed using the JTAG boundary-scan chain as well.

The Spartan-3 requires multiple supply voltages. This Spartan-3 core logic runs at 1.2 V, the I/O banks at 3.3 V, and the FPGA configuration circuitry at 2.5 V. This complicates the power-supply routing when using a two-layer PCB, but decoupling capacitors were used in accordance to Xilinx specifications and no problems were experienced, even though Xilinx recommends separate power-planes for each supply.

B.2.2 The DVI and VGA Video Encoders

The DVI TMDS encoder and the VGA video DAC are used to convert the parallel, digital video signals generated by the FPGA into the signals (analogue for VGA, TMDS for DVI) used by an attached DVI or VGA computer monitor. The TMDS encoder and video DAC operate at high frequencies for a two-layer board so termination resistors are used to improve signal quality. An early development version of OpenVGA (as shown in Figure 3.3) had significant signal integrity issues. Neither the DDR SDRAM or the TMDS encoder operated correctly. The signal waveforms, as displayed on an oscilloscope, had significant undershoot and overshoot. The current OpenVGA board with the terminating resistors has very good signal integrity for a two-layer board.




Figure B.2: OpenVGA solder-side components.

B.2.3 USB UART

A logic core was developed which allows the OpenVGA processor to read to, and write from, the USB UART. This can be used for debugging as well as interfacing with the many devices which use this standard.

B.2.4 PCI Bus Switches

The PCI Local Bus uses half-wave reflection signalling [33, 36] which causes voltage transients that exceed the maximum operating conditions specified for the Spartan III family[45]. Two 24 pin bus

switches were used to prevent PCI signals from exceeding 3.4 V and -0.7 V. These switches add a propagation delay of 250 ps [23].

Appendix C

Wishbone Interconnect Overview

C.1 Wishbone Interconnect Introduction

Wishbone was designed as an internal interconnection standard for System-on-a-Chip (Soc) applications (see www.OpenCores.org, the current home of Wishbone project, and the contains the full specification). Wishbone is an intended as a solution to provide a standard interface for communicating between logic cores from multiple vendors. It can be difficult to integrate logic cores if they each use a different interconnect scheme, and this can also add significant overhead (logic and latency).

C.1.1 Features

The Wishbone design was inspired by traditional microcomputer buses[20], like PCI and VME, as these are general purpose interconnects that are flexible and robust.

A brief summary of important Wishbone features:

- Supports reads and writes, burst and single-word transfers.
- Simple, logic resource requirements are low. At its simplest, an interface to a synchronous SRAM requires just one logic gate.
- User-defined tags allows the addition of extra signals.
- Support for data bit-widths of 8, 16, 32, and 64.
- Any address bit-width is supported, even zero, depending on the device.
- Synchronous; all transfers occur on clock edges.
- Choice of topology, i.e. bus, point-to-point, crossbar switch, is up to the system designer.

- Not encumbered by patents, free for all to use.
- A single device can be both a master and a slave.
- Offers some support for handling errors.

C.1.2 A Simple Example

In the timing diagram shown in Figure C.1.2, the simultaneous assertion of CYC and STB select the slave device, the WE signal indicates this is a write transaction.




Figure C.1: This is an example implementation of a basic master-slave, point-to-point Wishbone interconnect. The “SYSCON” design is not part of the Wishbone specification, it maybe implemented however the designer sees fit.

Source: 'WISHBONE System-on-Chip (SoC) Interconnection Architecture for Portable IP Cores',
Revision: B.3, Released: September 7, 2002.

C.2 Description of Wishbone Signals

This section gives a brief summary of the main Wishbone signals. User-defined tags were not used for OpenVGA's internal bus so these signals are not described. Additionally, Wishbone bus ports to HDL modules used within OpenVGA include both a prefix and a suffix. For example, the

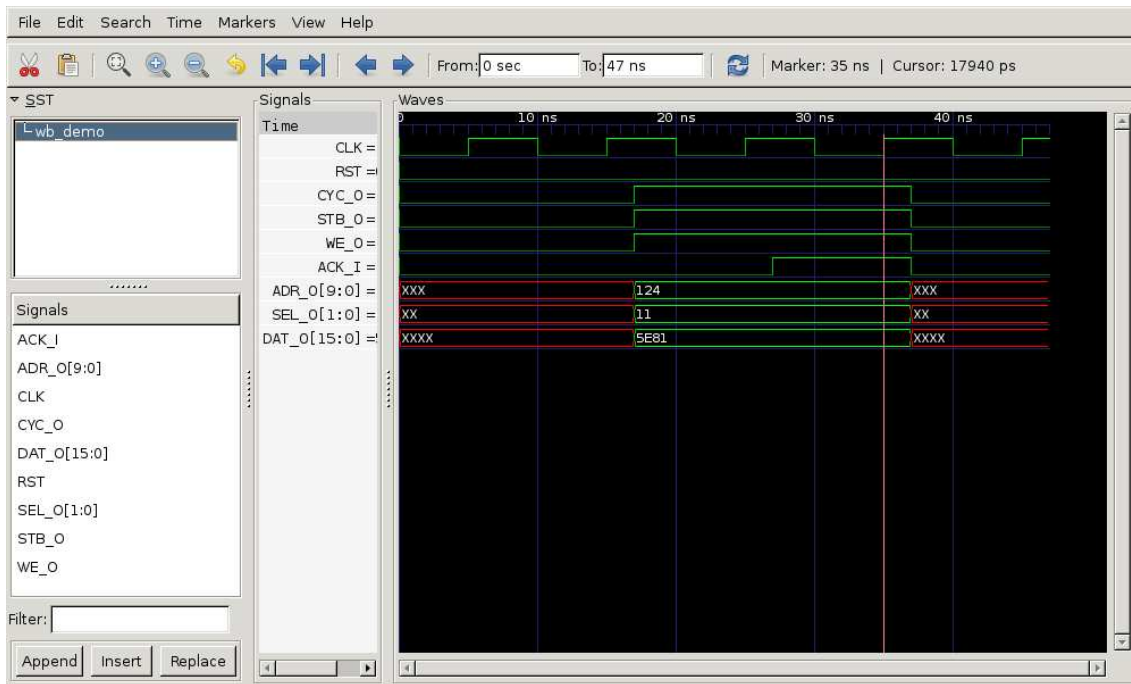


Figure C.2: One 16-bit word is written by a master to a slave device. The signals CYC.O and STB.O are asserted by the master to frame the transaction. The ACK.I response is received by the master and indicates that the slave has completed the transfer.

prefix “wb” of signal “wb_cyc_o” indicates that the signal is a Wishbone signal, and the suffix “o” indicates that this is an output. (The output direction of the CYC signal would also indicate that the module this signal belongs to is a Wishbone master.)

- **CYC:** This signal is asserted by a master and it claims ownership of the Wishbone bus. Other masters on a bus are prevented from asserting their CYC when a CYC is already asserted. This signal is still used for non-bus topologies, like point-to-point, and is used in combination with STB to select a slave device.
- **STB:** This signal is asserted by a master and frames a Wishbone transaction. If a slave device has its STB input signal asserted, then it has been selected for a Wishbone transaction.
- **WE:** This signal is driven by a master and selects whether the requested transaction is a write or a read (signal levels high and low respectively).
- **ACK:** This signal is asserted by a slave to indicate that it has either received data, if the transaction is a write, or has placed valid byte selects and data on its data outputs, for a read transaction. This signal is asserted just once for each data transfer, and can take many cycles to assert, depending on the latency of the device.

- **RTY**: This signal is asserted by a slave to indicate that the device is currently unable to fulfil the current Wishbone request, and to retry later.
- **ERR**: This signal is asserted by a slave to indicate that an error has occurred. Wishbone does not specify how errors are handled except that the slave must leave the bus in a usable state at the end of the transaction.
- **CTI[2:0]**: This signal array is driven by the master and can request a burst transfer. Even if the master specifies a burst mode, the slave is not required to perform a burst transfer. Also, if the slave supports burst transfers but not the mode requested, it responds with single transfers. It is optional to implement this signal, it has to be set to zero for a slave that supports this signal when a master does not. A full description of the modes is available in the Wishbone specification[20].
- **BTE[1:0]**: This signal array is driven by the master and specifies the Burst Type Extension. This signal can only be implemented if CTI is also implemented, and is required to be implemented if CTI is implemented and the device supports burst transfers. A full description of the modes is available in the Wishbone specification[20].
- **ADR**: This signal array is driven by the master and is the address of the data requested for the transaction. This signal is optional, and its width is device dependent, a FIFO may not need an address, for example.
- **SEL_O[n-1:0] & SEL_I[n-1:0]**: These signal arrays are driven by both masters and slaves and are used to select valid bytes within a data word. A 64-bit transfer has eight select bits, so $n = 8$, for example. During a write transaction, the master asserts the valid bytes on its own SEL_O outputs, and during a read transaction, the slave asserts the byte-valid selects on its SEL_O outputs. A device that supports reads and writes is required to implement both SEL_O and SEL_I. These signals have to be valid when the device responsible for driving them has its STB or ACK signals asserted.
- **DAT_O[m-1:0] & DAT_I[m-1:0]**: These signal arrays are driven by both masters and slaves and are used to transfer data across the Wishbone interconnect. The width, m , must be one of 8, 16, 32, or 64. A master performing a write must have valid write data on its DAT_O signals whenever its STB signal is asserted. A slave responding to a read request must have valid data on its DAT_O outputs when its ACK signal is asserted. For point-to-point applications the DAT_O of a device is typically connected to the DAT_I of the other device.

All wishbone control signals are active when high (or “1”) and Figure C.1.2 shows an example transfer using many of the signals described above.

Appendix D

TTA16 Assembly Language Programming Guide

Programming for the TTA16 architecture, and accessing the available OpenVGA hardware using the processor, requires significant documentation. This is because many modules have been created specifically for OpenVGA, so no documentation exists elsewhere, and the TTA16 architecture is unusual. This appendix does not attempt to provide exhaustive information, but hopefully enough to understand the source-code in the examples, and the firmware routines if necessary.

D.1 TTA16 Overview

Programming for the TTA16 architecture is unlike programming for more traditional architectures since a TTA with multiple transports has explicit-ILP (Instruction Level Parallelism). TTA16 instruction consists of multiple micro-instructions, and each being simple data moves between registers. A traditional CPU instruction typically specifies just a single operation to be performed – the *opcode* – and the arguments: registers, immediate constants and memory data.

A TTA processor contains three directly accessible register types: trigger registers, operand registers, and results registers. A write to a trigger register initiates an operation within its associated FU. A write to an operand register has no side-effect, that value is simply changed. Operand registers are used when multiple inputs to a FU are required. Subtract, for example, has two inputs, the minuend and the subtrahend. With TTA16, the minuend is the trigger register and the subtrahend input is an operand register.

The output of the subtract unit is a result register. These registers are modified by its associated FU a certain number of cycles, depending on the FU, after the trigger register is written to. These results can then put back onto a data transport to be used for subsequent operations.

Here is an example subtract operation for a TTA processor with one transport:

```
{r0      ->add    }
{r1      ->addt   }
{diff    ->r2      }
```

As part of the TTA assembler’s syntax, curly braces ({}) surround an instruction. The reason for syntax is to make obvious the explicit-ILP of the processor when the processor has multiple transports (see the `memcpy` programming example in Figure D.4 at the end of this appendix for assembly code written for TTA16 which has multiple transports).

On the first line, the argument `r0` is a general purpose register from the RF (Register File), one of 16 (`r0-r15`). The arrow `->` indicates the data direction, from `r0` to `add` in this case. The last argument is the destination register, which is the operand register for the addition FU. Since `add` is an operand register, this is a simple data move, there is no side-effect, so the start of an addition will not be triggered.

The second line moves another register to the addition FU, but this time to the trigger register. This instruction will cause an addition operation to begin. The sum of the two input registers will be stored back to the RF by the third instruction.

The programming tasks with TTA16, especially with high-performance code, are more explicitly concerned with scheduling the flow of data to and from FUs (Functional Units) via ‘transports’. TTA16 is not fully connected either, each transport can only read certain registers, and direct data to just a subset of total FUs, which increases programming complexity.

D.1.1 TTA16 Datatypes

The only native data types supported by TTA16 are 16-bit, signed and unsigned words. There is a little hardware support for larger word sizes though. This is achieved using the borrow flag for long subtractions, and the multiplier calculates a 32-bit product from two 16-bit inputs. Either half of the product can be read, just as any other 16-bit register.

Immediate constants are either 8-bits, sign-extended to 16-bits, or 11-bits sign-extended (which is only available to transport 0 as it is needed for branching). Producing longer immediate values requires multiple operations, for example using the multiplier as a shifter.

D.1.2 Instruction Format

Like RISC processors, TTAs tend to have fixed width instruction words [6]. Typically within an instruction word, all fields are predefined, whereas RISC tends to have multiple instruction formats[32, 37, 3, 46]. A disadvantage with TTAs is that instruction words are often wider, or

loading large immediate data values is more complex¹.




Figure D.1: TTA16 instruction format, see text for a complete explanation of the bit-fields.

TTA instructions consist of a list of data moves to complete during the transport stage of the pipeline, and often some extra information like guards[6], register file indices, and/or immediate data.

The TTA16 instruction word format is shown in Figure D.1. The fields labelled *SRCx*, where *x* is 0, 1, or 2, are the bit-fields which select the source register for the data-transport. The *DSTx* bit-fields specify the destination registers. The *COM* register is a special source register that is visible to all other transports, and is an operand register for most FUs.

The next fields select registers from the RF to read and/or write. The single *B* bit selects the bank, i.e. registers 0-7 or 8-15 for the instruction. Since register write index bit-fields follow the current instruction, and read bit-fields precede it, a register read and write instruction involving different banks is still possible, like the following part-instruction (a valid move for transport 3):

r0 -> r15

The final bit-field is the immediate data field. This is 8-bits wide, with the upper bit representing the sign for sign-extending the integer to 16-bits. Data transport 0 uses an 11-bit immediate, by using the value of *REG1* as well, so that branches can reach all locations within the TTA16 instruction block RAMs.

Only a subset of TTA16's registers are available to each data transport (see Table D.1). This is so that the internal multiplexers can be smaller, and fewer logic layers, resulting in a more efficient implementation. Since the *com* register can read most registers, and all transports can read *com*, it is believed that this restricted set of registers, accessible to each transport, is a good trade-off between size and clock-speed, and code performance.

An example move for transport 0 could be:

0x023 -> bra

Which means move the hexadecimal value 0x023 (in base-10, 35) to the unconditional branch register, i.e. jump to address 0x023 and start executing instructions there. Since there is a latency of two cycles, the instruction following this branch instruction would be executed before the branch is completed too.

¹For example, MOVE32INT, a TTA CPU designed at TU Delft by H. Corporaal et. al. has only 6-bit immediates, but upto four can be used per instruction[6].

From transport 0, it would not be possible to write 0x023 to the *sub* register, for example, since this register is not visible to this transport. But it would be a valid move when performed in transport 1 .

SRC0	DST0	SRC1	DST1	SRC2	DST2	COM
com	nop	com	nop	com	mem/nop	com/nop
IMM11	bra	IMM8	sbb	diff	RF0	IMM8
RF1	rad	RF1	sub	RF0	mul	RF0
mem	wad	mem	cmp	pc	msr	mem
	jb	nand				bits
	jnb	and				diff
	jz	or				plo
	jnz	xor				phi

Table D.1: TTA16 registers and the transports which they are accessible.

D.1.3 Data Transport Scheduling

All FUs, except the Wishbone interface FU, have a fixed latency, so a result is available a deterministic number of cycles following triggering, and will be over-written by subsequent FU triggering, after its fixed latency has elapsed. For example, the subtract FU has a latency of two cycles, as do all FUs except the register file (RF), this has a latency of three due to pipelining. Even the branch FU has a latency of two, causing the instruction following the branch to be executed. The instruction slot directly after a branch instruction is called a branch delay slot[32].

If a FU is read one cycle after it was triggered, triggering is by a write to a FUs trigger register (TR), it will still contain the previous value. Traditional architectures often stall the CPU until the result is available, or use data forwarding to get early access to the result. This approach was not taken with TTA16 for two reasons:

1. The extra hardware required for hazard detection, CPU stalling, and/or data forwarding would lead to a larger CPU, reducing one the key benefits of TTA processors.
2. It is not always desirable, especially with RF bandwidth being comparatively more scarce than other architectures. Taking advantage of this latency allows ‘time-sharing’ of a FU, like a subtractor, for the decrementing of two variables without having to write either result back to the register file.

There is one exception though, TTA16 does stall when a Wishbone access is initiated, resuming upon completion of the transaction. Though the interlock logic required to implement this carries

a performance penalty, it is less than a more general data hazard detection scheme, and Wishbone accesses are non-deterministic. A low silicon-cost solution would be to poll a machine status register until a Wishbone transaction completion signal is received, but this would lead to a large increase in code size.

Traditional architectures use the RF extensively, as does RISC16 (see Appendix E), typically two register reads and one write per instruction. While TTA16 has a RF, with a capacity of 16, 16-bit values, typical code only accesses the RF only about once per instruction (see Listing D.4 for some typical assembly code). This is because moving data between the RF and FUs uses up available processor data-move bandwidth. It is more efficient to move data from one FU straight to another.

While not often used, TTA16 is still capable of up to two register reads, one write, and using one unique immediate data value within a single instruction, similar to more traditional architectures. There are additional restrictions though, both register reads have to be from the same bank, register field sharing, and register fields from the preceding and following instructions determine the registers read and written respectively (see Figure D.2).

Examining the instruction word diagram in Figure D.1, it may be noted that there are only two register fields within the instruction word. The first two transports (0 and 1) share *REG1* and the last two transports (2 and *COM*) share *REG0*, with *REG0* optionally specifying the write register.




Figure D.2: TTA16 shared bit-fields between instructions. A TTA16 instruction uses the *B*, *REG0* and *REG1* fields from the previous instruction for register reads, the *REG0* field from the next instruction for register writes, and there is overlap of the *IMMED* and *REG1* bit-fields when using immediates with transport 0. Long immediates are so branches can reach every address in block RAM.

Additionally, it is the register fields of the preceding instruction word that contain the indices of registers to be read for the current instruction, and the following instruction contains the index of the register to be written. Because of these architectural peculiarities, register operations can be difficult to schedule correctly. Often it is not until code assembly is attempted that conflicts are discovered.

Due to the explicit ILP, interleaved instruction bit-fields, and only partial data-transport connectivity, tightly-packed, well-scheduled code is very time consuming to construct by hand. No

compiler or other tools for instruction scheduling were written or readily available to ease this task either. The main goal when implementing the OpenVGA firmware was to ensure that critical loops are highly optimised (like with the memory copy subroutine, see Listing D.4).

D.1.4 Exceptions

TTA16 has no support for exceptions due to the complexity of saving and restoring the state of a TTA processor. For example, some registers of the TTA16 processor are write only which means that these functional units would be difficult to use during an interrupt. The multiply FU, for example, uses a dedicated, write-only, triggering register as one input, the TTA16 COM register for the other, so it could be possible to work backwards to deduce the value within the triggering at the time of the interrupt, but the processor would need its own set of working registers while performing the calculations.

Exception complexity was avoided by avoiding exceptions/interrupts altogether. The TTA16 within OpenVGA does not need to respond immediately to any I/O requests, PCI requests go straight to local memory, so exceptions are not needed

D.2 A Simple Example

The following is a piece of TTA16 assembly code which sets OpenVGA's LEDs. Braces surround a complete instruction, commas separate the different execution transports, which control the data transports, and there are four streams per instruction. Each stream has access to just a subset of TTA16's total register set.

```
;-----
; set_leds - The two LSBs of 'r0' determines the LED outputs.
set_leds:    {          ,          ,          ,1      } ; 1
              ; Set segment to the memory-mapped LEDs.
              {          ,          ,com  ->msr   ,0x48  } ; 2
              {\r0    ->wad  ,          ,\r0    ->mem  ,1      } ; 3
              ; Restore SS and return.
              {\r15    ->bra  ,          ,com  ->msr   ,\r12  } ; 4
              {          ,          ,          ,        } ; 5
```

The first two lines are just comments, since the line begins with the ';' character, and any text after this – until the end of the line – is ignored. The assembler supports two types of comments, single-line and multi-line comments, which is any text between the comment opening sequence '/' and the closing sequence '*/'.

This is a line-by-line break-down of the above example, only lines with instructions ({...}) are counted:

1. This line begins with a label so that branches to this routine can use a convenient moniker, rather than using the exact numerical value of the address. It would be difficult to determine this address while writing code too, as its value is only generated during assembly (using a pseudo-random count order calculated by the MFSR used).

The purpose of this line is to load the number one into the `com` register, to be used by the following line. Note that fields can be left blank and the assembler automatically inserts no-operations (the `nop` opcode can be explicitly used too) for that transport. The commas are required though.

2. The numerical value of one is now within `com`, placed there by the previous instruction, and by moving the contents of `com` to the `msr` register, the desired MSR (Machine Special Register) is selected, in this case it is the write address segment register. This register is used whenever a Wishbone write transaction is initiated, the lower seven bits of this register form the upper seven bits of the 23-bit Wishbone address.

The final field in the instruction is the segment address for the LEDs FU, and this is the value that will be loaded into MSR #1, the write address segment register.

3. There are three used fields within this instruction, the left-most operation is a data move to the write address register (`wad`), which is a trigger register, to initiate a Wishbone write transaction. This value will be ignored by the LEDs FU, as it does not use the lower 16-bits of the Wishbone address field, so any value can be used. It is the segment address value that is important.

The third field writes the argument passed to the `set_leds` routine, which is passed via `r0`, to the LEDs FU. The data value '3' would set both LEDs, for example.

Notice that the registers are preceded by a '`\`' character, this is optional and is simply a mnemonic to remind the programmer that the RF index bit-field is actually stored in the previous instruction, for RF reads. This fact would prevent any RF reads in the preceding instruction (`#1`) from operating as expected.

The final field performs the same function as the first line, sets `com` to '1' for use with following instruction.

4. This instruction achieves two purposes, the first field causes the program flow to branch back to the location after the instruction that called this routine (which means that the instruction within the branch delay slot after the instruction calling the `set_leds` routine is executed twice). By convention, this location is stored within register `r15`.

The last two fields of this instruction are used to restore the value of the stack segment address register. This was modified at the start of this routine, and subsequent instructions executed after the completion of this routine will likely use this register and expect it to have

the correct contents. This is another programming convention, as is using `r12` to store the stack segment register value.

5. This is the branch delay slot instruction, but nothing is needed to be done in this slot so this is a simple `nop` instruction.

D.2.1 Functional Units

Since TTA processors can typically execute multiple operations simultaneously, the FUs need to be able to operate in parallel. TTA16 has bitwise logic unit, a multiply unit, a subtractor, a branch unit, a Wishbone interface unit, and a Machine Special Register (MSR) unit. Upto three² of these can be triggered, and therefore execute in parallel, each clock cycle. This is because there are three data-transports that can write to trigger registers.

Table D.1 shows the which data transports can access which registers (or aliases), and Table D.2 shows which FUs each register (and its aliases) maps to. For example, the branch unit can only be accessed from transport-0, and the *COM* transport is also connected to the operand register port of most of the FUs, as well as all the other transports. In practice, this tends to be the most heavily used transport.

Name	Inputs	Outputs	Aliases/Modes
Multiply	com, mul _t	plo, phi	mul
Branch	bra _t	pc	bra, jz, jnz, jb, jnb
Subtract	com, sub _t	diff	sub, sbb, cmp ⁺
Bitwise Ops	com, bits _t	bits	and, or, xor, cmp ⁺
Wishbone	mem, ad _t	mem	rad, wad
MSR	com, msr _t	N/A	msr

⁺See the ALU section and its associated footnote.

Table D.2: Functional unit summary.

Arithmetic Logic Unit

Traditional RISC processors have an ALU (Arithmetic Logic Unit) that perform arithmetic and logical operations on integers, but TTAs use a multiple FUs attached to data transports to achieve the same functionality. This is so multiple integer operations can occur simultaneously without requiring multiple ALUs.

²A write to *cmp* ‘register’ causes both the bitwise and subtract FUs to be triggered. This is so the zero flag (*ZF*) and borrow flag (*BF*) are modified simultaneously, which is useful for conditional branching.

The ALU functionality of TTA16 is obtained from just two FUs, the bitwise operation FU, and the subtract FU. The missing logical operators from these two FUs are the common left and right shift operators. Explicit shift FUs are not provided but the same functionality can be obtained using the multiplier³.

Memory Segment Registers

TTA16 has two parameterisable MSRs (Memory Segment Registers), each up to 16-bits in width, to allow memory addresses greater than 16-bits wide to be supported. The memory segment value is concatenated with the memory pointer value to give a total width of up to 32-bits, and with two bytes per address, this means up to 8 GB is addressable with this architecture.

The two MSRs are called the Data Segment (DS) and the Stack Segment (SS) registers which indicates their possible uses, though no such limit upon such use is incorporated within the hardware. When performing large block copies, for example, they both could be used as data registers.

By default, though this is set by the assembler/convention, not by the hardware, the DS is used with load and store instructions. To use SS, instead of DS, the suffix ‘.sf’ is appended to the load/store instruction.

D.3 OpenVGA Memory-Mapped I/O Peripherals

OpenVGA memory space is divided into two regions, Memory-Mapped I/O (MMIO) or SDRAM. Table D.3 lists the MMIO devices implemented within OpenVGA, and the associated segment to access the device.

I/O Device	Address Segment
CRTC	0x0040
SPROM	0x0044
LEDS	0x0048
DMA	0x0050
Cache Flush	0x0058

Table D.3: TTA16 memory-mapped I/O modules and there corresponding address segments.

³Multiplying by two is the same as a left shift. Multiplying by 32,768 and discarding the low 16 bit word of the 32 bit product is the same as a right shift of one place.

D.3.1 Cache Flush Peripheral

Module:	wb_cache_flush	
Description:	Flushes all data from the data cache.	
Related Files:	/rtl/cache/wb_cache_flush.v	
Testing Files:	/sim/cache/wb_cache_flush_tb.v	
Author:	Patrick Suggate	License: GPL

It is possible for the contents of the cache and the contents of OpenVGA's main memory to become incoherent. For example, if OpenVGA's processor is performing a data conversion, reading from an arbitrary block 'A', and writing to block 'B', and a PCI write transaction modifies the contents of block 'A' in main memory, but data from this location has already been cached, the contents of main memory and the data cache are now incoherent.

The solution used within OpenVGA is for the processor to manually flush the data cache, so that fresh data is retrieved from main memory. This peripheral allows the manual flushing of each of the 16 cache-lines.

Writing to this device 16 times, to each of the 16 cache-line addresses, completely flushes the cache. Any value for the write data will do since it is not used by this peripheral. The 16 addresses which invalidate the cache-lines are: {0x0000, 0x0020, ..., 0x01e0}

D.3.2 CRT Controller

Module:	wb_crtc	
Description:	Generates the timing signals for a CRT or LCD display.	
Related Files:	/rtl/video/wb_crtc.v /rtl/video/crtc.v	
Testing Files:	/sim/video/crtc_tb.v	
Author:	Patrick Suggate	License: GPL

The CRTC generates the display timing signals using the character clock, which runs at one eighth of the dot-clock frequency. This means that the horizontal timing signals (the first four within Table D.4) need to be multiplied by eight to obtain the values in dot-clocks/pixels.

The default values correspond to the 640x480, with a refresh rate of 60Hz, video mode, and the dot-clock is 25 MHz. These values are a running total so each value in the table, for either the horizontal or vertical totals, is the previous value plus the value of the current parameter. For example, the width in characters of the horizontal back porch is six, or 48 pixels, and this is added to the 11 character clock ticks of the horizontal sync signal, to give the running total of 17 ticks. Other modes are listed in Table 6.3.

CRTC Port	Default Value	Register Function
0	11	Hsync duration
1	17	Hroiz. back porch
2	97	Horiz. active
3	99	Hroiz. front porch
4	1	Vsync duration
5	34	Vert. back porch
6	514	Vert. active
7	524	Vert. front porch

Table D.4: CRTC ports and their defaults values.

D.3.3 DMA Unit

Module:	wb_dma	
Description:	Writes buffered data to main memory in burst-mode.	
Related Files:	/rtl/wb_dma.v /rtl/util/pre_read.v	
Testing Files:	/sim/wb_dma_tb.v /sim/xilinx/RAMB16_S18_S36.v	
Author:	Patrick Suggate	License: GPL

Since all direct memory writes performed by the CPU pass through the cache, which is a write-through design, they have a high latency penalty. Since the memory bus is shared by the PCI and display redraw units too, numerous CPU memory writes will significantly increase memory congestion. This is because a CPU to memory write is 16-bits wide and atomic, whereas the memory controller is designed for 32-bit wide burst transfers.

To increase CPU write performance, and reduce memory bus contention, the CPU can write data to a buffer, which later can have its contents written, as 32-bit wide bursts, to main memory. This buffer, a DMA, stores up to 2 kB of data which can be written, as a sequential burst, to any location in main memory.

To use the DMA, the destination segment and pointer registers are set, bit masks in the control register allow byte enables to be set/cleared, and data is written to the write data register. Once all desired data has been written to the DMA, setting the start transfer bit in the control register initiates the direct memory access.

The DMA does not modify the data in the buffer, and the buffer is implemented as a circular queue, so the same 2 kB of data can be written multiple times, like to clear the frame-buffer. Listing D.4 demonstrates using the DMA to achieve functionality similar to the C programming

language library function *memcpy*.

DMA Port	Default Value	Register Function
0	0	Write data
1	0	Address pointer
2	0	Address segment
3	0	Control register

D.4 Assembly Coding Conventions

All 16 of the TTA16, and RISC16, registers that are within the RF are general purpose, but several programming conventions were used when writing routines. These are guidelines concerning register usage and are listed in Table D.5. Registers zero through three are for argument passing, and returning, and if more than four arguments are required, they should be passed on the stack. These registers can be modified without backing up first, but all other should be saved if modified. This is not strictly enforced but obeying them will ensure routines work as expected.

Stack operations are somewhat clumsy⁴, requiring the stack segment to be set, and explicit stack-pointer incrementing and decrementing. Passing arguments as pointers to arrays would probably be more efficient if possible.

Name	Free to Use?	Usage
r0	Y	Parameter 0 for function calls
r1	Y	Parameter 1 for function calls
r2	Y	Parameter 1 for function calls
r3	Y	Parameter 3 for function calls
r4-r11	B	Backup before modifying, restore afterwards
r12	N	Stack Segment
r13	N	Zero Register
r14	N	Stack Pointer
r15	N	Link Register

Table D.5: Register usage conventions.

⁴A future version of TTA16 would probably have a stack FU implemented using a BRAM, and with MFSRs as the incrementers and decrementers. This would be fast while using very few general-purpose FPGA logic resources.

D.5 Executing the Generated Object Code

The TTA16 assembler reads in an assembly file (‘.s’) and produces an ASCII text output file (‘.out’). The format of the output file contains many lines of the form:

< address >: < instruction >

Each field is ASCII encoded hexadecimal, without a prefix or suffix. An additional processing step is required to get this object code to the processor so it can be executed. There are three available options:

1. Include the object code in the synthesis step so that it is placed into a BRAM.
2. Append the object code to the ROM file so that it can be retrieved after FPGA initialisation, and transferred to main memory.
3. The host PC can transfer a binary image to the main memory, via the PCI bus.

The last two options require that the processor explicitly fetch instructions from main memory and then store them in the local instruction BRAM. The code can then be executed.

For the first option, placing code within the BRAMs at synthesis, a post-processor tool `out2v.py` was written that converts a ‘.out’ file into a file that is included into a Verilog source file containing a BRAM. This sets the contents of the BRAM at synthesis. This code is ready to be executed as soon as the FPGA has finished initialising and clocks stabilised. Because TTA16 has two BRAMs, one containing the upper 16-bits of each instruction, the other the lower 16-bits. The post-processor can generate two files to satisfy this. For the generated file `test.out` the following two commands are issued:

```
out2v.py -n 2 -s 0 test.out tta_asm0.v
out2v.py -n 2 -s 1 test.out tta_asm1.v
```

The two generated files can then be included, using the ‘`include`’ directive, in the initialisation field of a BRAM (see the source file `/rtl/cpu/tta16/tta16.v` for an example).

D.5.1 TTA16 Programming Example: ‘memcpy’

The data cache used with TTA16 is a write-through design, and since the SDRAM is shared with PCI and video redraw modules too, many write-through accesses will reduce memory throughput. This is because each write is a 16-bit atomic write, and the memory controller is most efficient with 32-bit burst reads and writes. To achieve high performance, the TTA16 can (and should) write blocks of data using the DMA controller.

The DMA controller is connected to the TTA16 via the I/O bus, not the memory bus, so TTA16 can write data to the DMA controller without causing memory contention with the other

modules. Once all desired data has been transferred to the DMA controller's local memory (2 kB), a write command can be issued which will cause all stored data to be written to the SDRAM. The write is a burst, and 32-bits wide, so the efficiency of this technique is good.

The following is a programming example showing how to program OpenVGA's DMA controller to perform a memory copy. The data cache will perform data prefetches using 32-bit bursts as well, so all memory accesses are 32-bit wide, burst transfers.

The following is a 'memcpy' routine, written in the C programming language, of how a naive memory copy routine could be written:

```
void* memcpy(void* dst, void* src, int n)
{
    while(n--)
        (int*)dst++ = (int*)src++;
    return dst;
}
```

Figure D.3: A memory copy routine that copies 'n' integers from 'src' to 'dst'.

Similar code written in TTA16 assembly language is shown in Figure D.4

```

;-----
; memcpy_dma - Optimised 'memcpy' using the DMA to burst write data to the
; SDRAM controller. There are five arguments, read seg+pointer, write seg+
; pointer, and number of words. The fifth value is TOS.
; This procedure does not check 'n' until after the DMA is setup, or whether
; the DMA is currently in use (since this CPU has no interrupts), i.e. it
; assumes that the programmer is not a moron.
;      r0      - src_ptr*
;      r1      - src_seg
;      r2      - dst_ptr*
;      r3      - dst_seg
;      [TOS]   - n
; Return value: none
memcpy_dma:    ; Setup the DMA controller.
               {
               ,
               ,
               ,1      }
               {
               ,
               ,com     ->msr ,0x50 } ; DMA addy
               {
               ,-1      ->sub ,
               ,\r14    } ; Start POP
               {3      ->wad ,
               ,\r13    ->mem ,      }
               {2      ->wad ,
               ,\r3     ->mem ,diff  }
               {com     ->rad ,
               ,diff    ->r14\ ,      } ; Fetch for POP
               {1      ->wad ,1      ->sub ,\r2     ->mem ,mem   } ; Clear masks
               {
               ,
               ,
               ,      }
               {
               ,
               ,diff    ->r2\ ,0      } ; Set read seg
               {
               ,
               ,com     ->msr ,\r1    }
               {
               ,
               ,
               ,      }
               ; Setup complete, now spam data at the DMA.
memcpy_10:     {\r0     ->rad ,-1      ->sub ,
               ,\r0    } ; read+inc src*
               {
               ,1      ->sub ,
               ,\r2    }
               {
               ,
               ,diff    ->r0\ ,mem    }
               {\r13    ->wad ,
               ,com     ->mem ,      }
               {memcpy_10->jnb ,
               ,diff    ->r2\ ,      }
               {
               ,
               ,
               ,-128   }
               {3      ->wad ,
               ,com     ->mem ,      } ; start DMA write
               {
               ,
               ,
               ,      }
               {\r15    ->bra ,
               ,
               ,
               ,      }
               {
               ,
               ,
               ,
               ,      } ; BDS

```

Figure D.4: A memory copy routine, written in TTA16 assembly language, using the DMA module of OpenVGA.

Appendix E

RISC16 Assembly Language Programming Guide

RISC16 is a simple 16-bit RISC processor and this appendix is an introductory programming guide for it. Contained in this guide are some code fragments and simple routines, and an explanation of the RISC16 registers and instructions. Related topics that are covered in significant detail in the TTA16 programming guide (see Appendix D) will be covered only briefly here.

E.1 Tools: The RISC16 Assembler

File:	r16asm.py	
Description:	A simple RISC16 assembler.	
Related Files:	/src/r16asm.py, /src/CodeCleaner.py, /src/AsmParse.py, /src/Emit.py	
Testing Files:	/src/fw_risc16/*.S	
Author:	Patrick Suggate	License: GPL

An assembler, written in Python, is used to assemble RISC16 assembly files. Since the RISC16 has just one execution thread, and some special instructions, the TTA assembler was not ideal for RISC16. The biggest problem was using the *i12* instruction with labels (for long jumps). This was solved with *r16asm* as it can evaluate Python mathematical expressions involving bitwise operators, arithmetic operators, numbers, and named constants.

Example assembler usage:

```
r16asm.py in_file.S -o out_file.out
```

And then to make a .v file to include in synthesis:

```
out2v.py -n 1 -s 0 out_file.out risc_asm.v
```

E.2 Instruction Format Overview

Most RISC16 instructions are of a typical two-operand RISC form, though there are several single-operand and three-operand instructions too. The supported RISC16 instruction formats are listed in Figure E.1.

	15	12	8	4	0
rri	OP	SF	RD	RS	IMM4
rr	0x0	SF	RD	RS	FN CR
ri	0x1	SF	RD	FN CR	IMM4
i12	0x7	X	IMM12		
bx	0x6	CND	IMM10		

Figure E.1: RISC16 instruction formats.

The following is an **rr**-format, (see Figure E.1) subtract instruction where **r0** gets the value of **r0** minus **r2**:

```
sub r0, r2
```

The destination register (in this case **r0**) precedes the source registers and immediate constants in the operand list of RISC16 instructions. The two-operand RISC16 instructions use the first operand as a source and destination.

This is an example of a three operand (**rri**-format) instruction:

```
subi r0, r0, 1
```

With this format, all three operands are specified, even though **r0** repeated in this case (it is typically different and is encoded in a separate bit-field). This instruction subtract 1, a 4-bit, signed immediate constant, from **r0**. This is equivalent to a decrement of **r0**, and it has short-hand version that is also accepted by the assembler:

```
dec r0
```

This instruction has exactly the same function and encoding as the preceding **subi** example and this is just a convenient short-hand form. There is also a corresponding **inc** mnemonic, which increments a register, as well.

Most ALU operations do not have a three operand form so to use immediate constants, there is another instruction format, **ri**:

```
sub r0, 1
```

The result, stored in **r0**, is 1 minus **r0**. Note that this is the reverse order to **subi**. This allowed a negate short-hand instruction to be easily implemented too, simply subtract the desired register from zero.

If immediate constants are required that are greater than can be stored as a signed, 4-bit number (-8 to 7), an instruction prefix can be used:

```
i12 0x123
sub r0, 0x4
```

This instruction sequence uses the 16-bit value 0x1234 (the '0x' is to indicate that the following value is represented as hexadecimal) as the immediate constant. This immediate prefix can be used with all **rrr** instructions, like **subi**, and **ri** instructions, like the **sub** example above.

Conditional jump instructions are of the form:

```
jxx somewhere
```

The **jxx** represents any one of the eight supported jump types (see E.2). The **somewhere** is a destination address (within the instruction SRAM), and it can be a number, a label, or a simple mathematical expression (which can contain labels as these will be evaluated).

Load and store instructions are of type **rrr** and have several syntaxes. An example of the syntax used later in this appendix is:

```
lw.d r0, [r1-3]
```

The mnemonic **lw.d** means load-word, using the **d** memory segment¹. The square brackets following **r0** are simply to help the programmer remember that the second and third operands are used to calculate a memory address, from which a 16-bit word is loaded. The same instruction syntax of **subi** can be used instead if desired, as these instructions share the same encoding. The same load instruction could then be rewritten as:

```
lw.d r0, r1, 3
```

¹There are two segment registers, **d** and **s**. Their names were chosen to represent either source and destination segments, or stack and data segment registers, depending on the situation. This is purely a way to remember their names, and make code more readable, as these registers function identically.

Name	Free to Modify?	Usage
r0	Y	Parameter 0 for function calls
r1	Y	Parameter 1 for function calls
r2	Y	Parameter 1 for function calls
r3	Y	Parameter 3 for function calls
r4-r11	B	Backup before modifying, restore afterwards
r12	N	Stack Segment
r13	N	Zero Register
r14	N	Stack Pointer
r15	N	Link Register

Note that the sign has changed, but this is exactly the same operation. This is because the ALU performs a subtract to calculate the memory address for load and store instructions. This has to be remembered when using the `i12` instruction prefix as well.

E.2.1 RISC16 Registers

RISC16 has 16 general purpose registers. While they all are functionally identical, by convention several of the registers have function. There is no hardware restriction on which registers can be used, it is simply meant to increase interoperability between assembly routines.

E.2.2 ALU Instructions

The general form of RISC16 arithmetic instructions is two input values, one always a register, and the other a register or immediate, and a destination register. For most arithmetic instructions (except `subi` and `subic`) the destination register has to be one of the source registers as well.

The list of available ALU instructions is shown in Table E.1. Many of these instructions have two forms, one which modifies the processor condition-code flags, and one that does not. It is usually as simple as appending a `c` to the end of an ALU instruction to cause it to modify the flags. Bitwise operations only set the Zero Flag, and subtract-based operations (`sub`, `sbb`, `inc`, `dec`, `subi`, and `neg`) only set the Negative Flag (NF) and Borrow Flag (BF). The `cmp` instruction modifies all of these three flags².

Some example ALU instructions:

²The `cmp` instruction performs both an XOR and a subtract on the operands.

```

neg    r0
and    r1, -4
cmp    r2, 0x2
cmp    r3, r15
incc   r4
subi   r5, r6, 0

```

Normal	Set Flags	Encoding	Description
subi	subic	rrr	Subtract, 3-operand format
inc	incc	rrr	Increment
dec	decc	rrr	Decrement
sub	subc	rr, ri	Subtract
sbb	sbbc	rr, ri	Subtract using the borrow flag
neg	negc	ri	Negate
N/A	cmp	rr, ri	Compare
and	andc	rr, ri	Bitwise logical AND
nand	nandc	rr, ri	Bitwise logical NAND
or	orc	rr, ri	Bitwise logical OR
xor	xorc	rr, ri	Bitwise logical XOR
mull	N/A	rr, ri	Multiply, store low 16-bits
mulh	N/A	rr, ri	Multiply, store high 16-bits

Table E.1: RISC16 ALU instructions that can optionally set condition code flags.

E.2.3 Processor Condition Code Flags

RISC16 has three processor state flags which are set by the ALU. These are borrow, negative, and zero flags (BF, NF, ZF). When an instruction has an additional *c* appended to the instruction mnemonic the processor condition codes are then to be modified by that instruction. The processor flags are used with conditional jumps.

E.2.4 Multiply Instructions

Only 16-bit unsigned multiplication is directly supported, producing a 32-bit product. Only either the upper or lower 16-bits of the 32-bit product can be written back to the RF, not both. The *mull* instructions stores the lower 16-bits of the product in the destination register, the *mulh* stores the upper 16-bits. The following assembly fragment is an example of the syntax:

```

mull r0, 0x4

```

E.2.5 Program Flow Control

RISC16 takes three cycles to perform a jump or branch, but unconditional branches typically use the `i12` instruction, adding one more clock cycle. This is to set the upper bits for the branch address. This approach does not give a significant performance penalty since conditional branching is far more common than unconditional jumps [28].

Conditional jumps have the 10-bit destination address encoded into the instruction word, requiring just one instruction. Register indirect branches takes just one instruction as well, since the entire destination address is stored in a register. All of the supported jump and branch operations are listed in Table E.2.

Mnemonic	Encoding	Description
<code>jnz</code>	<code>bx</code>	Jump if ZF not set
<code>jz</code>	<code>bx</code>	Jump if ZF set
<code>j1</code>	<code>bx</code>	Jump if NF set
<code>jg</code>	<code>bx</code>	Jump if NF not set
<code>jb</code>	<code>bx</code>	Jump if BF set
<code>jbe</code>	<code>bx</code>	Jump if BF and ZF set
<code>ja</code>	<code>bx</code>	Jump if BF not set
<code>jae</code>	<code>bx</code>	Jump if BF not set and ZF set
<code>br</code>	<code>ri, rr</code>	Unconditional branch
<code>brl</code>	<code>ri, rr</code>	Branch with link

Table E.2: RISC16 branch and jump instructions.

E.3 Loads and Stores

RISC16 has a minimalist FU that initiates transfers over the Wishbone bus, either data loads or stores³. Main memory is accessed through a data cache connected between the processor and the system Wishbone bus.

There are two instructions for initiating Wishbone bus operations, load word and store word (`lw` and `sw` respectively). Each of these instructions causes the processor to stall until an acknowledge is received from the Wishbone bus. This potentially wastes cycles, but the trade-off is that data forwarding is not needed for memory operations, so no dependency checking logic is needed either.

³Since the processor operates at much higher frequencies than the system Wishbone bus, a data cache is used which also functions as a synchroniser. It additionally changes the data width from 16-bits, of RISC16, to 32-bits, of the system bus, and vice versa as well.

The ALU, specifically the subtract FU, is used to calculate the final address from the supplied arguments. The format is always register minus an immediate, but this can be zero (or the zero register which is `r13` by definition). Since the ALU operation is subtraction, immediate operands have to be negated when calculating the address. Listing E.4 contains examples of these instructions.

The ALU bypass register is used for sending data to the wishbone bus. The only valid source for this bypass register is `RS`, and there is no data forwarding provided by hardware for this register either. Any potential data hazards that are detected will result in pipeline interlocking.

E.3.1 Memory Segment Registers

RISC16 has two parameterisable, each up to 16-bits in width, MSRs (Memory Segment Registers) to allow memory addresses greater than 16-bits wide to be supported. The memory segment value is concatenated with the memory pointer value to give a total width of upto 32-bits, and with two bytes per address, this means up to 8 GB is addressable with this architecture.

The two MSRs are called the Data Segment (DS) and the Stack Segment (SS) registers which indicates their possible uses, though no such limit upon such use is incorporated within the hardware. When performing large block copies, for example, they both could be used as data registers.

By default, though this is set by the assembler/convention, not by the hardware, the DS is used with load and store instructions. To use SS, instead of DS, the suffix `‘.sf’` is appended to the load/store instruction.

E.3.2 OpenVGA MMIO Device Segments

RISC16 and TTA16 are interchangeable when synthesising OpenVGA, and all peripherals external to the processor are therefore identical. The MMIO devices are the same for both RISC16 and TTA16. For a complete list of MMIO devices, see Appendix D, and the *memcpy* listing later in this section demonstrates how to access MMIO devices (see Figure E.4).

E.4 Programming Example: ‘memcpy’

This code has the same functionality as the TTA16 assembly code in Listing D.4. The routine copies a block of memory from the source to the destination. The DMA is used to improve memory throughput as it allows the memory controller to operate in 32-bit wide, burst-mode reducing the cycles lost due to multiple memory access-latency penalties.

```

;-----
; memcpy_dma - Optimised ‘memcpy’ using the DMA to burst write data to the
; SDRAM controller. There are five arguments, read seg+pointer, write seg+

```

```

; pointer, and number of words. The fifth value is TOS.
;      r0      - src_ptr*
;      r1      - src_seg
;      r2      - dst_ptr*
;      r3      - dst_seg
;      [TOS]   - n
; Return value: none
DMA_start:    equ      0x0080
memcpy_dma:   ; Setup the DMA controller.
              sw.s      [r14], r4          ; PUSH r4, reg needed
              i12      -DMA_SEG>>4
              subi     r4, r13, -DMA_SEG
              nop
              nop                      ; No bypassing on
              msr.d    r4                ; MSR (the 'dat' path)
              msr.d    r4                ; Set output seg to DMA
              sw.d      [r13+1], r2       ; Setup DMA, lo-word of addr
              sw.d      [r13+2], r3       ; hi-word of addr
              sw.d      [r13+3], r13      ; clear bit-masks
              lw.s      [r14-1], r2       ; Get 'n' from stack
              msr.s     r1                ; Set src seg
memcpy_lo:    lw.s      [r0], r3
              decc      r2
              inc       r0                ; Don't modify flags
              sw.d      [r13+0], r3       ; Write to DMA
              jnz       memcpy_lo

              i12      -DMA_start>>4
              subi     r3, r13, -DMA_start
              msr.s     r12              ; Restore SS and POP r4
              lw.s      [r14], r4         ; While 'r3' writes back
              sw.d      [r13+3], r3       ; Start DMA xfer

              ret

```